

Bachelor Thesis

Julius-Maximilians-
UNIVERSITÄT
WÜRZBURG

Design, Implementation and evaluation of different strategies for playing Pokémon battles

Julian Schubert

Institute for Computer Science
Chair for Computer Science VI

Prof. Dr. Frank Puppe

First Reviewer

Jonathan Krebs

First Advisor

Submission

XX. Month 20YY

www.uni-wuerzburg.de

Abstract

(Pokémon is a popular role playing video game franchise) (Can i just copy the first sentence of another abstract?) (This work gives a detailed introduction to Pokémon battles and summarizes research done so far) (Additionally, a RuleBased player is implemented and evaluated against the state of the art search based agent) (Furthermore challenges for evaluation of random battles are covered)

ToDo
ToDo
ToDo
ToDo
ToDo
ToDo

Zusammenfassung

Contents

1. Introduction	3
2. Background	7
2.1. Basic rules	7
2.1.1. Types	7
2.1.2. Moves	7
2.1.3. Pokémon	8
2.1.4. Switching	10
2.1.5. Status condition	10
2.1.6. Items	10
2.1.7. Field effects	11
2.1.8. Order of events	13
2.1.9. Damage calculation	13
2.1.10. Effective Stats	15
2.2. Dynamaxing	15
2.3. Hazards	15
2.3.1. List of entry hazards	15
2.3.2. Hazard counterplay	17
2.4. Showdown random battles	17
2.4.1. Sets	17
2.4.2. Items	18
2.5. Pokémon Matchups	19
2.5.1. Check and Counter	19
3. Related Work	21
3.1. Baseline Agents	21
3.2. Breadth-first search	21
3.3. Minimax	21
3.4. Rule based agents	22
3.5. Other approaches	22
3.6. Self-Play Policy Optimization Approach	23
3.6.1. State of the art	25
3.7. Supervised Approach	27
4. Approach	29
4.1. Communication with Pokémon Showdown	29
4.2. Gathering Information about the enemy Pokémon	29
4.3. Scoring the current game state	30
4.4. Stages of the game	31
4.5. Determining matchups	32
4.6. Predictions	33
4.7. Dynamaxing	33

5. Evaluation	35
5.1. Challenges for evaluation	35
5.1.1. Randomness of battles	35
5.1.2. Evaluation against human opponents	36
5.2. Agents	36
5.2.1. HerrDonner	36
5.2.2. HerrGewitter	38
5.3. Results	40
5.3.1. Ranked results	40
5.4. Investigating Replays	42
5.4.1. Pro Replays	42
6. Conclusion	45
List of Figures	47
List of Tables	49
Listings	51
Acronyms	53
Bibliography	55
Appendix	59
A. First Appendix Section	59

(Fix quotes: Quotes were messed up when converting to apalike. Dates need to be added. **ToDo**
Quote style: The agent developed by AUTHOR, YEAR, manages... Results are surprising
(AUTHOR, YEAR)) (Avoid usage of I / WE. Use indirect speech instead) **ToDo**

1. Introduction

Pokémon (an abbreviation for **Pocket Monsters**) is a media franchise managed by *The Pokémon Company*, a company founded by *Nintendo*, *Game Freak* and *Creatures* (Wikipedia, 2022b). Figure 1.1 shows an image of *Pikachu*, one of the most popular Pokémons. Among multiple movies and series, a large variety of Pokémons games was released, starting with *Pokémon Red* and *Pokémon Blue* which were released on 28. September 1998, also known as generation 1 (gen1)-games. In order to promote trading between players, *Nintendo* published two very similar games at the same time. (**Fix Pikachu image, properly cite ingame screenshots**) Both releases differ in the available Pokémons and some minor changes to the story. Therefore, in order to collect all available Pokémons, players either had to buy two copies of the game or trade with a player that owns the counterpart. As of writing, there are eight major releases, also called *mainline games* with the latest being *Pokémon Sword* and *Pokémon Shield*. While the graphics evolved a lot over the years (see figures 1.2, 1.3, 1.4 and 1.5), the key concept of the game remained mostly unchanged. The player starts his journey in his hometown to become the *Pokémon Champion*, which is the highest known level of rank for a Pokémon trainer (https://bulbapedia.bulbagarden.net/wiki/Pok%C3%A9mon_Champion). In order to achieve this goal, the player has to catch wild Pokémons and use them to create a team of up to six individual Pokémons which he then trains to unleash their full potential. This Thesis focuses exclusively on the battling aspect of the game as there are detailed lists of locations and secrets there is to explore within the games. Pokémons battles are turn based where both players, unlike in for example chess, make their decisions at the same time. While the core battle mechanics are very simple, the game provides a lot of depth which led to the formation of a strong competitive battling scene. As catching training Pokémons to a



ToDo

Figure 1.1.: Pikachu, one of the most popular Pokémons (Fandom, 2016)

ToDo

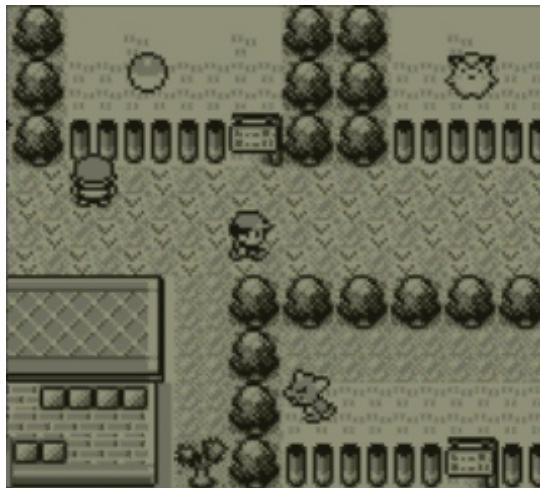


Figure 1.2.: Exploring the map in
Pokémon Red



Figure 1.3.: Fighting another trainer in
Pokémon Red

Image source: nintendo.de

Release: 1996



Figure 1.4.: Exploring the map in
Pokémon Sword.
Image source: pokemon.com



Figure 1.5.: Fighting another trainer in
Pokémon Sword
Image source: nintendo.de

Release: 2019

To Do level where they are competitive viable is a time-intensive task, battle simulators such as the open source project *Pokémon Showdown* ([Cite Showdown Link](#)) have arisen. On these fan made platforms, players have access to all available Pokémons and can compete against each other in a large variety of formats. Figure 1.6 ([Move figure to appendix](#)) shows a battle between two players on *Pokémon Showdown*. The screenshot was taken from a *Gen 8 random battle* between *Buckfae* and another player¹. The screenshot was taken at the start of turn 19. Currently, the *Mewtwo* of *Buckfae* is at full health (indicated red) while the opposing *Kommo-o* has 33% of hit points (hp) left. Below the avatars of both players, the team is displayed, marked yellow for player two. While player one has three Pokémons remaining, his opponent has four members alive. The Pokéball (bottom right) indicates a to the player unknown Pokémons. Moreover, the blue box highlights the status modifications of the enemy Pokémons. The effects of status conditions will be covered in detail in section ([Link to status](#)). Below the game window, the possible choices of the player are displayed. *Mewtwo* has access to the moves *Fire Blast*, *Recover*, *Psystrike* and *Nastyplot*. In addition to *Dynamaxing* which will be covered in ([Link to dynamax](#)) the player also has

¹The other name was blurred

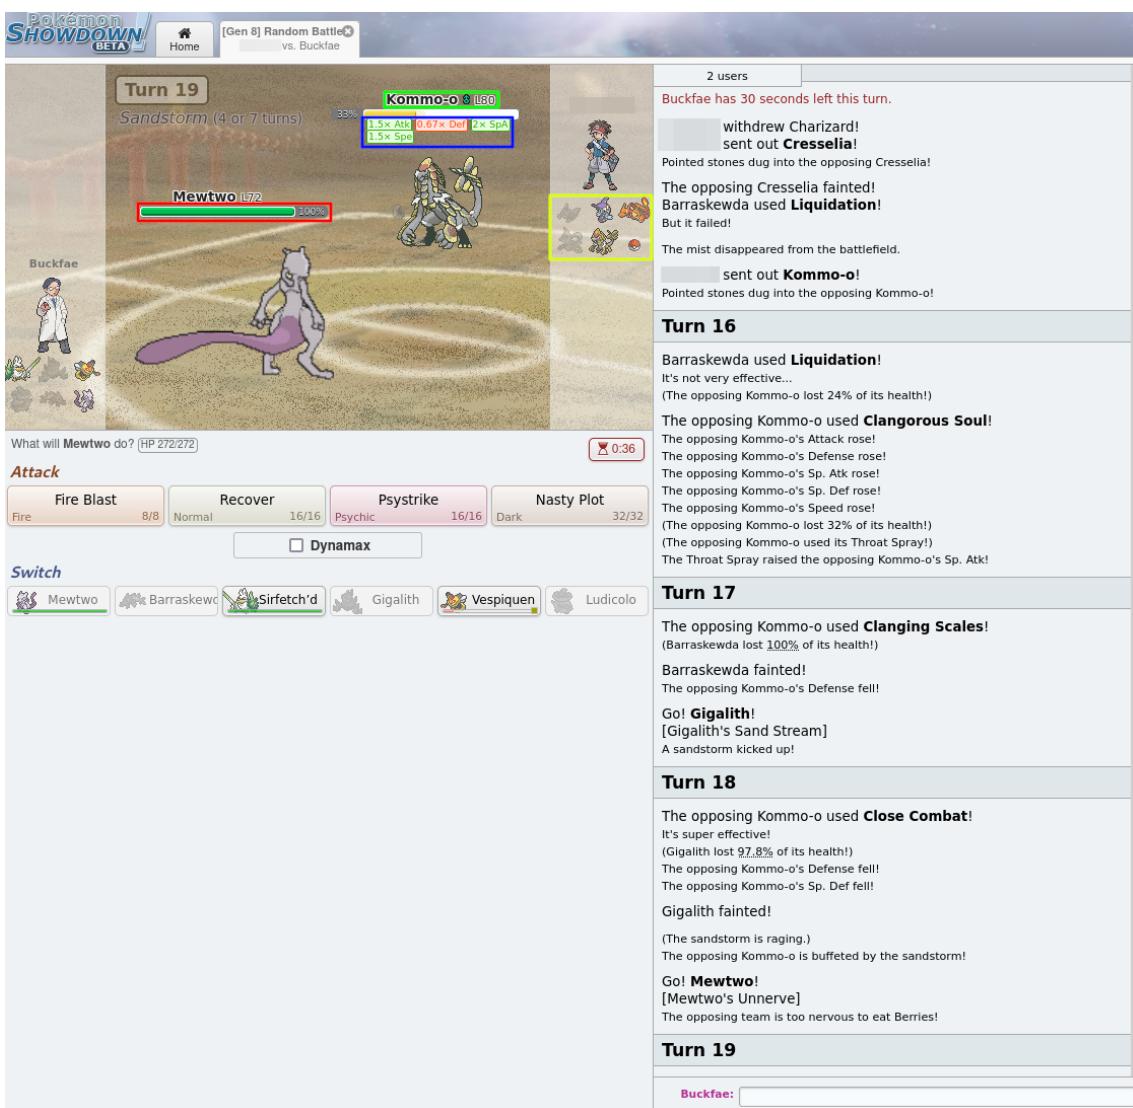


Figure 1.6.: Battle between two players on *Pokémon Showdown*

the option to switch to either *Sirfetch'd* or *Vespiquen*. Lastly, on the right-hand side a log of the previous turns can be found. In addition to the moves a Pokémon can use, a Pokémon has one ability and can hold an item that yields advantages in battle.

The high dimensional feature vectors of states in combination with the complex rules of the game result in Pokémon presenting an interesting challenge for AI to tackle.

This thesis gives an introduction to the main mechanics of the Pokémon genre while no prior knowledge of the game is required. Chapter Related Work introduces and compares different approaches and provides an in depth summarization of the current state-of-the-art battling AI. Furthermore, an agent combining previously existing rule- and Minimax-based implementations is developed to improve long term planning. Challenges of evaluating Pokémon AI and the results of this research are presented in Evaluation.

2. Background

2.1. Basic rules

One of the key aspects of the Pokémon game is to battle other Pokémon. In the mainline games, you can have up to six Pokémon in your team, also known as party. There is the option to swap a Pokémon with another Pokémon, but you can not have more than six Pokémon at any point in your team. When playing the original games, you explore the world to find more Pokémon and use your team to defeat wild Pokémon and other Pokémon trainers. This thesis focuses on random battles taking place on Pokémon Showdown. In a random battle, both competitors get a team of six random Pokémon. At the start of the battle, you know each of your six Pokémon but only the currently active enemy Pokémon. In every turn, both players can choose to either use a move of their currently active Pokémon or switch their active Pokémon to another Pokémon. Moves can either deal direct damage to the enemy Pokémon or yield other advantages like increasing the damage dealt by the next move. Moves will be covered in more detail in section 2.1.2. Each Pokémon has an amount of hp. The hp of a Pokémon can be dropped by attacking it with a move. If the hp of a Pokémon drops to zero, it faints and can not be used in this battle anymore. A player wins, once all enemies fainted.

Note that in the mainline games there is the possibility to heal or even revive a fainted Pokémon during battle using *Healing Items* like *Revive* or *Hyper Potion*. In competitive Play, only *Held items* like *Leftovers* are allowed. Items will be explained in depth in section 2.1.6.

2.1.1. Types

Pokémon implements a rock-paper-scissors (rps)-like system. Each Pokémon has either one or two of 18 types. Like *rock* beats *scissors* but loses to *paper*, *Fire*-type Pokémon have an *advantage* against *Grass*-type Pokémon but are on a *disadvantage* against *Water*-types. This is also called a type being *strong / weak* against another type. Figure 2.1 shows how different Pokémon types interact with each other. Unlike in rps, type modifiers will be multiplied if a Pokémon has two types. For instance, a *Fire*-type attack will deal 4 times the damage against *Parasect* as *Parasect* has the types *Grass* and *Bug* (Veekun, 2021).

2.1.2. Moves

Moves can be split up into three categories: *Physical*-, *Special*- and *Status*-Moves. While *Physical*- and *Special*-Moves usually deal damage to the opponent Pokémon, *Status*-Moves

DEFENSE → ATTACK ↴	NOR	FIR	WAT	ELE	GRA	ICE	FIG	POI	GRO	FLY	PSY	BUG	ROC	GHO	DRA	DAR	STE	FAI
NORMAL													½	0			½	
FIRE		½	½		2	2						2	½		½		2	
WATER		2	½		½				2				2		½			
ELECTRIC			2	½	½				0	2					½			
GRASS		½	2		½			½	2	½		½	2		½		½	
ICE		½	½		2	½			2	2				2		½		
FIGHTING	2				2			½			½	½	½	2	0	2	2	½
POISON					2			½	½					½	½	0	2	
GROUND		2		2	½			2		0		½	2			2		
FLYING			½	2		2					2	½					½	
PSYCHIC					2	2					½				0	½		
BUG		½			2		½	½		½	2			½	2	½	½	
ROCK		2				2	½		½	2		2				½		
GHOST	0									2			2		½			
DRAGON													2			½	0	
DARK							½				2		2		½		½	
STEEL		½	½	½		2						2				½	2	
FAIRY		½				2	½						2	2	½			

Figure 2.1.: Advantages of different types (Database, 2022)

can for example change the weather, which plays a role in damage calculation explained in section 2.1.9, inflict status effects, raise or lower the stats of a Pokémon. In addition, a move also has exactly one of the 18 possible types. Often, *Status*-moves are also referred to as *buffing* moves if they grant a benefit to the user or as *debuffing* if they have a negative effect on the opponent.

2.1.3. Pokémon

A key-concept of Pokémon battles are the *stats* of a Pokémon. Figure 2.2 displays information about possible stat-combinations of *Charizard*.

Explanation of stats

HP: The hp determines how much damage a Pokémon can receive before fainting.

Attack: The attack stat (atk) determines how much damage a Pokémon will deal when using a *Physical*-Move.

Defense: The defense stat (def) determines how well a Pokémon can resist against *physical* attacks.

Sp. Atk: The special Attack stat (spa) determines how much damage a Pokémon will deal when using a *Special*-Move.

Sp. Def: The special Defense stat (spd) determines how well a Pokémon can resist against special attacks.

Speed: The speed stat (spe) determines how fast a Pokémon can act. Usually, the Pokémon with a higher spe will move before the slower one. The exact order of actions in battles is covered in section 2.1.8. (**Cover evasion / accuracy, context to showdown**)

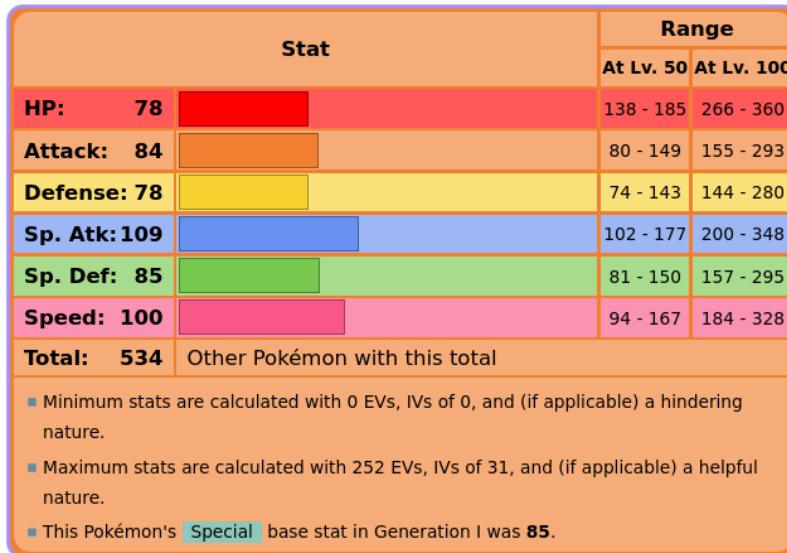


Figure 2.2.: Possible stats of *Charizard* (Bulbapedia, 2022)

Determination of stats

The total stat of a Pokémon is calculated as described in equation 2.1 and equation 2.2 (Bulbapedia, 2021w).

$$HP = \left\lfloor \frac{(2 \times \text{Base} + IV + \lfloor \frac{EV}{4} \rfloor) \times \text{Level}}{100} \right\rfloor + \text{Level} + 10 \quad (2.1)$$

$$\text{OtherStat} = \left\lfloor \left(\frac{2 \times \text{Base} + IV + \lfloor \frac{EV}{4} \rfloor) \times \text{Level}}{100} + 5 \right) \times \text{Nature} \right\rfloor \quad (2.2)$$

Base: Refers to the base stat of a Pokémon. Two Pokémons of the same species will always have the same base-stats. As seen in figure 2.2, a *Charizard* will always have a base-atk of 84.

Level: As mentioned in section 2.1, the goal of the mainline games is to create a team of six Pokémons and to make that team stronger by fighting other Pokémons. If a Pokémon defeats enough other Pokémons, it grows a Level. The maximum level of a Pokémon is 100. In Pokémon Showdown, the level of a Pokémon is set at the start of the battle and won't increase (Mitarai, 2021).

Nature: A Pokémon has a nature. Most natures enhance the growth of one stat, while hindering the growth of another. After all other calculations are finished, the stat that the Nature enhances will be 100% of what it would be without the Nature, and the stat hindered will be 90% of its normal value (Bulbapedia, 2021w). In this thesis nature can be neglected as all Pokémons in random battles have a neutral nature, meaning no stat is enhanced or hindered (Mitarai, 2021).

IV: Refers to the individual values (iv) of a Pokémon. These cause two Pokémons of the same species to have different Stats (Bulbapedia, 2021w). Pokémons in Pokémon Showdown will always have the best possible iv stat, 31, unless it is a disadvantage for the Pokémon, then it will be zero (Mitarai, 2021). For example, having a high spe is undesirable for a Pokémon that uses the move *Trick Room* as this move reverses the move order so that Pokémons with a lower spe stat attack first while those with a higher spe will attack last.

EV: These are the effort values (ev) of the Pokémon. ev are what causes a trained Pokémon to have higher stats than an untrained counterpart of the same level. For every 4 ev gained,

a level 100 Pokémons will have 1 extra point in the given stat. A Pokémon can earn up to 510 ev, but can not have more than 255 ev in a single stat (Bulbapedia, 2021w). Random Pokémons on Showdown will always have 85 ev in each stat, or 0 in the case that having a high stat being detrimental (Mitarai, 2021).

2.1.4. Switching

Instead of using a move with the current Pokémon, the player also has the option to switch out the active Pokémon for another Pokémon in his party. Switching always takes place before the execution of moves. However, the player does not know whether the opponent is switching or using a Move. Therefore, if the player decides to switch out a non-fainted Pokémon, the enemy gets to use his move on the new Pokémon. If a Pokémon faints, the player has to switch in a new Pokémon and then the next turns start. This means that the opponent gains a one turn advantage if the player decides to switch out a healthy Pokémon, but won't get to attack an additional time if the Pokémon was defeated.

2.1.5. Status condition

Moves may inflict so-called *status conditions* that affect a Pokémon negatively. The most important status conditions are

- **Burn:** If a Pokémon suffers from the status condition burn (brn), it will lose 1/8 of its total hp every turn. In addition to that, a burned Pokémon will only deal half as much damage when using a *physical* move.
- **Freeze:** If a Pokémon suffers from the status condition freeze (frz) it won't, with a few exceptions, be able to use moves.
- **Paralysis:** If a Pokémon suffers from the status condition paralysis (par) it won't be able to use the selected move 25% of the time and their speed is halved.
- **Poison:** If a Pokémon suffers from the status condition poison (psn) it will, with a few exceptions, take damage equal to 1/8 of its total hp at the end of every turn. A Pokémon can also be *badly poisoned*. Badly poison initially inflicts damage equal to 1/16 of the Pokémon's maximum hp, with the damage inflicted increasing by 1/16 each turn. This means that the Pokémon will take 2/16 damage on the second turn, 3/16 on the third turn (and so on).
- **Sleep:** If a Pokémon suffers from the status condition sleep (slp) it won't be able to use moves, except *Snore* and *Sleep Talk*. In the mainline games, sleep lasts randomly between one and three turns. However, in Pokémon Showdown a Pokémon will *always* be asleep for exactly two turns.

At any point, a Pokémon can only suffer from one status condition at a time, this means that a burned Pokémon can not fall asleep. While a sleeping Pokémon wakes up after a few turns and a frozen Pokémon is unfrozen after a few turns, the other status conditions remain until they are cured. A status condition can either be cured by abilities like *Natural Cure*. This ability heals any status condition when the Pokémon is switched out. Items can also remove negative effects, for example, *Lum Berry* cures all status conditions but gets consumed in the process. ([How long are Pokémons asleep / frozen?](#))

ToDo

2.1.6. Items

A Pokémon can also hold an Item that yields benefits in battle. There are various purposes that items can fulfill. For example, the item *Life Orb* boosts damage dealt by the holder's

damaging move by 30%¹, but the holder takes damage equal to 10% of its maximum hp after it uses a damaging move² (Bulbapedia, 2021r). *Leftovers* restore 1/16 of the holder's maximum hp³ at the end of each turn whereas the item *Air Balloon* makes the holder *ungrounded*, which means that the holder is immune to *Ground*-type moves as well as several related effects (Bulbapedia, 2021a). The generation of items in Pokémon Showdown is described in more detail in section 2.4.2.

Important items

In this section, a quick introduction to the most important items is given.

- **Choice Band:** When held by a Pokémon, this item boosts the atk by 50%, but only allows the use of the first move selected. This effect resets when the holder is switched out (Bulbapedia, 2021d).
- **Choice Scarf:** When held by a Pokémon, this item boosts the spe by 50%, but only allows the use of the first move selected. This effect resets when the holder is switched out (**Bulbapedia:ChoiceScarf**).
- **Choice Specs:** When held by a Pokémon, this item boosts the spa by 50%, but only allows the use of the first move selected. This effect resets when the holder is switched out (Bulbapedia, 2021e).
- **Heavy-Duty Boots:** The holder is unaffected by the effects of entry hazards. Entry hazards are described in 2.3.
- **Assault Vest:** Raises the holders spd by 50% but also prevents the holder from selecting any status move⁴ (Bulbapedia, 2021b).
- **Focus Sash:** If the holder has full hp and is hit by an attack that would otherwise cause fainting, it survives with 1 hp (Bulbapedia, 2021l).

2.1.7. Field effects

There are multiple *field effects* that affect combat. Field effect are different from items as they are not tied to a specific Pokémon but rather always effect the currently battling Pokémon. There are field effects that only target the side of one player as well as effects that target both sides.

Terrain

Terrain is set up by the respective move with identical name and last for five turns and effects both players. All of them are beneficial to *grounded* Pokémon. A Pokémon is *not grounded* if any of the following conditions apply: The Pokémon

- has the *Flying*-type
- has the Ability *Levitate*
- is holding the item *Air Balloon*
- is under the effect of *Magnet Rise* or *Telekinesis*.

Grounded Pokémon are with a few exceptions those Pokémon, that are not *ungrounded*. A Pokémon will be grounded if any of the following conditions apply:

¹This boost is approximated as $5324/4096 \approx 1.29980$

²Rounded down, but not less than 1

³Rounded down, but not less than 1

⁴Except *Me First*

- The Pokémon is holding an *Iron Ball*
- The Pokémon is under the effect of *Ingrain*, *Smack Down* or *Thousand Arrows*.
- The *Field effect Gravity* is in effect.

More information about grounding can be found at (Bulbapedia, 2021n) There are five different possible *terrain*-states.

- **None:** The default state, no other effects are applied.
- **Electric Terrain:** Grounded Pokémons can not fall asleep and the power of *Electric*-type moves is increased by 50%.
- **Grassy Terrain:** The HP of grounded Pokémons is restored by 1/16 of their maximum HP at the end of each turn. In addition, the power of *Grass*-type moves is increased by 50% and the moves *Earthquake*, *Magnitude* and *Bulldoze* halve in power.
- **Misty Terrain:** Protects all grounded Pokémons from status conditions (including confusion) (**Does confusion exist in Showdown?**) The power of *Dragon*-type moves is halved while in effect.
- **Psychic Terrain:** Prevents grounded Pokémons from being hit by priority moves. Priority moves will be covered in section 2.1.8. The power of *Psychic*-type moves is also increased

It is important to note that only one *terrain* can be active at a time, yet, *terrain* can coexist with other *field effects* like *weather*.

Weather

The *weather* is a set of mechanics that change the battle environment, activating abilities, modifying certain moves and potentially damaging the Pokémons in battle or affecting their stats. Only one type of weather may be present at a time, and only the most recent type of weather will take effect Bulbapedia, 2021ae. List of different *weather*-conditions:

- *Clear skies*: (also: *None*) Absence of weather, this is the default state
- *(Harsh-) Sunshine*: (also: *(Harsh-) Sunlight*) Strong sunlight shines on the battlefield
- *(Heavy-) Rain*: Rain falls on the battlefield.
- *Sandstorm*: Stinging sand whips across the battlefield. At the end of each turn, each Pokémon, with a few exceptions, takes damage equal to 1/16 of its maximum hp unless it is a *Rock*-, *Steel*- or *Ground*-type. ([https://bulbapedia.bulbagarden.net/wiki/Sandstorm_\(weather_condition\)](https://bulbapedia.bulbagarden.net/wiki/Sandstorm_(weather_condition)))
- *Hail*: Pelting hail falls on the battlefield. At the end of each turn, each Pokémon, with a few exceptions, takes damage equal to 1/16 of its maximum hp unless it is an *Ice*-type. ([https://bulbapedia.bulbagarden.net/wiki/Hail_\(weather_condition\)](https://bulbapedia.bulbagarden.net/wiki/Hail_(weather_condition)))

There are a few additional special weather conditions which will not be covered in this thesis as they only occur in very specific scenarios.

Other

The list below contains the remaining field effects that cover the entire field not covered so far:

- *Wonder Room*: swaps the def and spd of all Pokémons, but stat changes remain on their respective stat. ([https://bulbapedia.bulbagarden.net/wiki/Wonder_Room_\(move\)](https://bulbapedia.bulbagarden.net/wiki/Wonder_Room_(move)))

ToDo

ToDo

ToDo

- *Magic Room*: suppresses the effect of all items held by the Pokémons on the field ([https://bulbapedia.bulbagarden.net/wiki/Magic_Room_\(move\)](https://bulbapedia.bulbagarden.net/wiki/Magic_Room_(move)))
- *Gravity*: causes the field to undergo intense gravity which multiplies the accuracy stat of all Pokémons by 5/3 ([https://bulbapedia.bulbagarden.net/wiki/Gravity_\(move\)](https://bulbapedia.bulbagarden.net/wiki/Gravity_(move)))
- *Trick Room*: reverses the move order within each priority bracket. Move priority will be covered in section 2.1.8

(According to Bulbapedia: All last for 5 turns)

ToDo

ToDo

Single Side Effects

Among *Hazards* which will be covered in section 2.3, the most important single side effects are *Reflect* which halves the damage done to Pokémons on the given side by *physical* moves ([https://bulbapedia.bulbagarden.net/wiki/Reflect_\(move\)](https://bulbapedia.bulbagarden.net/wiki/Reflect_(move))) and *Light Screen* which works equivalent but for *special* moves ([https://bulbapedia.bulbagarden.net/wiki/Light_Screen_\(move\)](https://bulbapedia.bulbagarden.net/wiki/Light_Screen_(move)))

ToDo

2.1.8. Order of events

In battle, switching always is executed before moves meaning that if a player switches and his opponent chooses to use a move, the move will always hit the newly brought in Pokémons. Move order is determined based on *priority*.

Priority

Priority is a characteristic of moves, such that any move with a higher priority will always be performed first. When two moves have the same priority, the user's SPE stat will determine which one is performed first in battle. Each move has a hidden priority value in the game data, with values ranging from +5 to -7. The vast majority of moves have the standard priority value of 0. A move with a positive priority is called *priority move*. Moves that have the same priority are said to be in the same priority bracket. An example for a move with priority +2 is *Extreme Speed*, a *Normal*-type move with a base-power of 80 ([https://bulbapedia.bulbagarden.net/wiki/Extreme_Speed_\(move\)](https://bulbapedia.bulbagarden.net/wiki/Extreme_Speed_(move))). Having a *priority*-move can be useful if the own active Pokémon has a lower SPE-stat than the opponent with low HP as then the own Pokémon does not have to take an additional hit before defeating his enemy.

ToDo

2.1.9. Damage calculation

The damage dealt by a move mainly depends on the *level* of the Pokémon that uses the move, its effective Attack or Special Attack stat, the opponent's effective Defense or Special Defense stat and the move's effective power.

Precisely, the damage is calculated as follows (Bulbapedia, 2021i):

$$\text{Damage} = \left(\frac{\left(\frac{2 \times \text{Level}}{5} \right) \times \text{Power} \times A / D}{50} + 2 \right) \times \text{Targets} \times \text{Weather} \times \text{Badge} \times \text{Critical} \times \text{random} \times \text{STAB} \times \text{Type} \times \text{Burn} \times \text{other} \quad (2.3)$$

The only exception for this are moves that deal direct damage. A list of these moves can be found at (Bulbapedia, 2021c). In the following paragraph, parameters of the formula are covered in detail. An explanation of the *Type*-modifier is omitted since it is covered in 2.1.1

Level

Level refers to the level of the attacking Pokémon (Bulbapedia, 2021i). In Pokémon Showdown, the level is displayed next to the name of the Pokémon.

A / D

A is the effective Attack stat of the attacking Pokémon if the used move is a physical move, or the effective Special Attack stat of the attacking Pokémon if the used move is a special move.

D is likewise the effective Defense stat of the target if the used move is a physical move, or the effective Special Defense of the target if the used move is a special move (Bulbapedia, 2021i).

There are four moves that use stats from different categories, more Information can be found at (Bulbapedia, 2016).

Power

Power is the effective power of the used move. The *Base Power* of a move in Showdown can be seen when hovering over a move in the move list.

Note: The same move will always have the same base power. For example, *Fire Punch* will always have a base power of 75 (Bulbapedia, 2021k).

Weather

The *Weather* modifier is 1.5 if a *Water-type* move is used during *rain* or a *Fire-type* move during *Harsh Sunlight*. The modifier is 0.5 if a *Water-type* move is used during *Harsh Sunlight* or a *Fire-type* move during *rain* (Bulbapedia, 2021i).

Critical

In the latest generation, a critical hit (crit) deals 1.5 times the damage compared to a normal hit. If the crit rate is not increased, the chance of landing a crit is 1/24 (Bulbapedia, 2021h). Increasing crit rate, as well as other stats, will be explained in chapter 2.1.10.

Note: In earlier games, crits worked different, see (Bulbapedia, 2021h) for more details.

Random

Random is a random integer percentage between 85% and 100%. Because of this, the same move may deal different damage in the same scenario (Bulbapedia, 2021i).

STAB

STAB stands for *Same Type Attack Bonus*. It is a multiplier of 1.5 if the current move is of the same type as the attacking Pokémon. Otherwise, it is 1.0 (Bulbapedia, 2021i).

Type

This is the in type modifier described in section 2.1.1 (Bulbapedia, 2021i).

Burn

Burn is 0.5 if the attacking Pokémon is burned, and the attack is physical⁵. Otherwise, it is 1.0 (Bulbapedia, 2021i).

⁵This does not apply if the attacking Pokémon has the Ability *Guts* or the used move is *Facade*

Other

The *other* modifier is usually 1. A list of exceptions can be found at (Bulbapedia, 2021i).

2.1.10. Effective Stats

When a stat is used in a calculation in battle, a number of modifiers may be applied during the calculation. During a battle, a Pokémon's effective stats may be raised or lowered by certain moves abilities and held items. Some attacks may only have a chance of lowering stats, while certain abilities and held items may require a triggering event to activate stat modifications.

The modifiers conferred by most moves operate on a sliding scale of *stages*. When a given stat is raised or lowered, its current stage is increased or decreased by the amount dictated by the move; Up to a maximum of +6 or a minimum of -6. A given stat corresponds to a given multiplier that will modify the stat when it is used in calculations. The exact multipliers for stages are described in table 2.1 (Table / Text source:https://bulbapedia.bulbagarden.net/wiki/Stat#Stage_Multipliers

-6	-5	-4	-3	-2	1	0	1	2	3	4	5	6
2/8	2/7	2/6	2/5	2/4	2/3	2/2	3/2	4/2	5/2	6/2	7/2	8/2

Table 2.1.: Stage multipliers for atk, def, spa, spd and spe

(Other table for accuracy an evasion) When playing the mainline games, the player has to memorize the current stat changes of a Pokémon. In contrast to that, Pokémon Showdown displays the current stat multiplier of both Pokémon. Stat changes caused by boosts do not persist between switches, meaning that the stages for all stats of a Pokémon are reset to zero upon being switched in.

ToDo
ToDo

2.2. Dynamaxing

Dynamaxing is a temporary transformation affecting Pokémon that was introduced in Generation eight. Dynamaxing increases a Pokémon's size drastically, as well as changing the moves of the Pokémon and doubling their max and current hp, but can only be used once per battle and ends after three turns or if the Pokémon is switched out.

2.3. Hazards

An *entry hazard* is a condition that affects one side of the field. It causes any Pokémon that is sent into battle on that side of the field to be afflicted by a negative effect. Entry hazards are created by moves, usually status moves (Bulbapedia, 2021s).

2.3.1. List of entry hazards

Currently, there are five moves that create an entry hazard

Spikes

Spikes is a *Ground*-type entry hazard that causes the opponent to lose 1/8% of their maximum hp when they enter the field. This effect can be stacked up to three times. Two layers of spikes will deal 1/6% and three layers will deal 1/4% of the enemies maximum hp. Spikes are set by the move *Spikes* (Bulbapedia, 2021v).

Stealth Rock

The move *Stealth Rock* sets an entry hazard around the target Pokémon causing Pokémon on the target's field to receive damage upon being switched in. The amount of damage inflicted is affected by the effectiveness of the type *Rock* against the target. Unlike Spikes, this entry hazard does not stack. The damage taken from the victim's maximum is denoted in table 2.2 (Bulbapedia, 2021x).

Type effectiveness	Damage (Max. hp)
0.25x	3.125%
0.5x	6.25%
1x	12.5%
2x	25%
4x	50%

Table 2.2.: Damage dealt to Pokémon by Stealth Rocks (Bulbapedia, 2021x)

Sticky Web

The entry hazard set by the *Bug*-type move *Sticky Web* lowers the opponents speed stat by one stage upon switching in (Bulbapedia, 2021y).

ToDo

(Pokémon that are not affected by this)

Poison spikes

Poison Spikes set by the *Poison*-type move *Toxic Spikes* cause the opponent to become poisoned. If two layers of spikes are set, the Pokémon instead becomes badly poisoned (Bulbapedia, 2021aa).

ToDo

(Pokémon not affected)

Sharp steel

This entry hazard works very similar to Stealth Rock described in 2.3.1. However, Sharp steel can only be set by the *Steel*-type move *G-Max Steelsurge* which is the exclusive G-Max Move of *Gigantamax Copperhead* (Add source). The damage dealt by Sharp steel does not stack, the amount of damage dealt is based on the Type effectiveness of the *Steel*-type against the target. Exact damage modifiers can be found in table 2.3 (Bulbapedia, 2021m).

Type effectiveness	Damage (Max. hp)
0.25x	3.125%
0.5x	6.25%
1x	12.5%
2x	25%
4x	50%

Table 2.3.: Damage dealt to Pokémon by Sharp Steel (Bulbapedia, 2021m)

2.3.2. Hazard counterplay

There are some moves that can remove entry hazards. *Rapid Spin* (Bulbapedia, 2021u) removes entry hazards from the user's side of the field and *Defog* (Bulbapedia, 2021j) removes entry hazards on both sides of the field⁶. In addition, *Court Change* (Bulbapedia, 2021g) will exchange the entry hazards on each side of the field, along with other one-sided field conditions. If a grounded⁷ *Poison*-type Pokémon enters the battle, it will remove Toxic Spikes, described in 2.3.1, from its side of the field. Pokémon holding the item *Heavy-Duty Boots* (Bulbapedia, 2021p) are unaffected by entry hazards, but grounded *Poison*-type Pokémon can still remove Toxic Spikes even if they hold the boots (Bulbapedia, 2021s). Additionally, there are various other exceptions to hazards which won't be covered in detail here.

2.4. Showdown random battles

Pokémon Showdown features a large variety of game modes, the most popular ones being *Sw / Sh Singles OU* and *Sw / Sh Singles Random Battle*. *Sw / Sh* refers to Pokémon *Sword* and *Shield*, the latest Pokémon Games. These two games form the eighth generation of Pokémon games and are thereby also often referred as *Gen8*-games. As Pokémon Showdown solely focuses on battling every player has access to every available Pokémon. This mainly plays a role in the *OU*-format. In this game mode, each player knows all enemy Pokémon but does not know their exact builds, that is, the abilities, iv-, ev-combinations and items are unknown. As the success of a player in this format heavily depends on his team as well as which other teams are currently popular, this thesis will focus on random battles. In this format, each player gets a team of six random Pokémon and unlike in *OU* only knows the currently active opposing Pokémon. It is important to note that there are many possible builds for a Pokémon and therefore, information management plays a huge role. An example for this will be provided in section 4.2.

The development for battling bots is explicitly allowed and encouraged by the developers. Furthermore, there is no way to figure out whether the current opponent is a real human or an AI. This is beneficial as humans might behave differently when they know they are facing a machine potentially leading to invalid results.

2.4.1. Sets

As described in section 2.1.3, Pokémon created for random battles usually have 85 evs and 31 iv in every stat with a neutral nature, meaning a nature that does neither boost nor hinder any stat (Mitarai, 2021). There are some cases where a high stat is not beneficial, an example would be the move *Gyro Ball*. Unlike most moves, the *Base Power* of this move described in the damage calculation described in 2.1.9 is not a fixed value. It is determined as described in 2.4 (Bulbapedia, 2021o).

$$\text{BasePower} = \min \left(150, \frac{25 \times \text{CurrentSpeed}_{\text{target}}}{\text{CurrentSpeed}_{\text{user}}} \right) \quad (2.4)$$

Since the damage dealt by *Gyro Ball* gets bigger, the lower the spe of the attacker, Pokémon using this move have 0 ev and 0 iv in the spe stat.

Note: Being able to outspeed the opponent is extremely valuable, but the only two Pokémon using *Gyro Ball*, *Stakataka* and *Ferrothron*, already have a very low spe stat and are

⁶In older games *Defog* would only remove Hazards on the target's side of the field. But as we only investigate the latest version, this won't be covered in detail.

⁷The term *grounded* is used to describe a Pokémon that can not be affected by damaging *Ground*-type moves and several other associated effects (Bulbapedia, 2021n).

slower than almost all other Pokémons in random battles. A complete list of Pokémons with their respective spe stat can be found at (Bulbapedia, 2021t).

This knowledge can be exploited to gather additional information about the enemy. Section 4.2 describes how this is achieved.

2.4.2. Items

Items in random battles are procedurally generated by showdown and depend on the Pokémons moves, base stats and ability. As stated in (Mitarai, 2021), the exact implementation is “changed frequently with the intention of optimizing set generation“, yet, item assignment follows these rules:

- Pokémons with 2 or fewer attacking moves will get *Leftovers*, or *Black Sludge* if *Poison*-type.
- Pokémons with 3 attacking moves will get **Life Orb**, if the sum of their base hp, def and spd is less than 275. Otherwise, these Pokémons get *Leftovers* or *Black Sludge*.
- Pokémons with 4 matching attacks get a *Choice* item which follows these rules:
 - Pokémons with four physical attacks or four special attacks, a base spe between 60 and 108 and base atk or spa of 100 or more can get a *Choice Scarf* 2/3 of the time. If the Pokémon does not meet one of the stat qualifications or does not get the 2/3 chance, they’ll get *Choice Specs* or *Choice Band* instead.
 - Pokémons with 3 special attacks and the move *U-turn* always get *Choice Specs*. *U-turn* is a physical, *Bug*-type move that switches the user out after damage is dealt (Bulbapedia, 2021ac).
 - Pokémons with *Trick* (Bulbapedia, 2021ab) or *Switcheroo* (Bulbapedia, 2021z), both moves that allow to switch items with the opponent, will always get a choice item. If they meet the above-mentioned speed range, they will always get a *Choice Scarf*. Otherwise, they will always get *Choice Specs* or *Choice Band*.
 - Having priority moves will always prevent a *Choice Scarf* from being generated in all situations.
- Pokémons with 4 attacks that do not qualify for choice items, will get an *Assault Vest* if their $hp + def + spd \geq 235$. Otherwise, *Expert Belt*, *Leftovers* or *Life Orb* is generated.
- Pokémons that are weak to Rock will get *Heavy-Duty Boots* if they don’t get a higher priority item, such as *Assault Vest* or a choice item. Pokémons that are four times weak to *Rock*, such as *Charizard*, will always get *Heavy-Duty Boots*. This is done as these Pokémons would otherwise lose up to 50% hp to the entry hazard *Stealth Rock* described in 2.3.1. The only exception is *Scyther*, which can get *Eviolite*.
- Pokémons in the lead slot will get *Focus Sash* if their $hp + def + spd < 255$, and they would otherwise get *Leftovers* or *Life Orb*.
- Pokémons that get a Speed-boosting move will be given a *Weakness Policy* if their $hp + def + spd \geq 300$, and they aren’t four times weak to *Ground*. This item boosts the atk and spa by two stages each if hit by a super effective move. After that, the item breaks (Bulbapedia, 2021ad).

There are also some species that will always roll the same item, either because it’s their signature item or because doing so supports a niche ability or set. For example, Pikachu always has *Light Ball*

2.5. Pokémon Matchups

Due to the typing system, there is no best Pokémon that is the best option in all situations. Therefore, we have to determine how good a Pokémon is against another Pokémon in a given situation. In this case, the *situation* refers to the current state of both Pokémon like current hp and status conditions as well as field effects like weather.

2.5.1. Check and Counter

There are multiple definitions of *check* and *counter* ([Cite multiple definitions](#)). In this thesis, we refer to a Pokémon *checking* another Pokémon if it can beat the enemy Pokémon in every scenario and can safely be switched in at any point. A *counter* is also capable of defeating the enemy Pokémon but may lose in some situations. The most notable being if switched in without the previous active Pokémon fainting as this grants the opponent an additional attack. ToDo

The key difference between *check* and *counter* is that a check is also stronger if it takes damage once more while a counter is not guaranteed to win in this situation.

Note: Every *check* to a Pokémon is also always a *counter* while *counter* could also be a *check*, but is not guaranteed to.

3. Related Work

In this section, previous research regarding the topic of battles on Pokémon Showdown is described and compared.

3.1. Baseline Agents

A good way to get a rough idea on how well an agent performs is to compare it against a baseline agent. There are two popular baseline agents, the *Random*-Player and the *MaxDamage*-Player. While the *Random*-Agent always chooses either a random move or a random switch, the *MaxDamage*-Agent always picks the move with the highest base power. If no move is available, the agent will switch to a random Pokémon. This is roughly equal to the skill level of an inexperienced beginner human.

3.2. Breadth-first search

Given a root battle object representing the current game state, breadth-first search (bfs) explores the outcomes of all possible choices, treating these resultant states as child nodes. This algorithm traverses the game tree until it finds a state in which the enemy Pokémon is fainted. As a non-adversarial algorithm, the agent assumes that the agent does not move at all. This agent won 75 out of 90 total games played against a random agent (Lee & Togelius, 2017).

3.3. Minimax

Minimax builds upon bfs as this algorithm deals with adversarial paradigms by assuming the opponents act in their best interest. There are multiple possibilities on how to implement the Minimax algorithm for Pokémon games. The main difference in implementations lies in state evaluation and the assumptions about the opponents. Additionally, the amount of features taken into consideration have a considerable impact. In the implementation proposed by Lee and Togelius, 2017, a node represents the worst case scenario that would occur as a result of the current choice. The agent also employs alpha-beta pruning, ignoring any node in which the agents Pokémon faints. One drawback of this procedure is that, a Pokémon would never use the move *Explosion*, a very powerful *Normal*-type move that also faints the user. This move can for example be used if the active member is already at very low hp. The tree itself is traversed using a greedy strategy, which terminates when a state in which the enemy Pokémon is fainted is reached. Both the traversal

order and the worst-case evaluation are performed using the evaluation function 3.1 (Lee & Togelius, 2017):

$$Eval = \frac{\text{current hp}_{\text{Own Pok}\acute{\text{e}}\text{mon}}}{\max \text{hp}_{\text{Own Pok}\acute{\text{e}}\text{mon}}} - 3 \cdot \frac{\text{current hp}_{\text{Enemy Pok}\acute{\text{e}}\text{mon}}}{\max \text{hp}_{\text{Enemy Pok}\acute{\text{e}}\text{mon}}} - 0.3 \cdot \text{depth} \quad (3.1)$$

A YouTube-Video released by *Rempton Games* “Programming AI for Pokemon Showdown + Bot Battle Royale!”, 2021 (**This quote is even stranger**) alters this evaluation function by increasing the rating of a game state when an enemy takes damage and decreasing the rating when a member of the own team takes damage. Both evaluation functions do not take hazards, status condition or boosts into account. Due to the similarity of both evaluation functions, both agents performed very similarly: The first agent described by Lee and Togelius, 2017 won 70 out of 90 total games against a random agent, resulting in a win-ratio of $\approx 86\%$. The second agent defeated the random player in 831 out of 1000 total games, yielding a win-ratio of $\approx 83\%$. Due to the large amount of random number generation (rng) in battles, we assume both agents to be at an equal level of play.

The current state of the art search-based algorithm was developed by *pmariglia* and is fully available at pmariglia, 2022. This approach uses the *Expectiminimax*-Algorithm and takes hazards, boosts and status into consideration. In addition to “min” and “max” nodes of a traditional *Minimax*-tree, this variant has “chance” nodes, which take the expected value of a random event occurring (Wikipedia, 2022a). Currently, the agent calculates two turns in advance, and reaches an Elo rating of 1461 as well as a Glicko-1 rating of 1633 in Generation VII random battles on the Showdown ladder. In addition, *pmariglia* implemented an algorithm to estimate the item a Pokémon is holding based on the damage it dealt.

3.4. Rule based agents

The YouTube-Video released by *Rempton Games* “Programming AI for Pokemon Showdown + Bot Battle Royale!”, 2021 also introduces two rule-based agents. *Smart Damage* was written by the author of the video and uses Pokémon type and the spe-stat to determine the favorability of a matchup. On a bad matchup, this agent will switch to the team member with the best matchup. Otherwise, a simplified damage calculation is used to determine the move dealing the most amount of damage. *Simple Heuristics* was developed by, *Haris Sahovic*, the author of the *Poke-Env* library Sahovic, 2022a. The implementation for this agent can be found at Sahovic, 2022b. This agent uses simple rules to decide the next move or switch and takes boosts as well as hazards into account. The results for both agents are denoted in table 3.1. The *Minimax*-agent in the table reffers to the implementation of *Rempton Games*.

	<i>Random</i>	<i>Max damage (dmg)</i>	<i>Smart dmg</i>	<i>Minimax</i>	<i>Heuristics</i>
<i>Random</i>	N / A	897 / 103	957 / 43	831 / 169	992 / 8
<i>Max dmg</i>		N / A	829 / 171	834 / 166	955 / 45
<i>Smart dmg</i>			N / A	331 / 669	720 / 280
<i>Minimax</i>				N / A	181 / 819

Table 3.1.: Denotes how many games the Agent in the column won against the agent in the row “Programming AI for Pokemon Showdown + Bot Battle Royale!”, 2021

3.5. Other approaches

The authors of the *Showdown AI Competition* (Lee & Togelius, 2017) compared many simple AI implementations with each other. Approaches not covered so far will be summarized in this chapter.

One Turn Lookahead

One Turn Lookahead is a heuristic agent designed to encapsulate a greedy strategy that prioritizes the damage output. The agent operates by estimating the damage dealt by all usable moves, including those usable by the agent's inactive but usable Pokémons. If the highest damaging move belongs to the active Pokémon, the agent will use that attack. If the most damaging move belongs to an inactive Pokémon, the agent will switch to that Pokémon (Lee & Togelius, 2017). Depending on implementation details, this agent is very similar to *MaxDamage* or *Rule based* agents.

Type Selector

This is a variation of the *One Turn Lookahead*-Agent that utilizes a short series of if-else statements in its decision-making. At first, if the current Pokémon knows a move that drains the opponents hp to zero, this move is selected. Otherwise, the favorability of the current matchup is evaluated. If the current type matchup is undesirable, the agent will switch to the Pokémon with an acceptable type matchup. If no such Pokémon exists, the agent will default to the most damaging move (Lee & Togelius, 2017).

Pruned Breadth-First Search

This agent is designed to demonstrate a simple way to utilize domain knowledge as a cost-cutting measure. This is achieved by making modifications to the Breadth First Search agent. First, the algorithm does not simulate any actions that involve using a damaging move with a resisted type, nor does it simulate any actions that involve switching to a Pokémon with a subpar type matchup. Additionally, rather than selfishly assuming the opponent skips their turn in each simulation, the agent assumes its opponent is a *One Turn Lookahead*-agent and simulates accordingly (Lee & Togelius, 2017).

Results

Table 3.2 displays the results of the agents described in this section.

	<i>Random</i>
<i>One Turn Lookahead</i>	77 / 90
<i>Type Selector</i>	67 / 90
<i>Pruned bfs</i>	75 / 90

Table 3.2.: Results of other approaches against a random agent (Lee & Togelius, 2017)

3.6. Self-Play Policy Optimization Approach

Researchers from *New York University* (Lee & Togelius, 2017) were the first ones to apply *Q-Learning* to the field of Pokémon battles. Two agents using *Q-Learning* were developed: A single layer perceptron as well as a multi layer perceptron. Both agents were used to output the expected reward of all current moves and switches. Based on this, the best action was picked. Both agents were rewarded for defeating opponent Pokémon and punished for allowing one of its own Pokémon to faint. Because decisions made tend to have long term consequences, weights are updated using the last three (State, Action) pairs rather than the most recent pair only. Additionally, in order to promote exploration, the agent employs an epsilon-greedy selection policy, causing it to randomly override its decision with a probability of 0.1. The single layer perceptron was trained using the Delta Rule, while the multilayer perceptron was trained using Delta Rule plus Backpropagation.

The agent using a single layer perceptron won 90 out of 180 games against a random agent whereas the multi layer perceptron won 86 out of 180 games. Due to the large amount of rng in Pokémon games, more games would need to be played in order to confirm the superiority of the single layer perceptron in this particular use case. Randomness of games will be covered in more detail in section 5.1.1

ToDo

One year after the publication of Lee and Togelius, 2017, the authors of ([Include GottaTrainEmAll in BIB!! \(Zotero\)](#)) **GottaTrainEmAll** improved this design. Schulman et al., 2017 used proximal policy optimization (ppo) to train an agent. They also used embeddings to improve the representation of a Pokémon. Data available at (alopez247, 2017) was used to create embeddings for each Pokémon. The Dataset contains stats of the first 721 Pokémon, each row contains the name, type(s), numerical stat data (such as hp, atk, spe) and some other data such as color, height and whether the Pokémon is considered legendary in game. To create embeddings for each Pokémon, the data was turned into a graph to be used with Node2Vec (Grover & Leskovec, 2016) which creates embeddings from graph data in similar to Word2Vec (Mikolov et al., 2013). This algorithm samples random walks of some number of nodes of a given graph. Using these random walks, a skip-gram model is created which can be trained to generate embeddings. The Pokémon graph consisted of the name, type(s), numerical attributes (total stats, hp, atk, def, spa, spd, spe) as well as two special nodes, *Legendary* and *Mega*¹. Lastly, Node2Vec was applied to the graph. Table 3.3 displays the discovered similarity using this approach. Here,

Pokémon	Most similar Pokémon
<i>Bulbasaur</i>	Chikorita, Turtwig, Nuzleaf, Petilil, Exegguate, Skiploom, Jumpluff, Oddish, Budew
<i>Caterpie</i>	Wurmple, Weedle, Kakuna, Metapod, Paras, Ledyba, Spinarak, Venonat, Silcoon
<i>Mewtwo</i>	Lugia, Mesprit, Mew, Victini, Celebi, Cresselia, Volcation, <i>Ho-Oh</i> , Uxie

Table 3.3.: Similar Pokémon within the embedding space of **GottaTrainEmAll**

a similar Pokémon to *Mewtwo*, a *Psychic*-Type, is *Ho-Oh*, a *Flying-Fire*-Type. However, these two have entirely different strengths and weaknesses. In addition to that, they do not share a single move in their move set: In generation 6, *Ho-Oh* has access to the following moves: *Aura Sphere*, *Calm Mind*, *Fire Blast*, *Ice Beam*, *Psystrike* and *Recover* whereas *Lugia*'s move pool consists of *Brave Bird*, *Earthquake*, *Flame Charge*, *Roost*, *Sacred Fire*, *Substitute* and *Toxic* (Honko, 2022). This inappropriate classification is likely caused by both *Pokémon* having similar stats and both being legendary. ([The paper states that they are confident in their embeddings. This critic is my own work. How do I note this properly?](#))

ToDo

Features are derived from the battle state in the simulator. At a high level, battles consist of two sides. Each side consists of a team of Pokémon, and each Pokémon has some set of moves. Each of these objects (battle, side Pokémon, and move) have attributes that are used to derive a feature vector. After experimenting with multiple network architectures, the authors settled on a three-layer fully connected neural network, each with 512 neurons and a ReLU activation function.

The authors of **GottaTrainEmAll** trained the agent against a random agent, a default

¹Mega is a mechanic similar to dynamaxing. Mega is not available in the latest version of the game and therefore won't be covered in detail

agent that always tries to pick a non-switching move as well as a *Minimax*² agent until the reward curve stabilized which was around 100 epochs where an epoch is a single battle between two players. A reward of +1 is given to the agent if it wins the battle, -1 if it loses the battle and 0 for all other cases. The average epoch reward after training convergence for this approach can be found in table 3.4. The authors of **GottaTrainEmAll** describe

Opponent	Average epoch reward
<i>Random</i>	0.85
<i>Default</i>	0.85
<i>Minmax</i>	-0.9

Table 3.4.: Average epoch rewards after training convergence for opponent agents (**GottaTrainEmAll**). Agent outperforms random and default agent but fails to defeat the Minimax-agent.

their final agent as flawed as while the agent learned to switch Pokémons when the active Pokémons reaches low health, it almost always chooses to switch to the Pokémons in the last slot. In addition, this agent preferentially chooses the fourth move.

3.6.1. State of the art

In 2019, Huang and Lee, 2019 published a paper titled *A Self-Play Policy Optimization Approach to Battling Pokémons*. Due to its fine-grained analysis and the resulting agent performing on par with the state-of-the-art search based AI, this paper will be summarized in more depth.

Similar to OpenAI, 2018, the agent is represented using an actor-critic neural network. Actor-critic RL methods (Konda & Tsitsiklis, n.d.) (**Citation broken**) combine policy-based and value-based RL methods by predicting both policy and value for a given state, and then using the value prediction, the “critic”, as an estimate of expected return when updating the policy prediction, the “actor”. The authors represent both actor and critic using a two-headed neural network which is trained via self-play RL (Huang & Lee, 2019).

ToDo

Neural Network

Input to the neural network is the current state of the game, from the point of view of the player, represented as multi-level tree-like structure:

1. The *battle* consists of two *teams*, along with weather effects.
2. Each *team* consists of six *Pokémon*, along with side conditions described in section 2.1.7
3. Each *Pokémon* has many features. Table 3.5 contains a partial list³

The network has two outputs: a probability distribution $\pi \in \mathbb{R}^n$ over actions to take, and an estimate of player strength in the current state $v \in \mathbb{R}$. Huang and Lee, 2019 compute π as follows: (**Move neural network figure to the appendix**)

ToDo

1. The network outputs an intermediate vector $p \in \mathbb{R}^n$. Each of the colored cells in figure 3.1 correspond to an element of p
2. A probability distribution $\pi' \in \mathbb{R}^n$ is computed by using the softmax function:

$$\pi' = \frac{\exp(p_i)}{\sum_i \exp(p_i)}$$

²The Authors don't provide implementation details of their *Minimax* implementation

³The authors state that this list is not complete, but no additional information is provided.

Feature	Type	Dims	Description
<i>species</i>	categorical	1×1023	e.g. Pikachu
<i>item</i>	categorical	1×368	e.g. Leftovers, Choice Band
<i>ability</i>	categorical	4×238	e.g. Rough Skin, Shadow Tag
<i>moveset</i>	categorical	4×731	e.g. Flamethrower, Surf
<i>lastmove</i>	categorical	1×731	The last move used
<i>stats</i>	continuous	6	hp, atk, def, spa, spd, spe
<i>boosts</i>	continuous	6	Temporary boosts for stats
<i>hp</i>	continuous	1	Current number of hp
<i>maxhp</i>	continuous	1	Number of hp at full health
<i>ppUsed</i>	continuous	4	# times a move was used
<i>active</i>	indicator	1	1 if Pokémon is active, else 0
<i>fainted</i>	indicator	1	1 if Pokémon has no hp, else 0
<i>status</i>	indicator	28	e.g. slp, brn, par
<i>types</i>	indicator	18	e.g. Bug, Fire
<i>volatiles</i>	indicator	23	e.g. Leech Seed, Perish Song

Table 3.5.: Features used to describe a single Pokémon battle (Huang & Lee, 2019)

3. As not every action is valid in every state, for example, a switch to a Pokémon is invalid if that Pokémon is already fainted, the authors ensure their agent has zero probability of taking invalid actions. To do this, they take a mask $s \in \{0, 1\}^n$ as part of the input, and renormalize probabilities to obtain $\pi : \pi_i = \frac{s_i \pi'_i}{s^T \pi'}$

The authors point out the following two key design decisions: First, a 128-dimensional entity embedding layer for each of the categorical variables is used. This enables capturing similarities between different moves, species and abilities without having to directly model their, often complicated, effects. Second, the parameters for computing p from above are shared among all n actions. The resulting network is described by figure 3.1 and contains 1,327,618 parameters in total (Huang & Lee, 2019).

Training the network

Training was done serially: After $m = 7680$ games per iteration, the neural network parameters are updated using the $2m$ self-play matches as training data to obtain new neural network parameters. A reward of +1 for a win and -1 for a loss is assigned at the end of the match. To speed up learning, a dense reward signal using reward shaping was constructed. Auxiliary rewards are assigned based on events that occur over the course of the match. For example, a reward of -0.0125 is added when a Pokémon of the agent faints, and a reward of +0.0025 whenever the player's Pokémon makes a super effective move⁴.

To update the neural network, the authors use *Proximal Policy Optimization* (Schulman et al., 2017), which optimizes an objective function that combines expected reward, accuracy of state value prediction, and a bonus for high entropy policies. To reduce the variance of policy gradient estimates, *Generalized Advantage Estimation* (Schulman et al., 2018) is used.

After 500 iterations of the training loop, 3,840,000 self-play matches had been played by the neural network. Training was performed using Google Cloud Platform over the course of 6 days with an approximated cost of \$91 USD.

⁴While the authors do not provide an exact definition, we assume that a move is regarded as *super effective*, if the type modifier described in Damage calculation is bigger than one

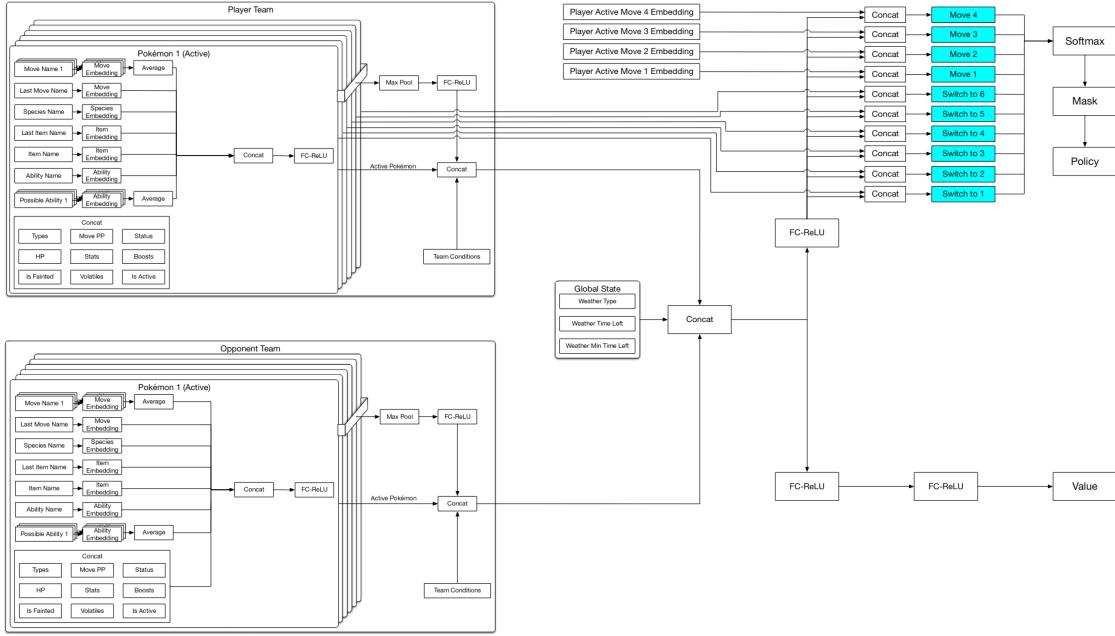


Figure 3.1.: The actor-critic neural network used by the authors in (Huang & Lee, 2019)

Evaluation

The refined embeddings as well as the complex network architecture lead to this agent outperforming the RL-Agent described in **GottaTrainEmAll**. Furthermore, the agent played 1.000 games against the open source agent of *pmariglia* mentioned in 3.3 Table 3.6 de-

Opponent	Wins	Losses
<i>random</i>	995	5
<i>max-damage</i>	929	71
<i>max-damage-typed</i>	829	171
<i>pmariglia</i>	612	388

Table 3.6.: Performance of the agent developed by Huang and Lee, 2019

scribes the performance of the agent developed by Huang and Lee, 2019. While the authors do not provide the resulting Elo of the agent, they state that a *Glicko-1* rating of 1677 was reached. It is important to note that this agent played against an older version of *pmariglia* which was released in May 2019. Results of this agent described in section 3.3 and section Evaluation used the latest release as of February 2022.

3.7. Supervised Approach

As of writing, there is only one publication researching supervised learning for Pokémon battles, a YouTube video uploaded to the channel “*The Third Build*“ “How an A.I. is Becoming the World’s Best Pokémon Player”, 2021. Unfortunately, the video does not cover much technical and implementation details, possibly to reach a broader audience on YouTube. The author obtained two million replays of human games from the creators of Pokémon Showdown. Using these replays, the input vector for a game contained the following features is created:

- **Pokémon attributes:** hp, whether the Pokémon is active, status condition and stat boosts.

- **Player’s side attributes:** side conditions (like *Light Screen*), entry hazards (like *Stealth Rocks*), volatile conditions like *Leech Seed* and the last used move
- **Battlefield attributes:** weather, pseudo-weather⁵ and terrain

Using these features, a neural network using TensorFlow for JavaScript was then trained to predict the win chance of a given player at a given state of the game. This network was able to predict the winner at a given state with an accuracy of 81%. In addition to that, the model was able to predict the next switch in with an accuracy of 86%, the next move with a certainty of 93% and whether the player would switch with an accuracy of 88%. Unfortunately, the author neither states the game type of the replays nor the exact architecture and evaluation method of the model.

This model was then used to pick the best move in a given scenario which allowed the agent to play games on the Pokémon Showdown ladder. Unlike other agents quoted in this thesis, this agent does not play *random* battles on Pokémon Showdown, but rather *Gen 8 OU Singles*.

Furthermore, *Inverse Damage Calculation* is introduced. As the exact item of enemies are unknown, the author predicts the given item of a Pokémon by comparing the damage dealt of a move with the expected damage for the move modified by possible items. Despite no concrete implementation details given, we assume this approach to work similar to the implementation of *pmariglia* described in 3.3. According to the video, this agent reached a maximum Elo of 1630 which ranks the agent among the best 3.6% of players.

⁵The video does not give an exact definition of the word *Pseudo-Weather*

4. Approach

(Short introduction. Describe idea. This builds on top of related work as combines minimax with rulebased) (Predictions: Use Move against most likely next enemy. No prediction switching (yet)) (Matchups only recalculated on new information.)

ToDo
ToDo
ToDo

4.1. Communication with Pokémon Showdown

The communication with Pokémon Showdown is handled using the python library *PokeEnv* (Sahovic, 2022a). This library provides a lot of the core functionality needed, like accessing the current Pokémon in battle as well as switch and move options. However, it does not provide functionality for damage calculation. We use the *Pokémon Damage Calculator* (Smogon, 2021), a node library written by the smogon-team for that. Communication between the two libraries is implemented by capturing stdout and stdin using the *subprocess* python library.

4.2. Gathering Information about the enemy Pokémon

As mentioned in 2.4, the same Pokémon can occur in various different builds, meaning the combination of moves, abilities and items. Knowing the exact build is crucial for the decision-making process. Consider the following example:

- **Player1** has an active *Charizard* with *Heavy-Duty Boots* and 150hp remaining on the field.
- **Player2** has just sent out a *Drapion* with 160hp remaining.
- The *Charizard* is faster but can not kill the enemy *Drapion* in one turn as his move *Fire Blast* deals between 127 and 151 damage to the *Drapion*.
- Therefore, if **Player1** decides to attack, *Drapion* is guaranteed to survive this turn and can attack *Charizard* as well.

In this scenario, the optimal strategy for **Player1** depends heavily on the move set of the enemy *Drapion*. Possible moves for *Drapion* are:

- *Aqua Tail*: A damaging *Water*-type move
- *Earthquake*: A damaging *Ground*-type move

- *Knock Off*: A damaging *Dark*-type move
- *Poison Jab*: A damaging *Poison*-type move
- *Swords Dance*: A move to raise the own atk by two stages
- *Taunt*: A move that makes the afflicted Pokémon unable to use status moves
- *Toxic Spikes*: A move that sets an entry hazard.

Hereby it is important to note that *Drapion* only knows the move *Aqua Tail*, if it knows four total damaging moves¹. In the given scenario, **Player1** should switch out his *Charizard* if the enemy *Drapion* knows the move *Aqua Tail* as this attack would kill *Charizard*. We can determine whether the enemy knows *Aqua Tail* based on his item:

- If *Drapion* rolls two status moves, it will have the item *Black Sludge* and therefore doesn't know *Aqua Tail*. Because *Drapion* is already damaged, we know that it has this item if it healed 1/16% of his max hp in his last turn.
- If *Drapion* rolls one status move, it will have the item *Life Orb*. If *Drapion* already attacked, we know if it has a *Life Orb* or not as this item causes it to lose 10% of his maximum hp after an attack.
- If *Drapion* has neither *Black Sludge* nor *Life Orb*, it has to have a *Choice Band* and as this item will only generate if the Pokémon knows four matching attacks, and therefore has to know the move *Aqua Tail*.

Therefore, the player should switch out his *Charizard* if the *Drapion* has neither *Black Sludge* nor *Life Orb* as then it is guaranteed to have a *Choice Band*, knows *Aqua Tail* and would therefore be able to defeat *Charizard*. In any other case, the **Player1** can safely kill *Drapion* using his active *Charizard*.

Implementation details

The first step to predicting enemy sets is to determine all possible sets and the likeliness of each individual set. In order to achieve this, I wrote a script that starts a battle between an information gathering player and a random agent. In the next step, the script extracts all builds of all Pokémon and stores them. Then, it forfeits, and a new battle is started. Once enough battles are fought, the script will store the builds as well as how often they appeared in text files, one file for each Pokémon.

In actual battles, if a new Pokémon enters the enemy side, we assume it to have the most likely build for its species. Once more information becomes available on items, moves and abilities, we rule-out non-matching builds and always assume the enemy to have the remaining most likely build.

4.3. Scoring the current game state

In order to not only rate the current board state, but also individual Pokémon, we implement the following scoring algorithm:

$$score(e_{i,j}) = \text{Expected Damage that Pokémon } i \text{ will deal to Pokémon } j \quad (4.1)$$

The expected damage is the damage dealt if both Pokémon behave optimal in the amount of turns that the bot looks into the future. Section 4.5 covers how the optimal moves are determined.

$$value(i) = \sum_{j \in \text{Enemy Pokémon}} score(e_{i,j}) \quad (4.2)$$

¹This information was extracted from all possible *Drapion* builds. Details on how these builds were obtained are covered in 4.2

Using this, we can also introduce a *value* for each of our Pokémons where a higher value implies a more important Pokémon. It is important to note that scores are determined independently of each other meaning that we do not take into account damage taken by the attacker. This does imply that this metric does not determine how good the Pokémon is if it has to battle *all* enemy Pokémons but rather against how many other Pokémons it *could* be used. This is done as the order in which a Pokémon battles multiple Pokémons plays a critical role. The reasoning behind this as well as the determination of an optimal order is explained in section 4.4. This metric also has multiple flaws because it only takes the damage dealt to the enemy into account, other important factors like damage received, healing, the availability of status moves and hazards is not considered. Using this rating, we can rank our team based on importance which is used in the switching routine described in section 4.4.

4.4. Stages of the game

We divide the game into two phases, the first one being the *Discover*-Phase whereas the second phase is called *Defeat*-Phase. Our goal and therefore our play style, is different in both phases.

Discover Phase

At the beginning of the game, we play safely until we know our opponent's entire team. In other words, we try to gather information about the enemy team while sacrificing as little hp as possible. In this stage, we act according to these rules:

1. Kill the opponent if we are guaranteed to kill him this turn. This either leads to us defeating the enemy Pokémons, or possibly new information if the enemy switches.
2. Healing our Pokémons. If we have a healing move that will heal us more than the expected damage we receive this turn, and we are not at full hp we will heal our Pokémons. Doing so will force the enemy to switch as we are otherwise gaining an advantage over him. (**Go into a bit more depth here?**) ToDo
3. If we have a hazard setting move available, we will use this move as they will help us in the *Defeat*-Phase. Other beneficial side-conditions like *Light Screen* will be set as well.
4. Using moves that inflict status to the opponent like brn.

If none of these conditions apply, we decide whether to switch out our Pokémons or not. If our current Pokémon is a *check* or *counter* against the enemy Pokémons, we will not switch. Otherwise, we switch to a *check*, or *counter* if present. Next, we check if the current matchup is unfavorable. This applies, if the enemy is expected to survive the current matchups for two turns longer than we do, meaning that our Pokémons would not be able to defeat the enemy, if it was allowed to attack two additional times. If this is the case, we determine the next action as follows:

We start by determining the *score* of each of our Pokémons as described in formula 4.1 and exclude our two most valuable Pokémons from the next steps. This is done as we assume to depend heavily on these Pokémons to defeat other enemies. Next up, we pick the Pokémons with the lowest *score* that fulfills the following criteria:

- The Pokémons has to survive for at least three turns against the enemy
- At least one of the following criteria apply:
 - It is able to set a *hazard* or other beneficial field effect that is not already present.

- It can inflict a status condition on the enemy.

If no Pokémon matches these criteria, the Pokémon with the lowest score is picked instead. Score is calculated as described in section Scoring the current game state.

In any other case, the currently active Pokémon will use the best move calculated in 4.5.

Defeat Phase

ToDo

(Move tree to appendix. If the tree is too small printed (when on appendix): Smaller Example using TikZ.) The *Defeat*-Phase starts as soon as we know all enemy Pokémon. At the beginning of this stage, we have to create a match plan for defeating the enemy. The main goal is to figure out which Pokémon to use against which enemies, especially the best order to send them into battle. Figure 4.1 shows a possible end-game scenario.

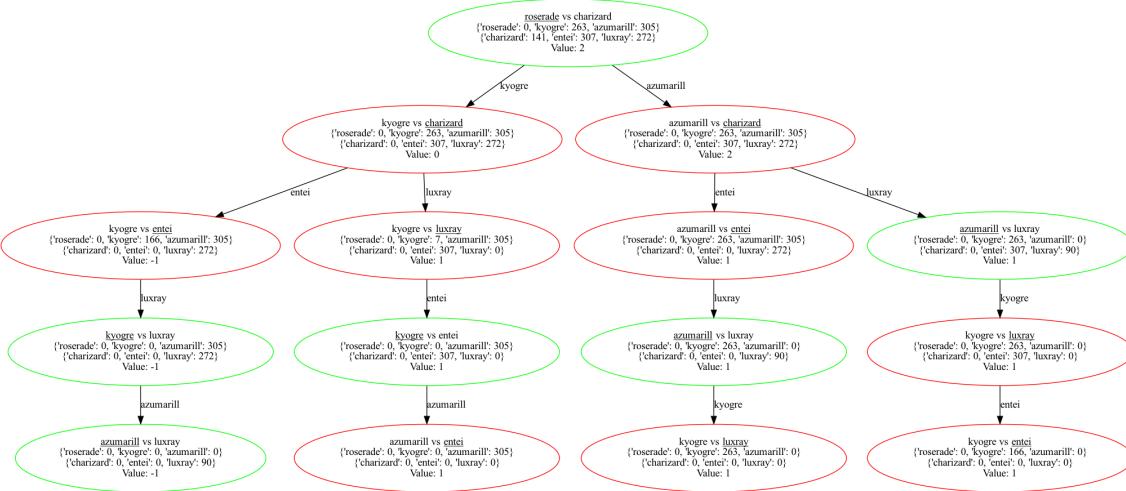


Figure 4.1.: Possible end game scenario

In this simple example, we assume both players to each have three Pokémon with full hp and no status conditions. Each circle represents two Pokémon battling each other. Therefore, the circle in the top center indicates, that our *Roserade* is currently fighting against the enemy *Charizard*. As *Charizard* is a *Fire*-type Pokémon which has a type advantage against his enemy, our *Roserade* will faint in this matchup, expressed by the underlined name. This means that Player 1 has to switch in a new Pokémon. A green circle around a node illustrates that the first player has to make a turn as his Pokémon fainted, and a red circle marks a decision for Player 2. The first line below the names of the two opposing Pokémon displays the amount of hp the first team has left after this battle took place. As *Roserade* fainted, it has zero hp left whereas the other two Pokémon in his team are still at full health. Below that, the remaining hp of the second team is displayed. *Charizard* is *expected* to survive the battle with an *expected* amount of 141 hp remaining. Next, we have two possible options remaining, we can either send *Kyogre* or *Azumarill* to defeat the enemy *Charizard*. Taking a look at the leaves of the tree, the *value* of a leaf node is -1 if the enemy wins the battle and 1 if we win the battle. The value of a non-leaf node is the sum of the *values* of the children nodes. The *value* of two means that we expect to win the battle. As a result of this, the best choice in our example scenario is to use *Azumarill* to defeat the already damaged *Charizard*.

ToDo

(Explain why MiniMax Tree is correct. Pokemon don't have expected items / moves!)

4.5. Determining matchups

In order to determine whether to attack or to switch, we need to determine how the optimal moves for a Pokémon against another Pokémon are calculated. As stated before, the

amount of possible combinations combined with the non-deterministic nature of the game makes it unfeasible get the optimal move combination by simulating every possible combination of actions and reactions over the span of multiple turns. Therefore, we determine the optimal moves for a Pokémon using this simplified method:

We start by generating all possible move combinations for a Pokémon with a given length. Then, we simulate the outcome of the battle. If the attacking Pokémon would use these moves in the defined order given the enemy would not do anything². Here, we also take boosting, items, status effects and the possibly changing field state into account. The combination resulting in the lowest amount of turns until the enemy faints is selected. It is important to note that this is not necessarily the move that the bot will use in the next turn as in the *Discover*-Phase healing, status and other beneficial effects are prioritized.

4.6. Predictions

(Explain predicting routine) Predicting the next enemy action plays a huge role in competitive Pokémon. In many situations, it is very likely that the opponent will switch to another Pokémon to avoid a bad matchup. For example, players often withdraw *Fire*-type Pokémon if the opponent knows a *Water*-type move as this will deal twice the amount of damage due to the type multiplier discussed in section 2.1.1. Therefore, players often choose to pick a move countering the expected Pokémon the opponent will send out or even switch themselves in anticipation of an enemy action. Our agent mimics this behavior in the following way:

ToDo

The agent keeps track how often the opponent chose to withdraw his active Pokémon when it is checked and how often he decided to stay in. In order to make the actions taken by the agent less predictable, rng is used on favorable matchups to decide how to act. The chance that the agent predicts an enemy switch is zero if this is the first time that our Pokémon counters an enemy and otherwise as described in equation 4.3.

$$\text{Chance make a prediction} = \frac{\# \text{ The enemy switched when countered}}{\# \text{ Total times the enemy was countered}} \quad (4.3)$$

Currently, the agent assumes the enemy to switch in a random Pokémon that counters the active Pokémon and will then pick the move, that will yield to the best outcome for this matchup. Flaws and possible improvements to this simple routine will be discussed in section 5.2.2.

4.7. Dynamaxing

At the time of writing, *Memezboi69* is the top ranked player on the Pokémon Showdown ladder. He described the use of dynamaxing as very situation dependent as it can either be used offensive or defensive and proposed a simple rule set that can be build upon: Dynamax if the current matchup is favorable, the entire enemy team is known and the active Pokémon has more than 70% hp remaining. In addition, the agent will dynamax if the current Pokémon is the last remaining Pokémon with full hp. *Memezboi69* also laid out a more sophisticated rule set, which is described in section 5.2.2, that can be applied in future iterations of the agent.

²The option of not doing anything in a turn does not exist in Pokémon, if possible, the player is always forced to either select a move or switch

5. Evaluation

(Lee-Paper: MinMax requires simulation)

ToDo

5.1. Challenges for evaluation

Different researchers use different metrics to evaluate the performance of their agents. There are multiple factors that increase the difficulty of properly evaluating the performance.

5.1.1. Randomness of battles

As teams are generated randomly, one player often ends up with a slightly better team than his opponent. In extreme cases, one player may not even have a chance at winning the battle. While battling our agent during the evaluation process, one particular game stood out as the first Pokémon of **Player1** was capable of defeating the entire enemy team. The Pokémon of **Player1** was a *Volcarona* with the following moves:

- *Fire Blast*, a damaging *Fire*-Type move
- *Quiver Dance*, a *Bug*-Type move that boosts the users spa, spd and spe by one stage each.
- *Bug Buzz*, a damaging *Bug*-Type move
- *Roost*, a move that restores half of the user's maximum hp

This Pokémon was able to defeat the entire enemy team with little to no possible counter play: The first enemy Pokémon, *Leafeon*, a *Grass*-Type Pokémon was killed in one hit using *Fire Blast* after damaging *Volcarona* using *X-Scissor*.

Next, *Glalie*, an *Ice*-Type was sent into battle. *Glalie* uses his best move, *Earth Quake* which brings *Volcarona* to 52% hp. As the enemy does not pressure *Glalie* much, **Player1** decided to boost using *Quiver Dance*. Now, *Volcarona* is faster than his enemy and kills it again in one hit using *Fire Blast*.

Then, *Mr. Mime (Galar)* is sent into battle. He fails to pressure *Volcarona* and therefore, **Player1** can heal his Pokémon using *Roost* and further boost using *Quiver Dance*. After defeating *Mr. Mime (Galar)*, *Volcarona* is back to 84% HP and boosts of 2.5 spa, 1.5 spd and 2.5 spe.

boosted this high, *Volcarona* can one-shot both the enemy *Volcarona*, *Pheromosa* and the

dynamaxed *Scraggy* using *Fire Blast*.

To eliminate the impact of these extreme cases, evaluation of agents against other agents should be done using multiple hundreds to thousands of games against each other.

In order to quantify and better visualize the randomness of battles, we started by generating (**HOW MANY GAMES IN TOTAL FOR GRAPH**) random battles and collected the team information. After both teams were stored, two *MaxDamage*-agents finished the game to determine a winner. (**Board rating was not done EXACTLY the same way as in the scoring chapter. This uses isCheck! Scoring does not use isCheck as in the bot itself we only need the score if we have NO check / counter!**) Next, we determined the matchups for each battle as described in 4.5. Finally, the board rating was calculated as follows:

```

1 # Amount of counters player 1 has against player 2
2 counter_p1_count = sum([sum([
3     1 if m.is_counter(m.pokemon_1.species, m.pokemon_2.species)
4     else 0 for m in matchups if m.pokemon_1.species == p.species])
5     for p in self.team_1])
6
7 # Amount of counters player 2 has against player 1
8 counter_p2_count = sum([sum([
9     1 if m.is_counter(m.pokemon_2.species, m.pokemon_1.species)
10    else 0 for m in matchups if m.pokemon_2.species == p.species])
11    for p in self.team_2])
12
13 board_rating = counter_p1_count - counter_p2_count

```

Listing 5.1: Calculate Board rating

Finally, figure 5.1 displays these results. The figure contains how often what board-rating occurs (red) as well as the win-rate for the given rating (green). Not only does this figure show a gaussian distribution of how fair games are (**Schlechte Formulierung!**), but also that our method of determining matchups works as intended. This is due to the fact that games in which we assume the player to be in a bad position also have a win-rate and games with a favorable board rating show the highest win ratio.

5.1.2. Evaluation against human opponents

As described in (**Link Showdown chapter**), Pokémon Showdown allows researchers to use bots on the ranked ladder.

5.2. Agents

During this thesis, two different agents were developed, *HerrDonner* and *HerrGewitter*.

5.2.1. HerrDonner

This agent was designed to establish a good baseline and to demonstrate the capabilities of a very simple rule set. The agent is capable of looking multiple turns into the future. To determine what moves to be used, the Agent generates every possible move combination with the specified amount of turns into the future and calculates the expected outcome while assuming that the enemy does not move at all, similar to the bfs-based algorithm described in paragraph 3.2. No drawback moves that heal the agent, set hazards or field conditions or inflict status conditions are not considered unless they result in the highest amount of damage dealt. Also, stat changes are not taken into account, neither for damage calculation nor for rating of matchups. This results in the bot often spamming moves like *Draco Meteor*, a special *Dragon*-type move that deal a lot of damage but also lower the

ToDo

ToDo

ToDo

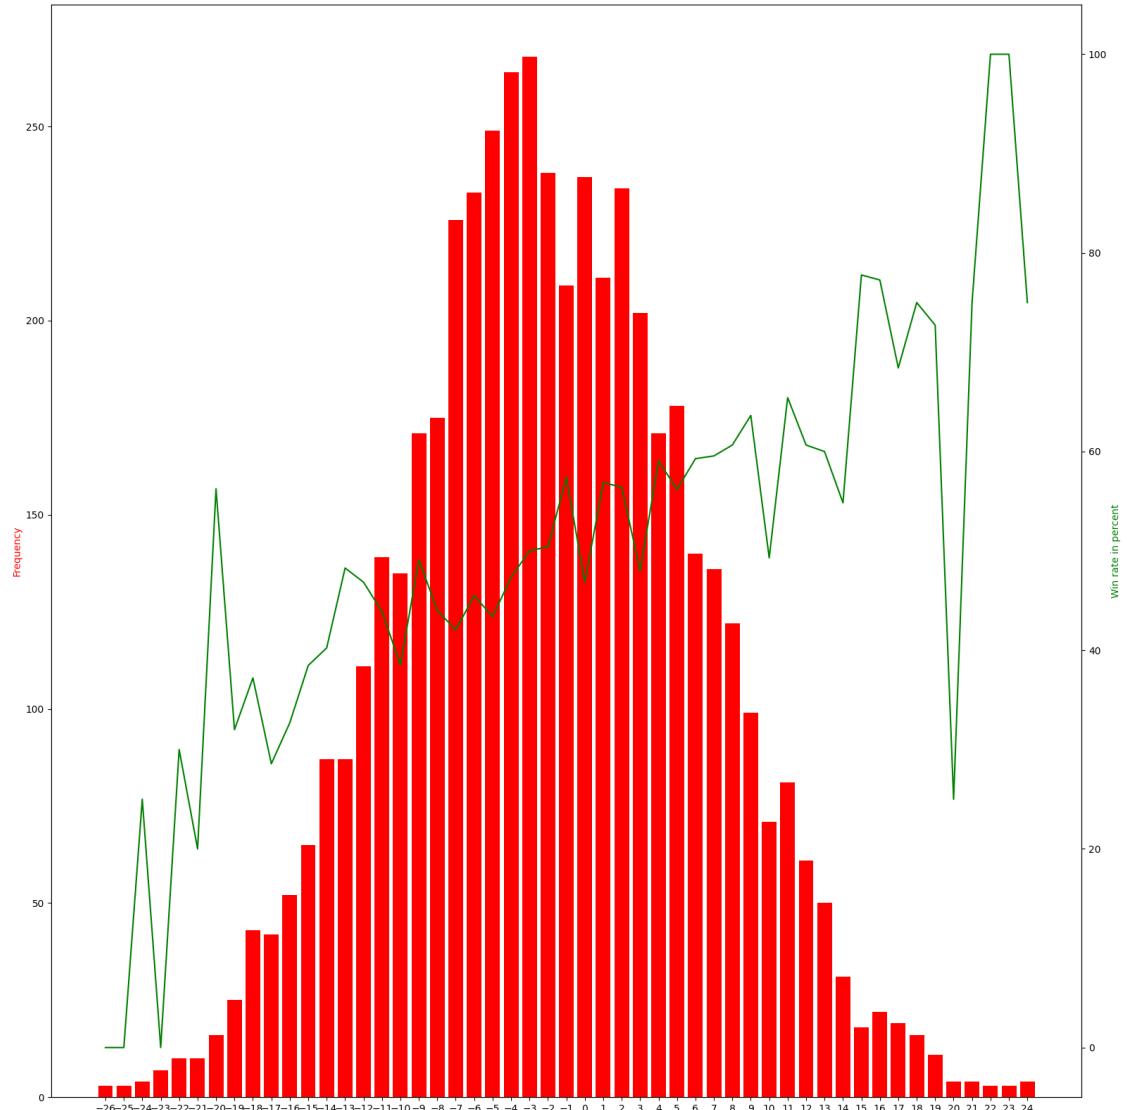


Figure 5.1.: Win rate for a given board rating

users spa by two stages resulting in the move dealing less damage every time it is carried out. When the agent is forced to switch, it will switch to a check if available. If no check is available, a counter is sent into battle if one exists. Otherwise, a random Pokémon will be picked.

At the start of each turn, *HerrDonner* will check if the current matchup is not favorable, a matchup is deemed unfavorable, if the current Pokémon is neither check nor counter to the current enemy. On a bad matchup, the bot will switch to an available check or counter. If neither is available, the bot won't switch and try to defeat the current opponent with his active team member.

Dynamaxing is implemented in a very simple and naive way: The agent will always dynamax the active Pokémon as soon as more than four enemy Pokémons are known. Lastly, if the current Pokémon is dynamaxed, the agent will not switch, even if the current matchup is not favorable. The challenges of properly using this mechanic will be explained in 5.2.2.

5.2.2. HerrGewitter

HerrGewitter behaves like described in section 4. Here, the most notable differences between both agents are highlighted, and limitations of this agent are discussed.

Firstly, *HerrGewitter* takes more factors into consideration. Now, damage calculation is done in regards to current stat changes and status conditions. Additionally, abilities and items are taken into account. Furthermore, recoil from moves, healing both from items like *Leftovers* (Bulbapedia, 2021q) and moves like *Recover* are not neglected anymore.

Switching and the selection of moves is done as described in section 4.

These improvements lead to *HerrGewitter* avoiding mistakes of *HerrDonner*. For example, this agent will burn a physical attacker using *Will-O-Wisp* (Bulbapedia:Will-O-Wisp) in order to reduce damage taken over the next turns. The agent will also boost and heal itself in favorable situations which stalls the game and forces the opponent to react. Another major improvement is that the agent switches out the current Pokémon if stat changes resulted in an unfavorable matchup which is especially important as stat changes reset on a swap. However, there are still a lot of features that *HerrGewitter* is lacking.

Weather and Field effects

One such feature is a proper support for weather and field effects in the damage calculation as well as in the *Minimax*-Algorithm. For instance, a consequence of this is the agent lacking awareness of the fact that a *Fire*-Type move deals 1.5 more damage during *Harsh Sunlight*.

Hazards

Currently, the agent will always try to set a non-present Hazard in the early game as this usually results in a long term benefit. There are however some notable exceptions to this that are not yet implemented:

- The agent will always set as many hazards as possible in the early game, even if the current matchup is unfavorable, including always setting up to two layers of spikes. A small test on human players indicated that this leads to slightly better results than only setting hazards on good matchups, but due to the very small sample size, future work is needed to determine the best strategy for setting hazards.
- The agent does not take the damage taken by hazards into account when switching Pokémon.
- The agent will always use *Toxic Spikes* even if the opponent has a *Poison*-Type Pokémon on his team that will remove this hazard upon being switched in.
- The agent will use Hazards even if the current enemy is known to have a hazard-clearing move like *Defog* (Bulbapedia, 2021j)
- The agent will not clear hazards

Choice Items

As mentioned in section 2.1.6, Pokémon holding a *Choice*-item are locked into using always the same move until they are switched out. The agent has two major flaws in regard to these items: When the active Pokémon of the agent is holding a choice item and is already locked into a move, the agent is not aware of the fact that once the Pokémon is switched out, it will regain access to his other moves which leads to an incorrect prediction for future matchups. As described in ([Link chapter about re-determining matchups](#)), the only matchups re-evaluated on a given turn are matchups that include one of the currently

active Pokémon. The following example illustrates how this design decision leads to issues on Pokémon with *Choice*-items:

In the given scenario, our agent has an active *Garchomp* which is locked into using *Earth Quake*. The *Garchomp* also has access to the *Rock*-Type move *Stone Edge*. This turn *Butterfree*, a *Bug* / *Flying*-Type Pokémon is sent into Battle. As the *Ground*-Type move *Earth Quake* has no effect on *Butterfree*, the agent will switch out *Garchomp* for another Pokémon. In the current implementation, matchups for *Garchomp* are not re-evaluated. While this won't lead to problems in the early game, this results in an incorrect *MinMax* calculation as for matchups involving *Garchomp* and any non-active opponent, *Garchomp* is still assumed to only have access to the move *Earth Quake*. In this scenario, the agent would fail to realize that *Garchomp* also has access to *Stone Edge* and would incorrectly assume *Garchomp* to lose all matchups against *Flying*-Type Pokémon.

While this behavior rarely effects battle, the agent failing to notice that an enemy is choice-locked has more often a negative impact on the battle: If the enemy is known to be choice-locked into *Earth Quake* we can safely switch a *Flying*-Type into battle. This applies especially if the enemy Pokémon is known to have the *Rock*-type super effective move *Stone Edge* as the enemy can not use this move until switched out and back in again. In this scenario, the agent wrongfully would not prefer to switch in a *Flying*-Type Pokémon due to the threat posed by *Stone Edge*. Switching in a Pokémon resisting *Earth Quake* in this scenario forces the enemy to switch to another Pokémon. This gained turn advantage can either be used to land an extra move on the next opponent, set hazards, beneficial field conditions, inflict status conditions or boost the current Pokémon.

Damage Calculator

The current implementation relies on the Pokémon Showdown Damage Calculator. As of now, this open source project does only support direct damage dealt by attacking and lacks functionalities like recoil, healing from items and moves. While we added these features to our implementation, some moves are still not properly implemented. For example, the move *Counter* has a move priority of -5 and works as follows: If the last mount of damage dealt before the use of *Counter* is greater than zero and was dealt by a *Physical*-Type move, *Counter* will do twice as much damage to the opponent. Otherwise, the move will miss. Additionally, *Counter* has a lot of extra rules regarding other special moves in place (Bulbapedia, 2021f). Issues like these are especially obvious on *Wobbuffet* as all of his four most likely moves, *Mirror Coat*, *Encore*, *Counter* and *Encore* are very useful yet do not deal any damage and are not implemented yet which leads the agent to believe that this Pokémon is bad in every possible matchup and has no good use scenarios whatsoever.

Dynamaxing

As previously stated, dynamaxing is a very complex mechanic that requires sophisticated long term planning for proper usage. According to *Memezboi69*, the safest usage of this mechanic is to dynamax in a late stage of the game on a good matchup. He stated multiple reasons for this:

Dynamaxing at a later stage of the game reduces the likeliness of an opponent being able to properly react due to the limited amount of Pokémon alive.

Another flaw of *HerrDonner* is the fact that defensive dynamaxing is not yet taken into consideration. If the opponent dynamaxes offensively, a good way to minimize damage when no counter is available is to dynamax defensive and use the move *Max Guard* which protects the user from moves. If a Pokémon has access to a status move, this move will be replaced by *Max Guard* when dynamaxing. Lastly, dynamaxing can under some circumstances also be used in an earlier stage of the game if the current matchup is favorable.

(**Improve this paragraph!**) (**Dynamaxing will likely be removed next generation?**)

ToDo
ToDo

Predictions

When play testing the bot, predictions were noticeable but only seemed to occur rarely. This is likely caused by the optimal move against the current Pokémon being equal to the best attack targeting the predicted switch. In order to compete with higher ranked players, more investigation on the prediction routine is required. Future work could also include switching in anticipation of the opponent withdrawing his active Pokémon.

Other special cases

This list contains more currently unhandled cases which will be addressed in future versions:

- The Pokémon *Ditto* can transform itself into the Pokémon of the current opponent.
- The Pokémon *Zoroark* can transform itself into another team member.
- The ability *Trace* changes the ability of a Pokémon to the ability of his opponent.

MiniMax

The *MiniMax* algorithm described in paragraph 4.4 only supports changes in health but ignores other important factors such as boosts and status conditions. Therefore, the agent will not recognize the possibility to weaken a very strong physical attacker like *Garchomp* by burning it first and then defeating it with another Pokémon. A simple way to include brn into this algorithm is to multiply the expected damage dealt by a burned Pokémon with 0.5.

5.3. Results

5.3.1. Ranked results

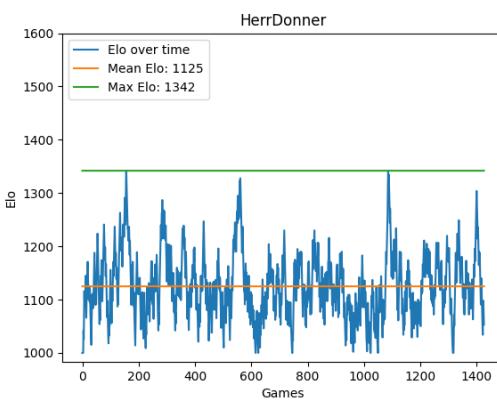


Figure 5.2.: Elo HerrDonner

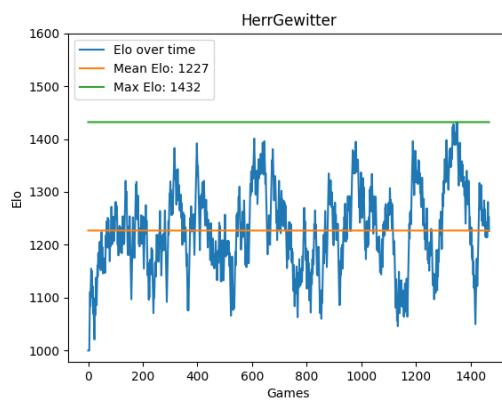


Figure 5.3.: Elo HerrGewitter

ToDo

(Print this page and check if the figures are to small) Both agents were evaluated at the same time over the span of multiple days by playing over 1400 ranked games each against human opponents on Pokémon Showdown. Figure 5.2 shows the Elo rating of *HerrDonner* over time, while figure 5.3 displays the Elo rating of *HerrGewitter*. In order to make a comparison of both agents more easy, figure 5.4 shows the smoothed Elo of both agents over time. Smoothing was achieved by dividing the Elo history in chunks of size 20 and then plotting the average Elo of each chunk.

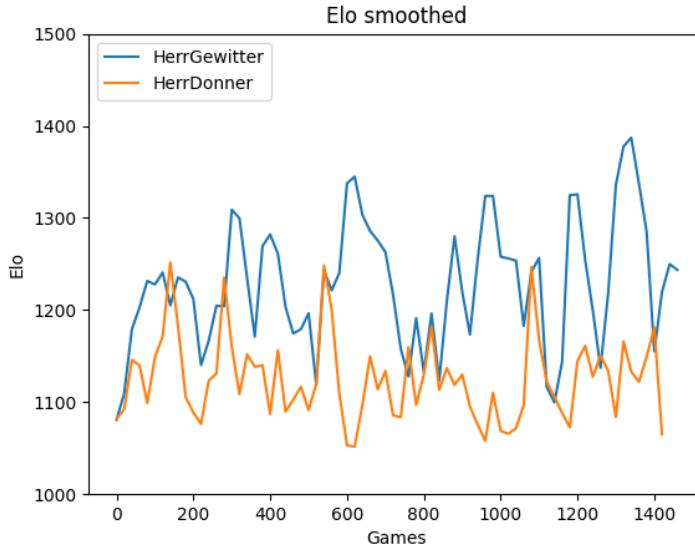


Figure 5.4.: Smoothed Elo

```

1 step = 20
2 smoothed_elo = []
3 for i in range(0, len(elo_history), step):
4     sec = elo_history[i: i + step]
5     smoothed_elo += [(i, sum(sec) / len(sec))]

```

Listing 5.2: Smoothing Elo values

The first thing to note is that *HerrGewitter* has a higher mean Elo (1227) and a higher max Elo (1432) than *HerrDonner* who achieved an average Elo of 1125 and peaked at an Elo of 1342. (**No footnote in table? How to properly format this?**) Table 5.1 shows the

	Random	MaxDmg	Pmariglia	DQN	Donner	Gewitter
Random	N / A	-	-	995 / 5	992 / 8	993 / 7
MaxDmg		N / A	-	929 / 71	906 / 94	951 / 49
Pmariglia			N / A	612 / 388 (*)	-	273 / 727
DQN				N / A	-	-
Donner					N / A	580 / 420
Gewitter						N / A

Table 5.1.: Battle results of 1,000 games between the agents.

Pmariglia: (pmariglia, 2022)

DQN: As described in (Huang & Lee, 2019)

performance of different agents when directly competing against each other.

ToDo

It is very important to note that while the agent of (Huang & Lee, 2019) played against an older version of against (2022) than our bot did and is therefore not included in this table as this possibly leads to misleading interpretations. (**Currently included. Figure this out!**)

ToDo

Each entry displays the results of a thousand games played between both agents. The first number indicates the amount of games the agent in the current column won. There are multiple things to point out here:

While *HerrDonner* and *HerrGewitter* achieved almost similar results when battling a

random agent, *HerrDonner* performed notably worse against the *MaxDmg* agent which always picks the move with the highest base power. Also, *HerrGewitter* only managed to win 58% against *HerrDonner* despite notably better results against human players, described in section 5.4. At the time of writing, neither *poke-env* nor replays created by Showdown allow the retrieval of the Glicko-1 rating. After playing for 1,400 games, *HerrDonner* ended up with a Glicko-1 rating of 1438 and *HerrGewitter* achieved a ranking of 1509.

5.4. Investigating Replays

In addition to playing against humans on the Showdown ranked ladder, the bot was also evaluated by playing against *Jonathan Krebs* and *Markus Krug*. While *Markus* was able to

	<i>HerrDonner</i>	<i>HerrGewitter</i>
<i>Jonathan</i>	28 / 23	37 / 23
<i>Markus</i>	N / A	4 / 11

Table 5.2.: Performance of both agents against human players. *HerrGewitter* posed a harder challenge but is unable to gain a lead against *Markus*, an experienced human player.

reliably beat *HerrGewitter*, the agent provided a significant challenge and required careful planning to overcome. Replays retrieved from these games confirm assumptions made about the bot in the previous sections. *HerrGewitter* is very good at picking the most damaging move, but can be exploited using mechanics not yet implemented.

5.4.1. Pro Replays

ToDo

(If *Memezboi69* is introduced in Dynamax, no introduction needed here!) *Memzboi69*, the top ranked player, agreed to play three games against *HerrGewitter*. While our agent lost all three games, we were able to gain some very interesting insights from these games:

Game One

In the first game, the bot made a huge miss play that ultimately lead to defeat. The agents *Shuckle* faced a *Mr. Mime (Galar)* which had access to *Rapid Spin*, a move that removes hazards. As *Shuckle* has access to the moves *Sticky Web* and *Stealth Rock*, the agent tried to set both hazards which were immediately removed. In the extra turn, *Mr. Mime (Galar)* boosted its speed stat and slowly killed his opponent using *Freeze-Dry*. This increased spe led to an advantage that the bot was not able to make up for. A future version of the agent will need to include a better hazard routine to prevent this kind of scenarios.

Game Two

Game two was lost due to a lack of counters against the opponents *Gengar*. Prior to sending in this Pokémon, *Memzboi69* set the entry hazard *Sticky Web* which decreases the spe by one stage on switch in. After boosting spa by two stages using *Nasty Plot*, *Gengar* was able to defeat the four remaining Pokémons of the agent. Further investigation revealed another flaw in the logic of the agent. When switching in the *Discover Phase* with no check or counter available, the agent ranks all possible options based on how much damage they are capable of dealing as described in section 4.3. If no Pokémon exists that is capable of surviving for more than two turns, the Pokémon with the lowest score will be sent into battle. In the given scenario, the agent sent out *Klefki* as his last Pokémon. This

Pokémon always has access to either *Thunderwave*, a move that applies par or *Toxic* that poisons the opponent. While the chances of winning would still be very low when sending out this Pokémon, bringing in *Klefki* would have increased the odds. Additionally, the agent did not take the entry hazard into account which lead to the wrong assumption that *Zamazenta* is faster than *Gengar* and could therefore kill the opponent. This resulted in *Zamazenta* getting killed in one shot.

Game Three

The last battle was the closest and ended with *Memezboi69* having only two Pokémons remaining. The agent even managed to set up a burned *Obstagoon*, boosting atk to stage two and def to stage three. Unfortunately, this Pokémon was then killed by *Buzzwole* with the super effective move *Close Combat*. Even after a lot of investigation, it is hard to determine how this match could have been won by the agent, one possibility might have been switching in a *check* to *Buzzwole* and dynamaxing next turn.

6. Conclusion

While simple rules like picking the most damaging move are sufficient to defeat *Random-Player* most of the time, more complex routines are required in order to rival human players.

Previously existing search-based implementations managed to yield good results against baseline agents but are outperformed by the current state of the art DQN-Agent described by Huang and Lee, 2019. Additionally, a way to implement long term planning by modifying the MiniMax-Algorithm was proposed. Chapter 5 revealed the difficulty in proper evaluation of agents created by the influence of rng.

HerrGewitter failing to rival *pmariglia* is likely caused by its inferior item prediction algorithm and handling of edge cases. Nevertheless, the agent is capable of defeating less experienced humans. Future work will address the flaws described in section 5.2.2. Additionally, a refined dynamaxing-routine and improved predictions of enemy actions are worth investigating.

Pointed out by Huang and Lee, 2019, the DQN-Agent might be improved by introducing “different network architectures or the addition of a recurrent element like an LSTM to better model human memory during the course of a game“. Experiments combining supervised learning and reinforcement learning similar to the *StarCraft II* AI *AlphaStar* (<https://deepmind.com/blog/article/alphastar-mastering-real-time-strategy-game-starcraft-Todo>) developed by *DeepMind* may improve the agent even further.

List of Figures

1.1.	Pikachu, one of the most popular Pokémon (Fandom, 2016)	3
1.2.	Exploring the map in Pokémon Red	4
1.3.	Fighting another trainer in Pokémon Red	4
1.4.	Exploring the map in Pokémon Sword. Image source: pokemon.com	4
1.5.	Fighting another trainer in Pokémon Sword Image source: nintendo.de	4
1.6.	Battle between two players on <i>Pokémon Showdown</i>	5
2.1.	Advantages of different types (Database, 2022)	8
2.2.	Possible stats of <i>Charizard</i> (Bulbapedia, 2022)	9
3.1.	The actor-critic neural network used by the authors in (Huang & Lee, 2019)	27
4.1.	Possible end game scenario	32
5.1.	Win rate for a given board rating	37
5.2.	Elo HerrDonner	40
5.3.	Elo HerrGewitter	40
5.4.	Smoothed Elo	41
A.1.	A figure	59

List of Tables

2.1.	Stage multipliers for atk, def, spa, spd and spe	15
2.2.	Damage dealt to Pokémon by Stealth Rocks (Bulbapedia, 2021x)	16
2.3.	Damage dealt to Pokémon by Sharp Steel (Bulbapedia, 2021m)	16
3.1.	Denotes how many games the Agent in the column won against the agent in the row “Programming AI for Pokemon Showdown + Bot Battle Royale!”, 2021	22
3.2.	Results of other approaches against a random agent (Lee & Togelius, 2017)	23
3.3.	Similar Pokémon within the embedding space of GottaTrainEmAll	24
3.4.	Average epoch rewards after training convergence for opponent agents (GottaTrainEmAll). Agent outperforms random and default agent but fails to defeat the Minimax-agent.	25
3.5.	Features used to describe a single Pokémon battle (Huang & Lee, 2019)	26
3.6.	Performance of the agent developed by Huang and Lee, 2019	27
5.1.	Battle results of 1,000 games between the agents. <i>Pmariglia</i> : (pmariglia, 2022) <i>DQN</i> : As described in (Huang & Lee, 2019)	41
5.2.	Performance of both agents against human players. <i>HerrGewitter</i> posed a harder challenge but is unable to gain a lead against <i>Markus</i> , an experienced human player.	42

Listings

5.1. Calculate Board rating	36
5.2. Smoothing Elo values	41

Acronyms

TBD To Be Done

hp hit points

atk attack stat

def defense stat

spa special Attack stat

spd special Defense stat

spe speed stat

crit critical hit

iv individual values

ev effort values

brn burn

frz freeze

par paralysis

psn poison

slp sleep

rng random number generation

bfs breadth-first search

ppo proximal policy optimization

dmg damage

rps rock-paper-scissors

gen1 generation 1

Bibliography

- alopez247. (2017). Pokémon for data mining and machine learning [[Online; accessed 24-January-2022]]. <https://www.kaggle.com/alopez247/pokemon-for-data-mining-and-machine-learning>
- Bulbapedia. (2016). Category:moves that use stats from different categories. https://bulbapedia.bulbagarden.net/wiki/Category:Moves_that_use_stats_from_different_categories
- Bulbapedia. (2021a). Air balloon. https://bulbapedia.bulbagarden.net/wiki/Air_Balloon
- Bulbapedia. (2021b). Assault vest. https://bulbapedia.bulbagarden.net/wiki/Assault_Vest
- Bulbapedia. (2021c). Category:moves that deal direct damage. https://bulbapedia.bulbagarden.net/wiki/Category:Moves_that_deal_direct_damage
- Bulbapedia. (2021d). Choice band. https://bulbapedia.bulbagarden.net/wiki/Choice_Band
- Bulbapedia. (2021e). Choice specs. https://bulbapedia.bulbagarden.net/wiki/Choice_Specs
- Bulbapedia. (2021f). Counter (move). Retrieved July 15, 2014, from [https://bulbapedia.bulbagarden.net/wiki/Counter_\(move\)](https://bulbapedia.bulbagarden.net/wiki/Counter_(move))
- Bulbapedia. (2021g). Court change (move). [https://bulbapedia.bulbagarden.net/wiki/Court_Change_\(move\)](https://bulbapedia.bulbagarden.net/wiki/Court_Change_(move))
- Bulbapedia. (2021h). Critical hit. https://bulbapedia.bulbagarden.net/wiki/Critical_Hit
- Bulbapedia. (2021i). Damage. <https://bulbapedia.bulbagarden.net/wiki/Damage>
- Bulbapedia. (2021j). Defog (move). [https://bulbapedia.bulbagarden.net/wiki/Defog_\(move\)](https://bulbapedia.bulbagarden.net/wiki/Defog_(move))
- Bulbapedia. (2021k). Fire punch (move). [https://bulbapedia.bulbagarden.net/wiki/Fire_Punch_\(move\)](https://bulbapedia.bulbagarden.net/wiki/Fire_Punch_(move))
- Bulbapedia. (2021l). Focus sash. https://bulbapedia.bulbagarden.net/wiki/Focus_Sash
- Bulbapedia. (2021m). G-max steelsurge (move). [https://bulbapedia.bulbagarden.net/wiki/G-Max_Steelsurge_\(move\)](https://bulbapedia.bulbagarden.net/wiki/G-Max_Steelsurge_(move))
- Bulbapedia. (2021n). Grounded. <https://bulbapedia.bulbagarden.net/wiki/Grounded>
- Bulbapedia. (2021o). Gyro ball (move). [https://bulbapedia.bulbagarden.net/wiki/Gyro_Ball_\(move\)](https://bulbapedia.bulbagarden.net/wiki/Gyro_Ball_(move))
- Bulbapedia. (2021p). Heavy-duty boots. https://bulbapedia.bulbagarden.net/wiki/Heavy-Duty_Boots
- Bulbapedia. (2021q). Leftovers. <https://bulbapedia.bulbagarden.net/wiki/Leftovers>
- Bulbapedia. (2021r). Life orb. https://bulbapedia.bulbagarden.net/wiki/Life_Orb
- Bulbapedia. (2021s). List of moves that cause entry hazards. https://bulbapedia.bulbagarden.net/wiki/List_of_moves_that_cause_entry_hazards
- Bulbapedia. (2021t). List of pokémon by base stats (generation vii). [https://bulbapedia.bulbagarden.net/wiki/List_of_Pok%C3%A9mon_by_base_stats_\(Generation_VII\)](https://bulbapedia.bulbagarden.net/wiki/List_of_Pok%C3%A9mon_by_base_stats_(Generation_VII))
- Bulbapedia. (2021u). Rapid spin (move). [https://bulbapedia.bulbagarden.net/wiki/Rapid_Spin_\(move\)](https://bulbapedia.bulbagarden.net/wiki/Rapid_Spin_(move))
- Bulbapedia. (2021v). Spikes (move). [https://bulbapedia.bulbagarden.net/wiki/Spikes_\(move\)](https://bulbapedia.bulbagarden.net/wiki/Spikes_(move))

- Bulbapedia. (2021w). Stat. <https://bulbapedia.bulbagarden.net/wiki/Stat>
- Bulbapedia. (2021x). Stealth rock (move). [https://bulbapedia.bulbagarden.net/wiki/Stealth_Rock_\(move\)](https://bulbapedia.bulbagarden.net/wiki/Stealth_Rock_(move))
- Bulbapedia. (2021y). Sticky web (move). [https://bulbapedia.bulbagarden.net/wiki/Sticky_Web_\(move\)](https://bulbapedia.bulbagarden.net/wiki/Sticky_Web_(move))
- Bulbapedia. (2021z). Switcheroo (move). [https://bulbapedia.bulbagarden.net/wiki/Switcheroo_\(move\)](https://bulbapedia.bulbagarden.net/wiki/Switcheroo_(move))
- Bulbapedia. (2021aa). Toxic spikes (move). [https://bulbapedia.bulbagarden.net/wiki/Toxic_Spikes_\(move\)](https://bulbapedia.bulbagarden.net/wiki/Toxic_Spikes_(move))
- Bulbapedia. (2021ab). Trick (move). [https://bulbapedia.bulbagarden.net/wiki/Trick_\(move\)](https://bulbapedia.bulbagarden.net/wiki/Trick_(move))
- Bulbapedia. (2021ac). U-turn (move). [https://bulbapedia.bulbagarden.net/wiki/U-Turn_\(move\)](https://bulbapedia.bulbagarden.net/wiki/U-Turn_(move))
- Bulbapedia. (2021ad). Weakness policy. https://bulbapedia.bulbagarden.net/wiki/Weakness_Policy
- Bulbapedia. (2021ae). Weather. <https://bulbapedia.bulbagarden.net/wiki/Weather>
- Bulbapedia. (2022). Charizard (pokémon). <https://bulbapedia.bulbagarden.net/wiki/Charizard>
- Database, P. (2022). Pokémon types & type chart. <https://pokemondb.net/type>
- Fandom. (2016). Ashs pikachu. Retrieved January 27, 2022, from https://pokemon.fandom.com/de/wiki/Ashs_Pikachu?action=history
- Grover, A., & Leskovec, J. (2016). Node2vec: Scalable feature learning for networks.
- Honko. (2022). Pokémon damage calculator (generation vi) [[Online; accessed 24-January-2022]]. <https://calc.pokemonshowdown.com/randoms.html?gen=6&mode=randoms>
- How an a.i. is becoming the world's best pokémon player. (2021). <https://www.youtube.com/watch?v=rhvj7CmTRkg>
- Huang, D., & Lee, S. (2019). A self-play policy optimization approach to battling pokémon. *2019 IEEE Conference on Games (CoG)*, 1–4. <https://doi.org/10.1109/CIG.2019.8848014>
- Konda, V. R., & Tsitsiklis, J. N. (n.d.). Actor-critic algorithms, 7.
- Lee, S., & Togelius, J. (2017). Showdown ai competition. *2017 IEEE Conference on Computational Intelligence and Games (CIG)*, 191–198. <https://doi.org/10.1109/CIG.2017.8080435>
- Mikolov, T., Sutskever, I., Chen, K., Corrado, G., & Dean, J. (2013). Distributed representations of words and phrases and their compositionality.
- Mitarai, R. (2021). Random battles competitive guide. <https://www.smogon.com/articles/randbats-guide-gen8>
- OpenAI. (2018). Openai five. <https://blog.openai.com/openai-five/>
- pmariglia. (2022). Github/pmariglia-showdown. <https://github.com/pmariglia/showdown>
- Programming ai for pokémon showdown + bot battle royale! (2021). <https://www.youtube.com/watch?v=rhyj7CmTRkg>
- Sahovic, H. (2022a). Github/poke-env. <https://github.com/hsahovic/poke-env>
- Sahovic, H. (2022b). Poke-env: Baselines.py. https://github.com/hsahovic/poke-env/blob/master/src/poke_env/player/baselines.py
- Schulman, J., Moritz, P., Levine, S., Jordan, M., & Abbeel, P. (2018). High-dimensional continuous control using generalized advantage estimation.
- Schulman, J., Wolski, F., Dhariwal, P., Radford, A., & Klimov, O. (2017). Proximal policy optimization algorithms.
- Smogon. (2021). Github/damage-calc. <https://github.com/smogon/damage-calc>
- Veekun. (2021). Parasect. <https://veekun.com/dex/pokemon/parasect>

- Wikipedia. (2022a). Expectiminimax — Wikipedia, the free encyclopedia [[Online; accessed 24-January-2022]]. <http://en.wikipedia.org/w/index.php?title=Expectiminimax&oldid=952355961>
- Wikipedia. (2022b). Pokémon. Retrieved January 27, 2022, from <https://en.wikipedia.org/wiki/Pok%C3%A9mon>

Appendix

A. First Appendix Section

ein Bild

Figure A.1.: A figure

...

Ich versichere, dass ich die vorstehende Arbeit selbstständig und ohne fremde Hilfe angefertigt und mich keiner anderer als der in den beigefügten Verzeichnissen angegebenen Hilfsmittel bedient habe. Alle Textstellen, die wörtlich oder sinngemäß aus Veröffentlichungen Dritter entnommen wurden, sind als solche kenntlich gemacht. Alle Quellen, die dem World Wide Web entnommen oder in einer digitalen Form verwendet wurden, sind der Arbeit beigefügt.

Weitere Personen waren an der geistigen Leistung der vorliegenden Arbeit nicht beteiligt. Insbesondere habe ich nicht die Hilfe eines Ghostwriters oder einer Ghostwriting-Agentur in Anspruch genommen. Dritte haben von mir weder unmittelbar noch mittelbar Geld oder geldwerte Leistungen für Arbeiten erhalten, die im Zusammenhang mit dem Inhalt der vorgelegten Arbeit stehen.

Der Durchführung einer elektronischen Plagiatsprüfung stimme ich hiermit zu. Die eingereichte elektronische Fassung der Arbeit ist vollständig. Mir ist bewusst, dass nachträgliche Ergänzungen ausgeschlossen sind.

Die Arbeit wurde bisher keiner anderen Prüfungsbehörde vorgelegt und auch nicht veröffentlicht. Ich bin mir bewusst, dass eine unwahre Erklärung zur Versicherung der selbstständigen Leistungserbringung rechtliche Folgen haben kann.

Place, XX. Month 20YY

.....
(Julian Schubert)