

Bachelor Thesis

Julius-Maximilians-
UNIVERSITÄT
WÜRZBURG

Design, Implementation and Evaluation of Different Strategies for Playing Pokémon Battles

Julian Schubert

Institute for Computer Science
Chair for Computer Science VI

Prof. Dr. Frank Puppe

First Reviewer

Dr. Markus Krug

First Advisor

Submission

07. February 2022

www.uni-wuerzburg.de

Abstract

Pokémon is the highest-grossing video game franchise and revolves around collecting Pokémons with the ultimate goal of becoming “the very best, like no one ever was”. This thesis provides a detailed introduction to competitive Pokémon battles and summarizes previous research on the topic. Current search-based agents lack proper long term planning as the large quantity of possible states limits the amount of turns that can be looked ahead. A combination of existing *Minimax*- and rule based-agents that allows for planning and execution of a match plan is introduced. After playing over 1,400 battles on the *Pokémon Showdown* ranked ladder against human players, the bot achieved a mean Elo of 1,227 and a maximum of 1,432. While the agent only won 273 out of 1,000 games against the best open source implementation, analysis of battles lead to the conclusion that defeats are often caused by lack of features that due to time restrictions were not fully implemented. Additionally, the initial state of over 70,000 games was rated and then played out by two baseline agents. Results showed that both a random agent and an agent that prioritizes the most damaging move have a lower win rate on unfavorable board ratings while winning more games on a positive rating.

Contents

1. Introduction	1
2. Background	5
2.1. Basic rules	5
2.1.1. Types	5
2.1.2. Moves	5
2.1.3. Pokémon	6
2.1.4. Switching	7
2.1.5. Status condition	8
2.1.6. Items	9
2.1.7. Field effects	9
2.1.8. Order of events	11
2.1.9. Damage calculation	11
2.1.10. Effective Stats	13
2.2. Dynamaxing	13
2.3. Hazards	13
2.3.1. List of entry hazards	13
2.3.2. Hazard counterplay	14
2.4. Showdown random battles	15
2.4.1. Sets	15
2.4.2. Items	15
2.4.3. Amount of states	16
2.5. Pokémon Matchups	17
2.5.1. Check and Counter	17
3. Related Work	19
3.1. Baseline Agents	19
3.2. Breadth-first search	19
3.3. MiniMax	19
3.4. Rule based agents	20
3.5. Other approaches	20
3.6. Self-Play Policy Optimization Approach	21
3.6.1. State of the art	23
3.7. Supervised Approach	25
4. Approach	27
4.1. Communication with Pokémon Showdown	27
4.2. Gathering Information about the enemy Pokémon	27
4.3. Scoring the current game state	28
4.4. Stages of the game	29
4.5. Determining matchups	30
4.6. Predictions	31

4.7. Dynamaxing	31
5. Evaluation	33
5.1. Challenges for evaluation	33
5.1.1. Randomness of battles	33
5.2. Agents	35
5.2.1. HerrDonner	35
5.2.2. HerrGewitter	35
5.3. Results	38
5.3.1. Ranked results	38
5.4. Investigating Replays	39
5.4.1. Pro Replays	40
6. Conclusion	41
List of Figures	43
List of Tables	45
Listings	47
Acronyms	49
Bibliography	51
Appendix	55
A. Figures	55

1. Introduction

Pokémon (an abbreviation for **Pocket Monsters**) is a media franchise managed by *The Pokémon Company*, a company founded by *Nintendo*, *Game Freak* and *Creatures* (Wikipedia, 2022b). Figure 1.1 shows an image of *Pikachu*, one of the most popular Pokémons. Among multiple movies and series, a large variety of Pokémons games was released, starting with *Pokémon Red* and *Pokémon Blue*, which were published on 28. September 1998, also known as generation 1 (gen1)-games. In order to promote trading between players, *Nintendo* published two very similar games at the same time. Both releases differ in the available Pokémons and some minor changes to the story. Therefore, in order to collect all available Pokémons, players either had to buy two copies of the game or trade with a player that owns the counterpart. As of writing, there are eight major releases, also called *mainline games* with the latest being *Pokémon Sword* and *Pokémon Shield*.

While the graphics evolved a lot over the years (see figures 1.2 and 1.3), the key concept of the game remained mostly unchanged. The player starts his journey in his hometown to become the *Pokémon Champion*, which is the highest known level of rank for a Pokémon trainer (Bulbapedia, 2022g). In order to achieve this goal, the player has to catch wild Pokémons and use them to create a team of up to six individual Pokémons which he then trains to unleash their full potential. This Thesis focuses exclusively on the battling aspect of the game as there are detailed lists of all locations and secrets there is to explore within the games. Pokémons battles are turn based where both players, unlike in for example chess, make their decisions at the same time. While the core battle mechanics are very simple, the game provides a lot of depth which led to the formation of a strong competitive battling scene. As catching and training Pokémons to a level where they are competitive viable is a time-intensive task, battle simulators such as the open source project *Poké-*



Figure 1.1.: Pikachu, one of the most popular Pokémons (Fandom, 2016)




Figure 1.2.: Screenshots of Pokémon Red, the earliest Pokémon game. The game was released in 1996 for the Game Boy and has since been re-released multiple times. Image source: (Nintendo, 2022)



Figure 1.3.: Screenshots of Pokémon Sword, the latest major Pokémon game released for the Nintendo Switch in 2019. While the graphics have evolved, the core concept of the game, exploring and collecting Pokémons, remained the same.

mon Showdown (Smogon, 2022b) have arisen. On these fan made platforms, players have access to all available Pokémons and can compete against each other in a large variety of formats. Figure A.1 shows a battle between two players on *Pokémon Showdown*. The screenshot was taken from a generation 8 random battle (gen8randbats) between *Buckfae* and another player¹. The screenshot was taken at the start of turn 19. Currently, the *Mewtwo* of *Buckfae* is at full health (indicated red) while the opposing *Kommo-o* has 33% of hit points (hp) left. Below the avatars of both players, the team is displayed, marked yellow for player two. While player one has three Pokémons remaining, his opponent has four members alive. The Pokéball (bottom right) indicates a to the player unknown Pokémons. Moreover, the blue box highlights the status modifications as well as boosts of the enemy Pokémons which will be covered in section 2.1.3 and 2.1.10 respectively. Below the game window, the possible choices of the player are displayed. *Mewtwo* has access to the moves *Fire Blast*, *Recover*, *Psystrike* and *Nastyplot*. In addition to *Dynamaxing* which will be covered in section 2.2 the player also has the option to switch to either *Sirfetch'd*

¹The other name was blurred

or *Vespiquen*. Lastly, on the right-hand side a log of the previous turns can be found. In addition to the moves a Pokémon can use, a Pokémon has one ability and can hold an item that yields advantages in battle.

Pokémon battles are simultaneous, meaning that both players move at the same time. Furthermore, each player has imperfect information about the possible team of his opponent. Currently, there are approximately 1.5×10^{39} (see section 2.4.3) possible starting positions for random battles which combined with the high dimensional feature vectors of states and the complex rules of the game result in Pokémon presenting an interesting challenge for AI to tackle.

This thesis gives an introduction to the main mechanics of the Pokémon genre while no prior knowledge of the game is required. Chapter 3 introduces and compares different approaches and provides an in depth summarization of the current state-of-the-art battling AI. Furthermore, an agent combining previously existing rule- and *Minimax*-based implementations is developed to improve long term planning. Challenges of evaluating Pokémon AI and the results of this research are presented in 5.

2. Background

2.1. Basic rules

One of the key aspects of the Pokémon games is to battle other Pokémon. In the mainline games, you can have up to six Pokémon in your team, also known as party. There is the option to swap a Pokémon with another Pokémon, but you can not have more than six Pokémon at any point in your team. When playing the original games, you explore the world to find more Pokémon and use your team to defeat wild Pokémon and other Pokémon trainers. This thesis focuses on random battles taking place on *Pokémon Showdown*. In a random battle, both competitors get a team of six random Pokémon. At the start of the battle, you know each of your six Pokémon but only the currently active enemy Pokémon. Every turn, both players can choose to either use a move of their currently active Pokémon or switch to another Pokémon. Moves can either deal direct damage to the enemy Pokémon or yield other advantages like increasing the damage dealt by the next move. Moves will be covered in more detail in section 2.1.2. Each Pokémon has an amount of hp. The hp of a Pokémon can be lowered by attacking it with a move. If the hp of a Pokémon drops to zero, it faints and can not be used in this battle anymore. A player wins, once all enemies fainted.

Note that in the mainline games there is the possibility to heal or even revive a fainted Pokémon during battle using *Healing Items* like *Revive* or *Hyper Potion*. In competitive Play, only *Held Items* like *Leftovers* are allowed. Items will be explained in depth in section 2.1.6.

2.1.1. Types

Pokémon implements a rock-paper-scissors (rps)-like system. Each Pokémon has either one or two of 18 types. Like *rock* beats *scissors* but loses to *paper*, *Fire*-type Pokémon have an *advantage* against *Grass*-type Pokémon but are on a *disadvantage* against *Water*-types. This is also called a type being *strong / weak* against another type. Figure 2.1 shows how different Pokémon types interact with each other. Unlike in rps, type modifiers will be multiplied if a Pokémon has two types. For instance, a *Fire*-type attack will deal 4 times the damage against *Parasect* as *Parasect* has the types *Grass* and *Bug* (Veeekun, 2021).

2.1.2. Moves

Moves can be split up into three categories: *physical*-, *special*- and *status*-moves. While *physical*- and *special*-moves usually deal damage to the opponent Pokémon, *status*-moves

DEFENSE → ATTACK ↴	NOR	FIR	WAT	ELE	GRA	ICE	FIG	POI	GRO	FLY	PSY	BUG	ROC	GHO	DRA	DAR	STE	FAI
NORMAL													½	0			½	
FIRE		½	½		2	2						2	½		½		2	
WATER		2	½		½				2			2	2		½			
ELECTRIC			2	½	½				0	2					½			
GRASS		½	2		½			½	2	½		½	2		½		½	
ICE		½	½		2	½			2	2				2		½		
FIGHTING	2				2			½			½	½	½	2	0	2	2	½
POISON					2			½	½					½	½	0	2	
GROUND		2		2	½			2		0		½	2			2		
FLYING			½	2		2					2	½					½	
PSYCHIC					2	2					½				0	½		
BUG		½			2		½	½		½	2			½		2	½	½
ROCK		2				2	½		½	2		2					½	
GHOST	0									2			2		½			
DRAGON													2			½	0	
DARK							½				2		2		½		½	
STEEL		½	½	½		2						2				½	2	
FAIRY		½				2	½					2	2				½	

Figure 2.1.: Damage multiplier for all types. The damage dealt will be multiplied with two if a move is *super effective* (highlighted green), halved if the defender *resists* the attacking type (highlighted red) or set to zero if the move has no effect on the attacker (highlighted gray) (Pokémon-Database, 2022)

can for example change the weather, which plays a role in damage calculation explained in section 2.1.9, inflict status effects, raise or lower the stats of a Pokémon. In addition, a move also has exactly one of the 18 possible types. Often, *status*-moves are also referred to as *buffing* moves if they grant a benefit to the user or as *debuffing* if they have a negative effect on the opponent. While some moves will never miss, most moves have an *accuracy* of less than 100%. If a move misses, the turn is effectively skipped.

2.1.3. Pokémon

A key-concept of Pokémon battles are the *stats* of a Pokémon. The stats, among other factors which will be covered in section 2.1.9, determine how much damage a Pokémon is capable of dealing and receiving before fainting as well as how fast it can act in battle.

Explanation of stats

HP: The hp determines how much damage a Pokémon can receive before fainting.

Attack: The attack stat (atk) determines how much damage a Pokémon will deal when using a *physical*-move.

Defense: The defense stat (def) determines how well a Pokémon can resist against *physical* attacks.

Sp. Atk: The special attack stat (spa) determines how much damage a Pokémon will deal when using a *special*-move.

Sp. Def: The special defense stat (spd) determines how well a Pokémon can resist against special attacks.

Speed: The speed stat (spe) determines how fast a Pokémon can act. Usually, the Pokémon

with a higher spe will move before the slower one. The exact order of actions in battles is covered in section 2.1.8.

Determination of stats

The total stat of a Pokémon is calculated as described in equation 2.1 and equation 2.2 (Bulbapedia, 2022l).

$$HP = \left\lfloor \frac{(2 \times Base + IV + \lfloor \frac{EV}{4} \rfloor) \times Level}{100} \right\rfloor + Level + 10 \quad (2.1)$$

$$OtherStat = \left\lfloor \left(\frac{(2 \times Base + IV + \lfloor \frac{EV}{4} \rfloor) \times Level}{100} + 5 \right) \times Nature \right\rfloor \quad (2.2)$$

Base: Refers to the base stat of a Pokémon. Two Pokémons of the same species will always have the same base-stats. As seen in figure 2.2, a *Charizard* will always have a base-atk of 84.

Level: As mentioned in section 2.1, the goal of the mainline games is to create a team of six Pokémons and to make that team stronger by fighting other Pokémons. If a Pokémon defeats enough other Pokémons, it grows a Level. The maximum level of a Pokémon is 100. In *Pokémon Showdown*, the level of a Pokémon is set at the start of the battle and won't increase (Mitarai, 2021).

Nature: A Pokémon has a nature. Most natures enhance the growth of one stat, while hindering the growth of another. After all other calculations are finished, the stat that the Nature enhances will be 110% of what it would be without the Nature, and the stat hindered will be 90% of its normal value (Bulbapedia, 2022l). In this thesis nature can be neglected as all Pokémons in random battles have a neutral nature, meaning no stat is enhanced or hindered (Mitarai, 2021).

IV: Refers to the individual values (iv) of a Pokémon. These cause two Pokémons of the same species to have different stats (Bulbapedia, 2022l). Pokémons in *Pokémon Showdown* will always have the best possible iv stat, 31, unless it is a disadvantage for the Pokémon, then it will be zero (Mitarai, 2021). For example, having a high spe is undesirable for a Pokémon that uses the move *Trick Room* as this move reverses the move order so that Pokémons with a lower spe stat attack first while those with a higher spe will attack last. Details regarding the order of events are explained in section 2.1.8.

EV: These are the effort values (ev) of the Pokémons. ev are what causes a trained Pokémon to have higher stats than an untrained counterpart of the same level. For every 4 ev gained, a level 100 Pokémon will have 1 extra point in the given stat. A Pokémon can earn up to 510 ev, but can not have more than 255 ev in a single stat (Bulbapedia, 2022l). Random Pokémons on *Showdown* will always have 85 ev in each stat, or 0 in the case that having a high stat being detrimental (Mitarai, 2021).

As an example, the stats for *Charizard* are denoted in figure 2.2.

2.1.4. Switching

Instead of using a move with the current Pokémon, the player also has the option to switch out the active Pokémon for another Pokémon in his party. Switching always takes place before the execution of moves. However, the player does not know whether the opponent is switching or using a move. Therefore, if the player decides to switch out a non-fainted Pokémon, the enemy gets to use his move on the new Pokémon. If a Pokémon faints, the player has to switch in a new Pokémon and then the next turns starts. This means that the opponent gains a one turn advantage if the player decides to switch out a healthy Pokémon, but won't get to attack an additional time if the Pokémon was defeated.

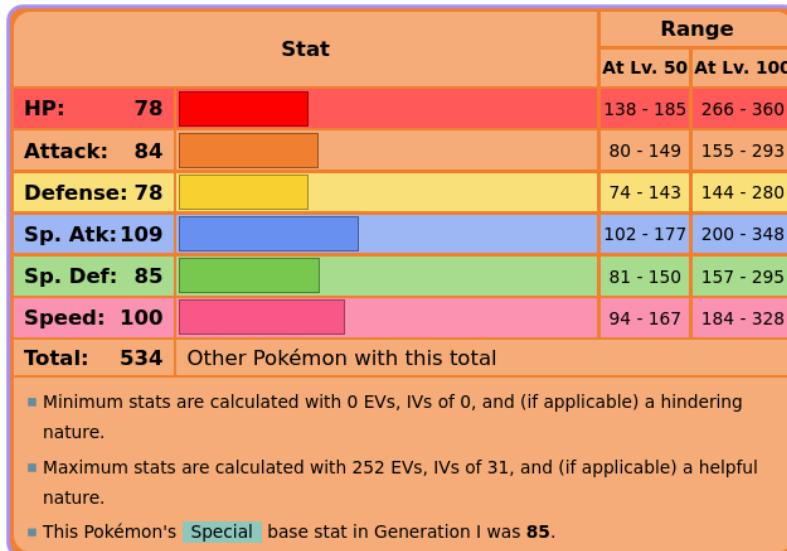


Figure 2.2.: Summary of the stats for *Charizard*. The left column shows the *base stats* of the Pokémons. To make comparison with other Pokémons easier, the second column shows these stats as bar chart. The last two columns show the possible ranges of final stats with ev and iv taken into account.

2.1.5. Status condition

Moves may inflict so-called *status conditions* that affect a Pokémon negatively. The most important status conditions are

- **Burn:** If a Pokémon suffers from the status condition burn (brn), it will lose 1/8 of its total hp every turn. In addition to that, a burned Pokémon will only deal half as much damage when using a *physical* move.
- **Freeze:** If a Pokémon suffers from the status condition freeze (frz) it won't, with a few exceptions, be able to use moves.
- **Paralysis:** If a Pokémon suffers from the status condition paralysis (par) it won't be able to use the selected move 25% of the time and their spe is halved.
- **Poison:** If a Pokémon suffers from the status condition poison (psn) it will, with a few exceptions, take damage equal to 1/8 of its total hp at the end of every turn. A Pokémon can also be *badly poisoned*. Badly poison initially inflicts damage equal to 1/16 of the Pokémon's maximum hp, with the damage inflicted increasing by 1/16 each turn. This means that the Pokémon will take 2/16 damage on the second turn, 3/16 on the third turn (and so on).
- **Sleep:** If a Pokémon suffers from the status condition sleep (slp) it won't be able to use moves, except *Snore* and *Sleep Talk*.

At any point, a Pokémon can only suffer from one status condition at a time, this means that a burned Pokémon can not fall asleep. While a sleeping Pokémon wakes up after one to three turns and a frozen Pokémon has a 10% chance to unfreeze each turn, the other status conditions remain until they are cured. A status condition can either be cured by abilities like *Natural Cure*. This ability heals any status condition when the Pokémon is switched out. Items can also remove negative effects, for example, *Lum Berry* cures all status conditions but gets consumed in the process.

2.1.6. Items

A Pokémon can also hold an item that yields benefits in battle. There are various purposes that items can fulfill. For example, the item *Life Orb* boosts damage dealt by the holder's damaging move by 30%¹, but the holder takes damage equal to 10% of its maximum hp after it uses a damaging move² (Bulbapedia, 2021o). *Leftovers* restore 1/16 of the holder's maximum hp³ at the end of each turn whereas the item *Air Balloon* makes the holder *ungrounded*, which means that the holder is immune to *Ground*-type moves as well as several related effects (Bulbapedia, 2021b). The generation of items in Pokémon Showdown is described in more detail in section 2.4.2.

Important items

In this section, a quick introduction to the most important items is given.

- **Choice Band:** When held by a Pokémon, this item boosts the atk by 50%, but only allows the use of the first move selected. This effect resets when the holder is switched out (Bulbapedia, 2021d).
- **Choice Scarf:** When held by a Pokémon, this item boosts the spe by 50%, but only allows the use of the first move selected. This effect resets when the holder is switched out (Bulbapedia, 2021e).
- **Choice Specs:** When held by a Pokémon, this item boosts the spa by 50%, but only allows the use of the first move selected. This effect resets when the holder is switched out (Bulbapedia, 2021f).
- **Heavy-Duty Boots:** The holder is unaffected by the effects of entry hazards. Entry hazards are described in 2.3.
- **Assault Vest:** Raises the holders spd by 50% but also prevents the holder from selecting any status moves⁴ (Bulbapedia, 2021c).
- **Focus Sash:** If the holder has full hp and is hit by an attack that would otherwise cause fainting, it survives with 1 hp (Bulbapedia, 2022e).

2.1.7. Field effects

There are multiple *field effects* that affect combat. Field effect are different from items as they are not tied to a specific Pokémon but rather always effect the currently battling Pokémon. There are field effects that only target the side of one player as well as effects that target both sides.

Terrain

Terrain is set up by the respective move with identical name and last for five turns and effects both players. All of them are beneficial to *grounded* Pokémon. A Pokémon is *not grounded* if any of the following conditions apply: The Pokémon

- has the *Flying*-type
- has the ability *Levitate*
- is holding the item *Air Balloon*

¹This boost is approximated as $5324/4096 \approx 1.29980$

²Rounded down, but not less than 1

³Rounded down, but not less than 1

⁴Except *Me First*

- is under the effect of *Magnet Rise* or *Telekinesis*.

Grounded Pokémons are with a few exceptions those Pokémons, that are not *ungrounded*. A Pokémon will be grounded if any of the following conditions apply:

- The Pokémon is holding an *Iron Ball*
- The Pokémon is under the effect of *Ingrain*, *Smack Down* or *Thousand Arrows*.
- The field effect *Gravity* is in effect.

More information about grounding can be found at (Bulbapedia, 2021k) There are five different possible *terrain*-states.

- **None:** The default state, no other effects are applied.
- **Electric Terrain:** Grounded Pokémons can not fall asleep and the power of *Electric*-type moves is increased by 50%.
- **Grassy Terrain:** The HP of grounded Pokémons is restored by 1/16 of their maximum HP at the end of each turn. In addition, the power of *Grass*-type moves is increased by 50% and the moves *Earthquake*, *Magnitude* and *Bulldoze* halve in power.
- **Misty Terrain:** Protects all grounded Pokémons from status conditions. The power of *Dragon*-type moves is halved while in effect.
- **Psychic Terrain:** Prevents grounded Pokémons from being hit by priority moves. Priority moves will be covered in section 2.1.8. The power of *Psychic*-type moves is also increased

It is important to note that only one *terrain* can be active at a time, yet, *terrain* can coexist with other *field effects* like *weather*.

Weather

The *weather* is a set of mechanics that change the battle environment, activating abilities, modifying certain moves and potentially damaging the Pokémons in battle or affecting their stats. Only one type of weather may be present at a time, and only the most recent type of weather will take effect (Bulbapedia, 2022p). List of different *weather*-conditions:

- *Clear skies*: (also: *None*) Absence of weather, this is the default state
- *(Harsh-) Sunshine*: (also: *(Harsh-) Sunlight*) Strong sunlight shines on the battlefield
- *(Heavy-) Rain*: Rain falls on the battlefield.
- *Sandstorm*: Stinging sand whips across the battlefield. At the end of each turn, each Pokémon, with a few exceptions, takes damage equal to 1/16 of its maximum hp unless it is a *Rock*-, *Steel*- or *Ground*-type (Bulbapedia, 2021t).
- *Hail*: Pelting hail falls on the battlefield. At the end of each turn, each Pokémon, with a few exceptions, takes damage equal to 1/16 of its maximum hp unless it is an *Ice*-type (Bulbapedia, 2021m)

There are a few additional special weather conditions which will not be covered in this thesis as they only occur in very specific scenarios.

Other

The list below contains the remaining field effects that cover the entire field not covered so far:

- *Wonder Room*: swaps the def and spd of all Pokémons, but stat changes remain on their respective stat (Bulbapedia, 2021y).
- *Magic Room*: suppresses the effect of all items held by the Pokémons on the field (Bulbapedia, 2021s).
- *Gravity*: causes the field to undergo intense gravity which multiplies the accuracy stat of all Pokémons by $5/3$ (Bulbapedia, 2022f).
- *Trick Room*: reverses the move order within each priority bracket. Move priority will be covered in section 2.1.8.

(According to Bulbapedia: All last for 5 turns)

ToDo

Single Side Effects

Among *Hazards* which will be covered in section 2.3, the most important single side effects are *Reflect* which halves the damage done to Pokémons on the given side by *physical* moves (Bulbapedia, 2022i) and *Light Screen* which works equivalent but for *special* moves (Bulbapedia, 2021p).

2.1.8. Order of events

In battle, switching always is executed before moves meaning that if a player switches and his opponent chooses to use a move, the move will always hit the newly brought in Pokémons. Move order is determined based on *priority*.

Priority

Priority is a characteristic of moves, such that any move with a higher priority will always be performed first. When two moves have the same priority, the users spe stat will determine which one is performed first in battle, the Pokémon with the higher spe moves first. Each move has a hidden priority value in the game data, with values ranging from $+5$ to -7 . The vast majority of moves have the standard priority value of 0. A move with a positive priority is called *priority move*. Moves that have the same priority are said to be in the same priority bracket. An example for a move with priority $+2$ is *Extreme Speed*, a *Normal*-type move with a base-power of 80 (Bulbapedia, 2022c). Having a *priority-move* can be useful if the own active Pokémon has a lower spe-stat than the opponent with low hp as then the own Pokémon does not have to take an additional hit before defeating his enemy.

2.1.9. Damage calculation

The damage dealt by a move mainly depends on the *level* of the Pokémon that uses the move, its effective atk or spa stat, the opponent's effective def or spd stat and the move's effective power.

Precisely, the damage is calculated as follows (Bulbapedia, 2021i):

$$\text{Damage} = \left(\frac{\left(\frac{2 \times \text{Level}}{5} \right) \times \text{Power} \times A / D}{50} + 2 \right) \times \text{Targets} \times \text{Weather} \times \text{Badge} \times \text{Critical} \times \text{random} \times \text{STAB} \times \text{Type} \times \text{Burn} \times \text{other} \quad (2.3)$$

The only exception for this are moves that deal direct damage. A list of these moves can be found at Bulbapedia, 2016. In the following paragraph, parameters of the formula are covered in detail. An explanation of the *Type*-modifier is omitted since it is covered in section 2.1.1.

Level

Level refers to the level of the attacking Pokémon. In *Pokémon Showdown*, the level is displayed next to the name of the Pokémon (Bulbapedia, 2021i).

A / D

A is the effective atk stat of the attacking Pokémon if the used move is a physical move, or the effective spa stat of the attacking Pokémon if the used move is a special move.

D is likewise the effective def stat of the target if the used move is a physical move, or the effective spd of the target if the used move is a special move (Bulbapedia, 2021i).

Power

Power is the effective power of the used move. The *Base Power* of a move in *Showdown* can be seen when hovering over a move in the move list.

Note: The same move will always have the same base power. For example, *Fire Punch* will always have a base power of 75 (Bulbapedia, 2022d).

Weather

The *Weather* modifier is 1.5 if a *Water*-type move is used during *Rain* or a *Fire*-type move during *Harsh Sunlight*. The modifier is 0.5 if a *Water*-type move is used during *Harsh Sunlight* or a *Fire*-type move during *Rain* (Bulbapedia, 2021i).

Critical

In the latest generation, a critical hit (crit) deals 1.5 times the damage compared to a normal hit. If the crit rate is not increased, the chance of landing a crit is 1/24 (Bulbapedia, 2022a). Increasing crit rate, as well as other stats, will be explained in chapter 2.1.10.

Random

Random is a random integer percentage between 85% and 100%. Because of this, the same move may deal different damage in the same scenario (Bulbapedia, 2021i).

STAB

STAB stands for *Same Type Attack Bonus*. It is a multiplier of 1.5 if the current move is of the same type as the attacking Pokémon. Otherwise, it is 1.0 (Bulbapedia, 2021i).

Burn

Burn is 0.5 if the attacking Pokémon is burned, and the attack is physical⁵. Otherwise, it is 1.0 (Bulbapedia, 2021i).

Other

The *other* modifier is usually 1. A list of exceptions can be found at (Bulbapedia, 2021i).

⁵This does not apply if the attacking Pokémon has the Ability *Guts* or the used move is *Facade*

2.1.10. Effective Stats

When a stat is used in a calculation in battle, a number of modifiers may be applied during the calculation. During a battle, a Pokémon's effective stats may be raised or lowered by certain moves abilities and held items. Some attacks may only have a chance of lowering stats, while certain abilities and held items may require a triggering event to activate stat modifications.

The modifiers conferred by most moves operate on a sliding scale of *stages*. When a given stat is raised or lowered, its current stage is increased or decreased by the amount dictated by the move; up to a maximum of +6 or a minimum of -6. A given stat corresponds to a given multiplier that will modify the stat when it is used in calculations. The exact multipliers for stages are described in table 2.1 When playing the mainline games, the

Table 2.1.: Stage multipliers that are applied for atk, def, spa, spd and spe. For example, a Pokémon with atk raised by two stages will have his existing atk stat multiplied by two (Bulbapedia, 2022l).

-6	-5	-4	-3	-2	1	0	1	2	3	4	5	6
2/8	2/7	2/6	2/5	2/4	2/3	2/2	3/2	4/2	5/2	6/2	7/2	8/2

player has to memorize the current stat changes of a Pokémon. In contrast to that, *Pokémon Showdown* displays the current stat multiplier of both Pokémons. Stat changes caused by boosts do not persist between switches, meaning that the stages for all stats of a Pokémon are reset to zero upon being switched in. Furthermore, a Pokémon also has an *accuracy* and *evasion* stat. Both can also be increased and lowered in stages but do not use the same modifiers. In gen8randbats the player very rarely has influence on these multipliers which therefore will not be covered in depth. Additional information can be found at (Bulbapedia, 2021a).

2.2. Dynamaxing

Dynamaxing is a temporary transformation affecting Pokémon that was introduced in generation 8 (gen8). Dynamaxing increases a Pokémon's size drastically, as well as changing the moves of the Pokémon and doubling their max and current hp, but can only be used once per battle and ends after three turns or if the Pokémon is switched out.

2.3. Hazards

An *entry hazard* is a condition that affects one side of the field. It causes any Pokémon that is sent into battle on that side of the field to be afflicted by a negative effect. Entry hazards are usually created by status moves (Bulbapedia, 2021q).

2.3.1. List of entry hazards

Currently, there are five moves that create an entry hazard

Spikes

Spikes is a *Ground*-type entry hazard that causes the opponent to lose 1/8 of their maximum hp when they enter the field. This effect can be stacked up to three times. Two layers of spikes will deal 1/6 and three layers will deal 1/4 of the enemies maximum hp. Spikes are set by the move *Spikes* (Bulbapedia, 2022k).

Stealth Rock

The move *Stealth Rock* sets an entry hazard around the target Pokémon causing Pokémon on the target's field to receive damage upon being switched in. The amount of damage inflicted is affected by the effectiveness of the type *Rock* against the target. Unlike Spikes, this entry hazard does not stack. The damage taken from the victim's maximum is denoted in table 2.2 (Bulbapedia, 2022m).

Type effectiveness	Damage (max. hp)
0.25x	3.125%
0.5x	6.25%
1x	12.5%
2x	25%
4x	50%

Table 2.2.: Damage dealt to Pokémon by *Stealth Rocks* (Bulbapedia, 2022m)

Sticky Web

The entry hazard set by the *Bug*-type move *Sticky Web* lowers the opponents spe by one stage upon switching in (Bulbapedia, 2022n).

Poison spikes

Poison Spikes set by the *Poison*-type move *Toxic Spikes* cause the opponent to become poisoned. If two layers of spikes are set, the Pokémon instead becomes badly poisoned (Bulbapedia, 2021v).

Sharp steel

This entry hazard works very similar to Stealth Rock described in 2.3.1. However, Sharp steel can only be set by the *Steel*-type move *G-Max Steelsurge* which is the exclusive G-Max Move of *Gigantamax Copperajah*. The damage dealt by Sharp steel does not stack, the amount of damage dealt is based on the Type effectiveness of the *Steel*-type against the target. Exact damage modifiers can be found in table 2.3 (Bulbapedia, 2021j).

Type effectiveness	Damage (Max. hp)
0.25x	3.125%
0.5x	6.25%
1x	12.5%
2x	25%
4x	50%

Table 2.3.: Damage dealt to Pokémon by *Sharp Steel* (Bulbapedia, 2021j)

2.3.2. Hazard counterplay

There are some moves that can remove entry hazards. *Rapid Spin* (Bulbapedia, 2022h) removes entry hazards from the user's side of the field and *Defog* (Bulbapedia, 2022b) removes entry hazards on both sides of the field⁶. In addition, *Court Change* (Bulbapedia, 2021h) will exchange the entry hazards on each side of the field, along with other one-sided

⁶In older games *Defog* would only remove hazards on the target's side of the field. But as we only investigate the latest version, this won't be covered in detail.

field conditions. If a grounded *Poison*-type Pokéémon enters the battle, it will remove *Toxic Spikes*, described in 2.3.1, from its side of the field. Pokéémon holding the item *Heavy-Duty Boots* (Bulbapedia, 2021n) are unaffected by entry hazards, but grounded *Poison*-type Pokéémon can still remove *Toxic Spikes* even if they hold the boots (Bulbapedia, 2021q). Additionally, there are various other exceptions to hazards which won't be covered in detail here.

2.4. Showdown random battles

Pokéémon Showdown features a large variety of game modes, the most popular ones being *Sw / Sh Singles OU* and *Sw / Sh Singles Random Battle*. *Sw / Sh* refers to *Pokéémon Sword* and *Shield*, the latest *Pokéémon Games*. These two games form the eighth generation of *Pokéémon* games and are thereby also often referred as gen8-games. As *Pokéémon Showdown* solely focuses on battling every player has access to every available *Pokéémon*. This mainly plays a role in the *OU*-format. In this game mode, each player knows all enemy *Pokéémon* but does not know their exact builds, that is, the abilities, iv-, ev-combinations and items are unknown. As the success of a player in this format heavily depends on his team as well as which other teams are currently popular, this thesis will focus on gen8randbats. In this format, each player gets a team of six random *Pokéémon* and unlike in *OU* only knows the currently active opposing *Pokéémon*. It is important to note that there are many possible builds for a *Pokéémon* and therefore, information management plays a huge role. An example for this will be provided in section 4.2.

The development for battling bots is explicitly allowed and encouraged by the developers. Furthermore, there is no way to figure out whether the current opponent is a real human or an AI. This is beneficial as humans might behave differently when they know they are facing a machine potentially leading to invalid results.

2.4.1. Sets

As described in section 2.1.3, *Pokéémon* created for random battles usually have 85 evs and 31 iv in every stat with a neutral nature, meaning a nature that does neither boost nor hinder any stat (Mitarai, 2021). There are some cases where a high stat is not beneficial, an example would be the move *Gyro Ball*. Unlike most moves, the *Base Power* of this move described in the damage calculation described in 2.1.9 is not a fixed value. It is determined as described in 2.4 (Bulbapedia, 2021l).

$$\text{BasePower} = \min \left(150, \frac{25 \times \text{CurrentSpeed}_{\text{target}}}{\text{CurrentSpeed}_{\text{user}}} \right) \quad (2.4)$$

Since the damage dealt by *Gyro Ball* gets bigger, the lower the spe of the attacker, *Pokéémon* using this move have 0 ev and 0 iv in the spe stat.

Note: Being able to outspeed the opponent is extremely valuable, but the only two *Pokéémon* using *Gyro Ball*, *Stakataka* and *Ferrothron*, already have a very low spe stat and are slower than almost all other *Pokéémon* in random battles. A complete list of *Pokéémon* with their respective spe stat can be found at (Bulbapedia, 2021r).

This knowledge can be exploited to gather additional information about the enemy. Section 4.2 describes how this is achieved.

2.4.2. Items

Items in random battles are procedurally generated by *Showdown* and depend on the *Pokéémon*'s moves, base stats and ability. As stated in Mitarai, 2021, the exact implementation is "changed frequently with the intention of optimizing set generation", yet, item assignment follows these rules:

- Pokémons with 2 or fewer attacking moves will get *Leftovers*, or *Black Sludge* if *Poison*-type.
- Pokémons with 3 attacking moves will get *Life Orb*, if the sum of their base hp, def and spd is less than 275. Otherwise, these Pokémons get *Leftovers* or *Black Sludge*.
- Pokémons with 4 matching attacks get a *Choice* item which follows these rules:
 - Pokémons with four physical attacks or four special attacks, a base spe between 60 and 108 and base atk or spa of 100 or more can get a *Choice Scarf* 2/3 of the time. If the Pokémon does not meet one of the stat qualifications or does not get the 2/3 chance, they'll get *Choice Specs* or *Choice Band* instead.
 - Pokémons with 3 special attacks and the move *U-turn* always get *Choice Specs*. *U-turn* is a physical, *Bug*-type move that switches the user out after damage is dealt (Bulbapedia, 2022o).
 - Pokémons with *Trick* (Bulbapedia, 2021w) or *Switcheroo* (Bulbapedia, 2021u), both moves that allow to switch items with the opponent, will always get a choice item. If they meet the above-mentioned speed range, they will always get a *Choice Scarf*. Otherwise, they will always get *Choice Specs* or *Choice Band*.
 - Having priority moves will always prevent a *Choice Scarf* from being generated in all situations.
- Pokémons with 4 attacks that do not qualify for choice items, will get an *Assault Vest* if their $hp + def + spd \geq 235$. Otherwise, *Expert Belt*, *Leftovers* or *Life Orb* is generated.
- Pokémons that are weak to *Rock* will get *Heavy-Duty Boots* if they don't get a higher priority item, such as *Assault Vest* or a choice item. Pokémons that are four times weak to *Rock*, such as *Charizard*, will always get *Heavy-Duty Boots*. This is done as these Pokémons would otherwise lose up to 50% hp to the entry hazard *Stealth Rock* described in 2.3.1. The only exception is *Scyther*, which can get *Eviolite*.
- Pokémons in the lead slot will get *Focus Sash* if their $hp + def + spd < 255$, and they would otherwise get *Leftovers* or *Life Orb*.
- Pokémons that get a Speed-boosting move will be given a *Weakness Policy* if their $hp + def + spd \geq 300$, and they aren't four times weak to *Ground*. This item boosts the atk and spa by two stages each if hit by a super effective move. After that, the item breaks (Bulbapedia, 2021x).

There are also some species that will always roll the same item, either because it's their signature item or because doing so supports a niche ability or set. For example, Pikachu always has *Light Ball*.

2.4.3. Amount of states

At the time of writing, 457 different Pokémons are available in gen8randbats. On average, each Pokémon has approximately nine different possible builds which leads to 1.2×10^{13} different possible teams for one player. As for one of these teams, each of the six Pokémons could be in the first slot, this number has to be multiplied by 6 to get the total amount of possible teams for one player. Equation 2.5 shows how the total amount of possible starting states.

$$\left(\left(\binom{457}{6} \cdot 9^6 \right) \cdot 6 \right)^2 \approx 1.5 \times 10^{39} \quad (2.5)$$

2.5. Pokémon Matchups

Due to the typing system, there is no best Pokémon that is the best option in all situations. Therefore, we have to determine how good a Pokémon is against another Pokémon in a given situation. In this case, the *situation* refers to the current state of both Pokémon like current hp and status conditions as well as field effects like weather.

2.5.1. Check and Counter

MattL, 2014 refers to a Pokémon *checking* another Pokémon if it can beat the enemy Pokémon in every scenario and can safely be switched in at any point. A *counter* is also capable of defeating the enemy Pokémon but may lose in some situations. The most notable being if switched in without the previous active Pokémon fainting as this grants the opponent an additional attack.

The key difference between check and counter is that a check is also stronger if it takes damage once more while a counter is not guaranteed to win in this situation.

Note: Every check to a Pokémon is also always a counter while counter could also be a check, but is not guaranteed to.

3. Related Work

In this section, previous research regarding the topic of battles on Pokémon Showdown is described and compared.

3.1. Baseline Agents

A good way to get a rough idea on how well an agent performs is to compare it against a baseline agent. There are two popular baseline agents, the *Random*-Player and the *MaxDamage*-Player. While the *Random*-Agent always chooses either a random move or a random switch, the *MaxDamage*-Agent always picks the move with the highest base power. If no move is available, the agent will switch to a random Pokémon. This is roughly equal to the skill level of an inexperienced beginner human.

3.2. Breadth-first search

Given a root battle object representing the current game state, breadth-first search (bfs) explores the outcomes of all possible choices, treating these resultant states as child nodes. This algorithm traverses the game tree until it finds a state in which the enemy Pokémon is fainted. As a non-adversarial algorithm, the agent assumes that the agent does not move at all. This agent won 75 out of 90 total games played against a random agent (Lee & Togelius, 2017).

3.3. MiniMax

MiniMax builds upon bfs as this algorithm deals with adversarial paradigms by assuming the opponents act in their best interest (Lee & Togelius, 2017). There are multiple possibilities on how to implement the *MiniMax*-Algorithm for Pokémon games. The main difference in implementations lies in state evaluation and the assumptions about the opponents. Additionally, the amount of features taken into consideration have a considerable impact. In the implementation proposed by Lee and Togelius, 2017, a node represents the worst case scenario that would occur as a result of the current choice. The agent also employes alpha-beta pruning, ignoring any node in which the agents Pokémon faints. One drawback of this procedure is that, a Pokémon would never use the move *Explosion*, a very powerful *Normal*-type move that also faints the user. This move can for example be used if the active member is already at very low hp. The tree itself is traversed using a

greedy strategy, which terminates when a state in which the enemy Pokéémon is fainted is reached. Both the traversal order and the worst-case evaluation are performed using the evaluation function 3.1 (Lee & Togelius, 2017):

$$Eval = \frac{\text{current hp}_{\text{Own Pokéémon}}}{\max \text{hp}_{\text{Own Pokéémon}}} - 3 \cdot \frac{\text{current hp}_{\text{Enemy Pokéémon}}}{\max \text{hp}_{\text{Enemy Pokéémon}}} - 0.3 \cdot \text{depth} \quad (3.1)$$

A YouTube-Video released by RemptonGames, 2021 alters this evaluation function by increasing the rating of a game state when an enemy takes damage and decreasing the rating when a member of the own team takes damage. Both evaluation functions do not take hazards, status condition or boosts into account. Due to the similarity of both evaluation functions, both agents performed very similarly: The first agent described by Lee and Togelius, 2017 won 70 out of 90 total games against a random agent, resulting in a win-ratio of $\approx 86\%$. The second agent defeated the random player in 831 out of 1,000 total games, yielding a win-ratio of $\approx 83\%$. Due to the large amount of random number generation (rng) in battles, both agents are assumed to be at an equal level of play.

The current state of the art search-based algorithm was developed by *pmariglia* and is fully available at *pmariglia*, 2022. This approach uses the *ExpectiMiniMax*-Algorithm and takes hazards, boosts and status into consideration. In addition to “min” and “max” nodes of a traditional *MiniMax*-tree, this variant has “chance” nodes, which take the expected value of a random event occurring (Wikipedia, 2022a). Currently, the agent calculates two turns in advance, and reaches an Elo rating of 1,461 as well as a Glicko-1 rating of 1,633 in generation 7 (gen7) random battles on the *Showdown* ladder. In addition, *pmariglia* implemented an algorithm to estimate the item a Pokéémon is holding based on the damage it dealt.

3.4. Rule based agents

The YouTube-Video released by RemptonGames, 2021 also introduces two rule-based agents. *Smart Damage* was written by the author of the video and uses Pokéémon type and the spe-stat to determine the favorability of a matchup. On a bad matchup, this agent will switch to the team member with the best matchup. Otherwise, a simplified damage calculation is used to determine the move dealing the most amount of damage. *Simple Heuristics* was developed by, *Haris Sahovic*, the author of the *Poke-Env* library (Sahovic, 2022a). The implementation for this agent can be found at Sahovic, 2022b. This agent uses simple rules to decide the next move or switch and takes boosts as well as hazards into account. The results for both agents are denoted in table 3.1. The *MiniMax*-Agent in the table refers to the implementation of *Rempton Games*.

	<i>Random</i>	<i>MaxDamage</i>	<i>SmartDamage</i>	<i>MiniMax</i>	<i>Heuristics</i>
<i>Random</i>	N / A	897 / 103	957 / 43	831 / 169	992 / 8
<i>MaxDamage</i>		N / A	829 / 171	834 / 166	955 / 45
<i>SmartDamage</i>			N / A	331 / 669	720 / 280
<i>MiniMax</i>				N / A	181 / 819

Table 3.1.: Denotes how many games the Agent in the column won against the agent in the row RemptonGames, 2021.

3.5. Other approaches

The authors of the *Showdown AI Competition* (Lee & Togelius, 2017) compared many simple AI implementations with each other. Approaches not covered so far will be summarized in this chapter.

One Turn Lookahead

One Turn Lookahead is a heuristic agent designed to encapsulate a greedy strategy that prioritizes the damage output. The agent operates by estimating the damage dealt by all usable moves, including those usable by the agent's inactive but usable Pokémons. If the highest damaging move belongs to the active Pokémon, the agent will use that attack. If the most damaging move belongs to an inactive Pokémon, the agent will switch to that Pokémon (Lee & Togelius, 2017). Depending on implementation details, this agent is very similar to *MaxDamage* or *Rule Based* agents.

Type Selector

This is a variation of the *One Turn Lookahead*-Agent that utilizes a short series of if-else statements in its decision-making. At first, if the current Pokémon knows a move that drains the opponents hp to zero, this move is selected. Otherwise, the favorability of the current matchup is evaluated. If the current type matchup is undesirable, the agent will switch to the Pokémon with an acceptable type matchup. If no such Pokémon exists, the agent will default to the most damaging move (Lee & Togelius, 2017).

Pruned Breadth-First Search

This agent is designed to demonstrate a simple way to utilize domain knowledge as a cost-cutting measure. This is achieved by making modifications to the Breadth First Search agent. First, the algorithm does not simulate any actions that involve using a damaging move with a resisted type, nor does it simulate any actions that involve switching to a Pokémon with a subpar type matchup. Additionally, rather than selfishly assuming the opponent skips their turn in each simulation, the agent assumes its opponent is a *One Turn Lookahead*-agent and simulates accordingly (Lee & Togelius, 2017).

Results

Table 3.2 displays the results of the agents described in this section.

	<i>Random</i>
<i>One Turn Lookahead</i>	77 / 90
<i>Type Selector</i>	67 / 90
<i>Pruned BFS</i>	75 / 90

Table 3.2.: Results of other approaches against a random agent (Lee & Togelius, 2017)

3.6. Self-Play Policy Optimization Approach

Researchers from *New York University* (Lee & Togelius, 2017) were the first ones to apply *Q-Learning* to the field of Pokémon battles. Two agents using *Q-Learning* were developed: A single layer perceptron as well as a multi layer perceptron. Both agents were used to output the expected reward of all current moves and switches. Based on this, the best action was picked. Both agents were rewarded for defeating opponent Pokémon and punished for allowing one of its own Pokémon to faint. Because decisions made tend to have long term consequences, weights are updated using the last three (State, Action) pairs rather than the most recent pair only. Additionally, in order to promote exploration, the agent employs an epsilon-greedy selection policy, causing it to randomly override its decision with a probability of 0.1. The single layer perceptron was trained using the *Delta Rule*, while the multilayer perceptron was trained using *Delta Rule* plus *Backpropagation*.

The agent using a single layer perceptron won 90 out of 180 games against a random agent whereas the multi layer perceptron won 86 out of 180 games. Due to the large amount of rng in Pokémon games, more games would need to be played in order to confirm the superiority of the single layer perceptron in this particular use case. Randomness of games will be covered in more detail in section 5.1.1.

One year after the publication of Lee and Togelius, 2017, Chen and Lin, 2018 used proximal policy optimization (ppo) (Schulman et al., 2017) to train an agent. They also used embeddings to improve the representation of a Pokémon. Data available at alopez247, 2017 was used to create embeddings for each Pokémon. The Dataset contains stats of the first 721 Pokémon, each row contains the name, type(s), numerical stat data (such as hp, atk, spe) and some other data such as color, height and whether the Pokémon is considered legendary in game. To create embeddings for each Pokémon, the data was turned into a graph to be used with Node2Vec (Grover & Leskovec, 2016) which creates embeddings from graph data in similar to Word2Vec (Mikolov et al., 2013). This algorithm samples random walks of some number of nodes of a given graph. Using these random walks, a skip-gram model is created which can be trained to generate embeddings. The Pokémon graph consisted of the name, type(s), numerical attributes (total stats, hp, atk, def, spa, spd, spe) as well as two special nodes, *Legendary* and *Mega*¹. Lastly, Node2Vec was applied to the graph. Table 3.3 displays the discovered similarity using this approach. Here, a similar

Table 3.3.: Similar Pokémon within the embedding space of Chen and Lin, 2018.

Pokémon	Most similar Pokémon
<i>Bulbasaur</i>	Chikorita, Turtwig, Nuzleaf, Petilil, Exeggucete, Skiploom, Jumpluff, Oddish, Budew
<i>Caterpie</i>	Wurmple, Weedle, Kakuna, Metapod, Paras, Ledyba, Spinarak, Venonat, Silcoon
<i>Mewtwo</i>	Lugia, Mesprit, Mew, Victini, Celebi, Cresselia, Volcation, Ho-Oh, Uxie

Pokémon to *Mewtwo*, a *Psychic*-Type, is *Ho-Oh*, a *Flying-Fire*-Type. However, these two have entirely different strengths and weaknesses. In addition to that, they do not share a single move in their move set: In generation 6 (gen6), *Ho-Oh* has access to the following moves: *Aura Sphere*, *Calm Mind*, *Fire Blast*, *Ice Beam*, *Psystrike* and *Recover* whereas *Lugia*'s move pool consists of *Brave Bird*, *Earthquake*, *Flame Charge*, *Roost*, *Sacred Fire*, *Substitute* and *Toxic* (Honko, 2022). This inappropriate classification is likely caused by both *Pokémon* having similar stats and both being legendary.

Features are derived from the battle state in the simulator. At a high level, battles consist of two sides. Each side consists of a team of Pokémon, and each Pokémon has some set of moves. Each of these objects (battle, side Pokémon, and move) have attributes that are used to derive a feature vector. After experimenting with multiple network architectures, the authors settled on a three-layer fully connected neural network, each with 512 neurons and a ReLU activation function.

Chen and Lin, 2018 trained the agent against a random agent as well as a *MiniMax-Agent*² until the reward curve stabilized which was around 100 epochs where an epoch is a single battle between two players. A reward of +1 is given to the agent if it wins the battle, -1 if it loses the battle and 0 for all other cases. The average epoch reward after training convergence for this approach can be found in table 3.4 The authors of Chen and Lin, 2018

¹Mega is a mechanic similar to dynamaxing. Mega is not available in the latest version of the game and therefore won't be covered in detail

²The Authors don't provide implementation details of their *MiniMax* implementation

Opponent	Average epoch reward
<i>Random</i>	0.85
<i>Default</i>	0.85
<i>MiniMax</i>	-0.9

Table 3.4.: Average epoch rewards after training convergence for opponent agents (Chen & Lin, 2018). Agent outperforms random and default agent but fails to defeat the *MiniMax*-Agent.

describe their final agent as flawed as while the agent learned to switch Pokémons when the active Pokémon reaches low health, it almost always chooses to switch to the Pokémon in the last slot. In addition, this agent preferentially chooses the fourth move.

3.6.1. State of the art

In 2019, Huang and Lee, 2019 published a paper titled *A Self-Play Policy Optimization Approach to Battling Pokémons*. Due to its fine-grained analysis and the resulting agent performing on par with the state-of-the-art search based AI, this paper will be summarized in more depth.

Similar to OpenAI, 2018, the agent is represented using an actor-critic neural network. Actor-critic RL methods combine policy-based and value-based RL methods by predicting both policy and value for a given state, and then using the value prediction, the “critic”, as an estimate of expected return when updating the policy prediction, the “actor” (Konda & Tsitsiklis, 2000). The authors represent both actor and critic using a two-headed neural network which is trained via self-play RL (Huang & Lee, 2019).

Neural Network

Input to the neural network is the current state of the game, from the point of view of the player, represented as multi-level tree-like structure:

1. The *battle* consists of two *teams*, along with weather effects.
2. Each *team* consists of six *Pokémon*, along with side conditions described in section 2.1.7.
3. Each *Pokémon* has many features. Table 3.5 contains a partial list³.

The network has two outputs: a probability distribution $\pi \in \mathbb{R}^n$ over actions to take, and an estimate of player strength in the current state $v \in \mathbb{R}$. Huang and Lee, 2019 compute π as follows:

1. The network outputs an intermediate vector $p \in \mathbb{R}^n$. Each of the colored cells in figure A.2 correspond to an element of p
2. A probability distribution $\pi' \in \mathbb{R}^n$ is computed by using the softmax function:

$$\pi' = \frac{\exp(p_i)}{\sum_i \exp(p_i)}$$
3. As not every action is valid in every state, for example, a switch to a Pokémon is invalid if that Pokémon is already fainted, the authors ensure their agent has zero probability of taking invalid actions. To do this, they take a mask $s \in \{0, 1\}^n$ as part of the input, and renormalize probabilities to obtain $\pi : \pi_i = \frac{s_i \pi'_i}{s^T \pi'}$

³The authors state that this list is not complete, but no additional information is provided.

Feature	Type	Dims	Description
<i>species</i>	categorical	1×1023	e.g. Pikachu
<i>item</i>	categorical	1×368	e.g. Leftovers, Choice Band
<i>ability</i>	categorical	4×238	e.g. Rough Skin, Shadow Tag
<i>moveset</i>	categorical	4×731	e.g. Flamethrower, Surf
<i>lastmove</i>	categorical	1×731	The last move used
<i>stats</i>	continuous	6	hp, atk, def, spa, spd, spe
<i>boosts</i>	continuous	6	Temporary boosts for stats
<i>hp</i>	continuous	1	Current number of hp
<i>maxhp</i>	continuous	1	Number of hp at full health
<i>ppUsed</i>	continuous	4	# times a move was used
<i>active</i>	indicator	1	1 if Pokémon is active, else 0
<i>fainted</i>	indicator	1	1 if Pokémon has no hp, else 0
<i>status</i>	indicator	28	e.g. slp, brn, par
<i>types</i>	indicator	18	e.g. Bug, Fire
<i>volatiles</i>	indicator	23	e.g. Leech Seed, Perish Song

Table 3.5.: Features used to describe a single Pokémon battle (Huang & Lee, 2019).

The authors point out the following two key design decisions: First, a 128-dimensional entity embedding layer for each of the categorical variables is used. This enables capturing similarities between different moves, species and abilities without having to directly model their, often complicated, effects. Second, the parameters for computing p from above are shared among all n actions. The resulting network is described by figure A.2 and contains 1,327,618 parameters in total (Huang & Lee, 2019).

Training the network

Training was done serially: After $m = 7680$ games per iteration, the neural network parameters are updated using the $2m$ self-play matches as training data to obtain new neural network parameters. A reward of +1 for a win and -1 for a loss is assigned at the end of the match. To speed up learning, a dense reward signal using reward shaping was constructed. Auxiliary rewards are assigned based on events that occur over the course of the match. For example, a reward of -0.0125 is added when a Pokémon of the agent faints, and a reward of +0.0025 whenever the player's Pokémon makes a super effective move⁴.

To update the neural network, the authors use ppo (Schulman et al., 2017), which optimizes an objective function that combines expected reward, accuracy of state value prediction, and a bonus for high entropy policies. To reduce the variance of policy gradient estimates, *Generalized Advantage Estimation* (Schulman et al., 2018) is used.

After 500 iterations of the training loop, 3,840,000 self-play matches had been played by the neural network. Training was performed using Google Cloud Platform over the course of 6 days with an approximated cost of \$91 USD.

Evaluation

The refined embeddings as well as the complex network architecture lead to this agent outperforming the RL-Agent described in Chen and Lin, 2018. Furthermore, the agent played 1,000 games against the open source agent of *pmariiglia* mentioned in 3.3 Table

⁴While the authors do not provide an exact definition, we assume that a move is regarded as *super effective*, if the type modifier described in Damage calculation is bigger than one

Opponent	Wins	Losses
<i>random</i>	995	5
<i>max-damage</i>	929	71
<i>max-damage-typed</i>	829	171
<i>pmariglia</i>	612	388

Table 3.6.: Performance of the agent developed by Huang and Lee, 2019

3.6 describes the performance of the agent developed by Huang and Lee, 2019. While the authors do not provide the resulting Elo of the agent, they state that a *Glicko-1* rating of 1,677 was reached. It is important to note that this agent played against an older version of *pmariglia* which was released in May 2019. Results of this agent described in section 3.3 and section 5 used the latest release as of February 2022.

3.7. Supervised Approach

As of writing, there is only one publication researching supervised learning for Pokémon battles, a YouTube video uploaded by TheThirdBuild, 2021. Unfortunately, the video does not cover much technical and implementation details, possibly to reach a broader audience on YouTube. The author obtained two million replays of human games from the creators of *Pokémon Showdown*. Using these replays, the input vector for a game contained the following features is created:

- **Pokémon attributes:** hp, whether the Pokémon is active, status condition and stat boosts.
- **Player's side attributes:** side conditions (like *Light Screen*), entry hazards (like *Stealth Rocks*), volatile conditions like *Leech Seed* and the last used move
- **Battlefield attributes:** weather, pseudo-weather⁵ and terrain

Using these features, a neural network using TensorFlow for JavaScript was then trained to predict the win chance of a given player at a given state of the game. This network was able to predict the winner at a given state with an accuracy of 81%. In addition to that, the model was able to predict the next switch in with an accuracy of 86%, the next move with a certainty of 93% and whether the player would switch with an accuracy of 88%. Unfortunately, the author neither states the game type of the replays nor the exact architecture and evaluation method of the model.

This model was then used to pick the best move in a given scenario which allowed the agent to play games on the *Pokémon Showdown* ladder. Unlike other agents quoted in this thesis, this agent does not play gen8randbats on *Pokémon Showdown*, but rather *Gen 8 OU Singles*.

Furthermore, *Inverse Damage Calculation* is introduced. As the exact item of enemies are unknown, the author predicts the given item of a Pokémon by comparing the damage dealt of a move with the expected damage for the move modified by possible items. Despite no concrete implementation details given, we assume this approach to work similar to the implementation of *pmariglia*, 2022 described in section 3.3. According to the video, this agent reached a maximum Elo of 1,630 which ranks the agent among the best 3.6% of players.

⁵The video does not give an exact definition of the word *Pseudo-Weather*

4. Approach

(Short introduction. Describe idea. This builds on top of related work as combines minimax with rulebased) (Predictions: Use Move against most likely next enemy. No prediction switching (yet)) ToDo
ToDo

4.1. Communication with Pokémon Showdown

The communication with *Pokémon Showdown* is handled using Sahovic, 2022a. This library provides a lot of the core functionality needed, like accessing the current Pokémon in battle as well as switch and move options. However, it does not provide functionality for damage calculation. We use the *Pokémon Damage Calculator* (Smogon, 2022a), a node library written by the smogon-team for that. Communication between the two libraries is implemented by capturing stdout and stdin using the *subprocess* python library.

4.2. Gathering Information about the enemy Pokémon

As mentioned in section 2.4, the same Pokémon can occur in various different builds, meaning the combination of moves, abilities and items. Knowing the exact build is crucial for the decision-making process. Consider the following example:

- **Player1** has an active *Charizard* with *Heavy-Duty Boots* and 150hp remaining on the field.
- **Player2** has just sent out a *Drapion* with 160hp remaining.
- The *Charizard* is faster but can not kill the enemy *Drapion* in one turn as his move *Fire Blast* deals between 127 and 151 damage to the *Drapion*.
- Therefore, if **Player1** decides to attack, *Drapion* is guaranteed to survive this turn and can attack *Charizard* as well.

In this scenario, the optimal strategy for **Player1** depends heavily on the move set of the enemy *Drapion*. Possible moves for *Drapion* are:

- *Aqua Tail*: A damaging *Water*-type move
- *Earthquake*: A damaging *Ground*-type move
- *Knock Off*: A damaging *Dark*-type move

- *Poison Jab*: A damaging *Poison*-type move
- *Swords Dance*: A move to raise the own atk by two stages
- *Taunt*: A move that makes the afflicted Pokémon unable to use status moves
- *Toxic Spikes*: A move that sets an entry hazard.

Hereby it is important to note that *Drapion* only knows the move *Aqua Tail*, if it knows four total damaging moves¹. In the given scenario, **Player1** should switch out his *Charizard* if the enemy *Drapion* knows the move *Aqua Tail* as this attack would kill *Charizard*. We can determine whether the enemy knows *Aqua Tail* based on his item:

- If *Drapion* rolls two status moves, it will have the item *Black Sludge* and therefore doesn't know *Aqua Tail*. Because *Drapion* is already damaged, we know that it has this item if it healed 1/16% of his max hp in his last turn.
- If *Drapion* rolls one status move, it will have the item *Life Orb*. If *Drapion* already attacked, we know if it has a *Life Orb* or not as this item causes it to lose 10% of his maximum hp after an attack.
- If *Drapion* has neither *Black Sludge* nor *Life Orb*, it has to have a *Choice Band* and as this item will only generate if the Pokémon knows four matching attacks, and therefore has to know the move *Aqua Tail*.

Therefore, the player should switch out his *Charizard* if the *Drapion* has neither *Black Sludge* nor *Life Orb* as then it is guaranteed to have a *Choice Band*, knows *Aqua Tail* and would therefore be able to defeat *Charizard*. In any other case, the **Player1** can safely kill *Drapion* using his active *Charizard*.

Implementation details

The first step to predicting enemy sets is to determine all possible sets and the likeliness of each individual set. In order to achieve this, a script that starts a battle between an information gathering player and a random agent was created. In the next step, the script extracts all builds of all Pokémon and stores them. Then, it forfeits, and a new battle is started. Once enough battles are fought, the script will store the builds as well as how often they appeared in text files, one file for each Pokémon.

In actual battles, if a new Pokémon enters the enemy side, we assume it to have the most likely build for its species. Once more information becomes available on items, moves and abilities, we rule-out non-matching builds and always assume the enemy to have the remaining most likely build.

4.3. Scoring the current game state

In order to not only rate the current board state, but also individual Pokémon, we implement the following scoring algorithm:

$$score(e_{i,j}) = \text{Expected Damage that Pokémon } i \text{ will deal to Pokémon } j \quad (4.1)$$

The expected damage is the damage dealt if both Pokémon behave optimal in the amount of turns that the bot looks into the future. Section 4.5 covers how the optimal moves are determined.

$$value(i) = \sum_{j \in \text{Enemy Pokémon}} score(e_{i,j}) \quad (4.2)$$

¹This information was extracted from all possible *Drapion* builds. Details on how these builds were obtained are covered in section 4.2

Using this, we can also introduce a *value* for each of our Pokémons where a higher value implies a more important Pokémon. It is important to note that scores are determined independently of each other meaning that we do not take into account damage taken by the attacker. This does imply that this metric does not determine how good the Pokémon is if it has to battle *all* enemy Pokémons but rather against how many other Pokémons it *could* be used. This is done as the order in which a Pokémon battles multiple Pokémons plays a critical role. The reasoning behind this as well as the determination of an optimal order is explained in section 4.4. This metric also has multiple flaws because it only takes the damage dealt to the enemy into account, other important factors like damage received, healing, the availability of status moves and hazards is not considered. Using this rating, we can rank our team based on importance which is used in the switching routine described in section 4.4.

4.4. Stages of the game

We divide the game into two phases, the first one being the *Discover*-Phase whereas the second phase is called *Defeat*-Phase. Our goal and therefore our play style, is different in both phases.

Discover Phase

At the beginning of the game, we play safely until we know our opponent's entire team. In other words, we try to gather information about the enemy team while sacrificing as little hp as possible. In this stage, we act according to these rules:

1. Kill the opponent if we are guaranteed to kill him this turn. This either leads to us defeating the enemy Pokémons, or possibly new information if the enemy switches.
2. Healing our Pokémons. If we have a healing move that will heal us more than the expected damage we receive this turn, and we are not at full hp we will heal our Pokémons. Doing so will force the enemy to switch as we are otherwise gaining an advantage over him. (**Go into a bit more depth here?**) ToDo
3. If we have a hazard setting move available, we will use this move as they will help us in the *Defeat*-Phase. Other beneficial side-conditions like *Light Screen* will be set as well.
4. Using moves that inflict status to the opponent like brn.

If none of these conditions apply, we decide whether to switch out our Pokémons or not. If our current Pokémon is a *check* or *counter* against the enemy Pokémons, we will not switch. Otherwise, we switch to a *check*, or *counter* if present. Next, we check if the current matchup is unfavorable. This applies, if the enemy is expected to survive the current matchups for two turns longer than we do, meaning that our Pokémons would not be able to defeat the enemy, even if it was allowed to attack two additional times. If this is the case, we determine the next action as follows:

We start by determining the *score* of each of our Pokémons as described in equation 4.1 and exclude our two most valuable Pokémons from the next steps. This is done as we assume to depend heavily on these Pokémons to defeat other enemies. Next up, we pick the Pokémons with the lowest *score* that fulfills the following criteria:

- The Pokémons has to survive for at least three turns against the enemy
- At least one of the following criteria apply:
 - It is able to set a *hazard* or other beneficial field effect that is not already present.

- It can inflict a status condition on the enemy.

If no Pokémon matches these criteria, the Pokémon with the lowest score is picked instead. Score is calculated as described in section 4.3.

In any other case, the currently active Pokémon will use the best move calculated in 4.5.

Defeat Phase

The *Defeat*-Phase starts as soon as we know all enemy Pokémon. At the beginning of this stage, we have to create a match plan for defeating the enemy. The main goal is to figure out which Pokémon to use against which enemies, especially the best order to send them into battle. Figure A.3 shows a possible end-game scenario. In this simple example, we assume both players to each have three Pokémon with full hp and no status conditions. Each circle represents two Pokémon battling each other. Therefore, the circle in the top center indicates, that our *Roserade* is currently fighting against the enemy *Charizard*. As *Charizard* is a *Fire*-type Pokémon which has a type advantage against his enemy, our *Roserade* will faint in this matchup, expressed by the underlined name. This means that *Player1* has to switch in a new Pokémon. A green circle around a node illustrates that the first player has to make a turn as his Pokémon fainted, and a red circle marks a decision for *Player2*. The first line below the names of the two opposing Pokémon displays the amount of hp the first team has left after this battle took place. As *Roserade* fainted, it has zero hp left whereas the other two Pokémon in his team are still at full health. Below that, the remaining hp of the second team is displayed. *Charizard* is *expected* to survive the battle with an *expected* amount of 141 hp remaining. Next, we have two possible options remaining, we can either send *Kyogre* or *Azumarill* to defeat the enemy *Charizard*. Taking a look at the leaves of the tree, the *value* of a leaf node is -1 if the enemy wins the battle and 1 if we win the battle. The value of a non-leaf node is the sum of the *values* of the children nodes. The *value* of two means that we expect to win the battle. As a result of this, the best choice in our example scenario is to use *Azumarill* to defeat the already damaged *Charizard*. The *Minimax*-Tree was generated using the most likely builds scraped from replays and does not use the default set of Smogon, 2022a. For example, the most likely build for *Charizard* contains the move *Roost* (Bulbapedia, 2022j) which heals the Pokémon.

4.5. Determining matchups

In order to determine whether to attack or to switch, the optimal moves for a Pokémon against another Pokémon need to be calculated. As stated before, the amount of possible combinations combined with the non-deterministic nature of the game makes it unfeasible get the optimal move combination by simulating every possible combination of actions and reactions over the span of multiple turns. Therefore, we determine the optimal moves for a Pokémon using this simplified method:

We start by generating all possible move combinations for a Pokémon with a given length². Then, we simulate the outcome of the battle. If the attacking Pokémon would use these moves in the defined order given the enemy would not do anything³. Here, we also take boosting, items, status effects and the possibly changing field state into account. The combination resulting in the lowest amount of turns until the enemy faints is selected. It is important to note that this is not necessarily the move that the bot will use in the next turn as in the *Discover*-Phase healing, status and other beneficial effects are prioritized. In order to improve performance, only matchups including the currently active Pokémon are re-evaluated each turn.

²The final implementation used a length of 3

³The option of not doing anything in a turn does not exist in Pokémon, if possible, the player is always forced to either select a move or switch

4.6. Predictions

Predicting the next enemy action plays a huge role in competitive Pokémon. In many situations, it is very likely that the opponent will switch to another Pokémon to avoid a bad matchup. For example, players often withdraw *Fire*-type Pokémon if the opponent knows a *Water*-type move as this will deal twice the amount of damage due to the type multiplier discussed in section 2.1.1. Therefore, players often choose to pick a move countering the expected Pokémon the opponent will send out or even switch themselves in anticipation of an enemy action. Our agent mimics this behavior in the following way:

The agent keeps track how often the opponent chose to withdraw his active Pokémon when it is checked and how often he decided to stay in. In order to make the actions taken by the agent less predictable, rng is used on favorable matchups to decide how to act. The chance that the agent predicts an enemy switch is zero if this is the first time that our Pokémon counters an enemy and otherwise as described in equation 4.3.

$$\text{Chance make a prediction} = \frac{\# \text{ The enemy switched when countered}}{\# \text{ Total times the enemy was countered}} \quad (4.3)$$

Currently, the agent assumes the enemy to switch in a random Pokémon that counters the active Pokémon and will then pick the move, that will yield to the best outcome for this matchup. Flaws and possible improvements to this simple routine will be discussed in section 5.2.2.

4.7. Dynamaxing

At the time of writing, *Memezboi69* is the top ranked player on the Pokémon Showdown ladder. He described the use of dynamaxing as very situation dependent as it can either be used offensive or defensive and proposed a simple rule set that can be build upon: Dynamax if the current matchup is favorable, the entire enemy team is known and the active Pokémon has more than 70% hp remaining. In addition, the agent will dynamax if the current Pokémon is the last remaining Pokémon with full hp. *Memezboi69* also laid out a more sophisticated rule set, which is described in section 5.2.2, that can be applied in future iterations of the agent.

5. Evaluation

5.1. Challenges for evaluation

Different researchers use different metrics to evaluate the performance of their agents. There are multiple factors that increase the difficulty of properly evaluating the performance.

5.1.1. Randomness of battles

As teams are generated randomly, one player often ends up with a slightly better team than his opponent. In extreme cases, one player may not even have a chance at winning the battle. While battling our agent during the evaluation process, one particular game stood out as the first Pokémon of **Player1** was capable of defeating the entire enemy team. The Pokémon of **Player1** was a *Volcarona* with the following moves:

- *Fire Blast*, a damaging *Fire*-Type move
- *Quiver Dance*, a *Bug*-Type move that boosts the user's spa, spd and spe by one stage each.
- *Bug Buzz*, a damaging *Bug*-Type move
- *Roost*, a move that restores half of the user's maximum hp

This Pokémon was able to defeat the entire enemy team with little to no possible counter play: The first enemy Pokémon, *Leafeon*, a *Grass*-Type Pokémon was killed in one hit using *Fire Blast* after damaging *Volcarona* using *X-Scissor*.

Next, *Glalie*, an *Ice*-Type was sent into battle. *Glalie* uses his best move, *Earth Quake* which brings *Volcarona* to 52% hp. As the enemy does not pose much threat to *Glalie*, **Player1** decided to boost using *Quiver Dance*. Now, *Volcarona* is faster than his enemy and kills it again in one hit using *Fire Blast*.

Then, *Mr. Mime (Galar)* is sent into battle. He fails to pressure *Volcarona* and therefore, **Player1** can heal his Pokémon using *Roost* and further boost using *Quiver Dance*. After defeating *Mr. Mime (Galar)*, *Volcarona* is back to 84% HP and boosts of 2.5 spa, 1.5 spd and 2.5 spe. Boosted this high, *Volcarona* can one-shot both the enemy *Volcarona*, *Pheromosa* and the dynamaxed *Scraggy* using *Fire Blast*.

To eliminate the impact of these extreme cases, evaluation of agents against other agents should be done using multiple hundreds to thousands of games against each other.

In order to quantify and better visualize the randomness of battles, over 70,000 gen8randbats were generated, and the team information was stored. Next, the board rating was calculated. Let e be a Pokémon and C the set of all enemy Pokémon that are countered by e . We compute the score of a team member e as follows:

$$\text{member-score}(\text{Pokémon: } e) = \sum_{c \in C} 1 \quad (5.1)$$

The score of a team p is defined by:

$$\text{team-score}(\text{Team: } p) = \sum_{e \in p} \text{member-score}(e) \quad (5.2)$$

Finally, the board rating of two teams p_1 p_2 is defined as the difference between their respective team score

$$\text{board rating}(\text{Team: } p_1, \text{Team: } p_2) = \text{team-score}(p_1) - \text{team-score}(p_2) \quad (5.3)$$

Therefore, a low board rating indicates that the player has fewer counters than his opponents. Finally, the generated games were played out by *Random-Player* as well as *MaxDamage-Player*. The win-rate of both agents at a given board rating is displayed in figure 5.1. Most notably, fair games, meaning games with a board rating close to zero,




Figure 5.1.: Win rate of the baseline agents at a given board rating. Both agents move have a lower win rate on unfavorable board ratings while winning more games on a positive rating.

are gaussian-distributed with the majority of the matchups being almost fair according to our metric. The blue line indicates the win rate of a *Random-Player* when playing against another *Random-Player* at the given board rating, indicated green is the win rate

of two *MaxDamage*-Players. There are multiple important observations: Firstly, at negative board ratings, both agents have a significantly lower win rate against the opponent while positive board ratings indicate a high chance of winning. Despite collecting over 70,000 data points, board ratings worse than -20 or better than $+20$ will be ignored in this analysis as the large irregularities in these areas indicate an insufficient amount of data points. The steeper slope of the green line is likely explained by the rng of the actions performed by random agents having a much higher influence than the actual team composition. As a random agent rarely makes the correct choice, it fails to utilize the full potential of a stronger team. **(More refined metric: Same results, longer to simulate as better agents required. Showdown replays do not contain all information necessary.)** ToDo

5.2. Agents

During this thesis, two different agents were developed, *HerrDonner* and *HerrGewitter*.

5.2.1. HerrDonner

This agent was designed to establish a good baseline and to demonstrate the capabilities of a very simple rule set. The agent is capable of looking multiple turns into the future. To determine the moves to be used, the agent estimates the damage dealt for every possible move, leading to *HerrDonner* being an improved *MaxDamage*-Player as it behaves similarly, but has a more refined damage calculation algorithm. No drawback moves that heal the agent, set hazards or field conditions or inflict status conditions are not considered unless they result in the highest amount of damage dealt. Also, stat changes are not taken into account, neither for damage calculation nor for rating of matchups. This results in the bot often spamming moves like *Draco Meteor*, a special *Dragon*-type move that deal a lot of damage but also lower the users spa by two stages resulting in the move dealing less damage every time it is carried out. When the agent is forced to switch, it will switch to a check if available. If no check is available, a counter is sent into battle if one exists. Otherwise, a random Pokémon will be picked.

At the start of each turn, *HerrDonner* will evaluate the current matchup, a matchup is deemed unfavorable, if the current Pokémon is neither check nor counter to the current enemy. On a bad matchup, the bot will switch to an available check or counter. If neither is available, the bot won't switch and try to defeat the current opponent with his active team member.

Dynamaxing is implemented in a very simple and naive way: The agent will always dynamax the active Pokémon as soon as more than four enemy Pokémon are known. Lastly, if the current Pokémon is dynamaxed, the agent will not switch, even if the current matchup is not favorable. The challenges of properly using this mechanic will be explained in section 5.2.2.

5.2.2. HerrGewitter

HerrGewitter behaves like described in section 4. Here, the most notable differences between both agents are highlighted, and limitations of this agent are discussed.

Firstly, *HerrGewitter* takes more factors into consideration. Now, damage calculation is done in regard to current stat changes and status conditions. Additionally, abilities and items are taken into account. Furthermore, recoil from moves, healing both from items like *Leftovers* and moves like *Recover* are not neglected anymore.

Switching and the selection of moves is done as described in section 4.

These improvements lead to *HerrGewitter* avoiding mistakes of *HerrDonner*. For example, this agent will burn a physical attacker using *Will-O-Wisp* (Bulbapedia, 2022q) in order to reduce damage taken over the next turns. The agent will also boost and heal itself

in favorable situations which stalls the game and forces the opponent to react. Another major improvement is that the agent switches out the current Pokémon if stat changes resulted in an unfavorable matchup which is especially important as stat changes reset on a swap. However, there are still a lot of features that *HerrGewitter* is lacking.

Weather and Field effects

One such feature is a proper support for weather and field effects in the damage calculation as well as in the *Minimax*-Algorithm. For instance, a consequence of this is the agent lacking awareness of the fact that a *Fire*-Type move deals 1.5 more damage during *Harsh Sunlight*.

Hazards

Currently, the agent will always try to set a non-present Hazard in the early game as this usually results in a long term benefit. There are however some notable exceptions to this that are not yet implemented:

- The agent will always set as many hazards as possible in the early game, even if the current matchup is unfavorable, including always setting up to two layers of spikes. A small test on human players indicated that this leads to slightly better results than only setting hazards on good matchups, but due to the very small sample size, future work is needed to determine the best strategy for setting hazards.
- The agent does not take the damage taken by hazards into account when switching Pokémon.
- The agent will always use *Toxic Spikes* even if the opponent has a *Poison*-Type Pokémon on his team that will remove this hazard upon being switched in.
- The agent will use Hazards even if the current enemy is known to have a hazard-clearing move like *Defog* (Bulbapedia, 2022b).
- The agent will not clear hazards.

Choice Items

As mentioned in section 2.1.6, Pokémon holding a *Choice*-item are locked into using always the same move until they are switched out. The agent has two major flaws in regard to these items: When the active Pokémon of the agent is holding a choice item and is already locked into a move, the agent is not aware of the fact that once the Pokémon is switched out, it will regain access to his other moves which leads to an incorrect prediction for future matchups. As described in section 4.5, the only matchups re-evaluated on a given turn are matchups that include one of the currently active Pokémon. The following example illustrates how this design decision leads to issues on Pokémon with *Choice*-items:

In the given scenario, our agent has an active *Garchomp* which is locked into using *Earth Quake*. The *Garchomp* also has access to the *Rock*-Type move *Stone Edge*. This turn *Butterfree*, a *Bug* / *Flying*-Type Pokémon is sent into Battle. As the *Ground*-Type move *Earth Quake* has no effect on *Butterfree*, the agent will switch out *Garchomp* for another Pokémon. In the current implementation, matchups for *Garchomp* are not re-evaluated. While this won't lead to problems in the early game, this results in an incorrect *Minimax* calculation as for matchups involving *Garchomp* and any non-active opponent, *Garchomp* is still assumed to only have access to the move *Earth Quake*. In this scenario, the agent would fail to realize that *Garchomp* also has access to *Stone Edge* and would incorrectly assume *Garchomp* to lose all matchups against *Flying*-Type Pokémon.

While this behavior rarely effects battle, the agent failing to notice that an enemy is

choice-locked has more often a negative impact on the battle: If the enemy is known to be choice-locked into *Earth Quake* we can safely switch a *Flying*-Type into battle. This applies especially if the enemy Pokémon is known to have the *Rock*-type super effective move *Stone Edge* as the enemy can not use this move until switched out and back in again. In this scenario, the agent wrongfully would not prefer to switch in a *Flying*-Type Pokémon due to the threat posed by *Stone Edge*. Switching in a Pokémon resisting *Earth Quake* in this scenario forces the enemy to switch to another Pokémon. This gained turn advantage can either be used to land an extra move on the next opponent, set hazards, beneficial field conditions, inflict status conditions or boost the current Pokémon.

Damage Calculator

The current implementation relies on the Pokémon Showdown Damage Calculator. As of now, this open source project does only support direct damage dealt by attacking and lacks functionalities like recoil, healing from items and moves. While most of these features were added to *HerrGewitter*, some moves are still not properly implemented. For example, the move *Counter* has a move priority of -5 and works as follows: If the last mount of damage dealt before the use of *Counter* is greater than zero and was dealt by a physical move, *Counter* will do twice as much damage to the opponent. Otherwise, the move will miss. Additionally, *Counter* has a lot of extra rules regarding other special moves in place (Bulbapedia, 2021g). Issues like these are especially obvious on *Wobbuffet* as all of his four most likely moves, *Mirror Coat*, *Encore*, *Counter* and *Encore* are very useful yet do not deal any damage and are not implemented yet which leads the agent to believe that this Pokémon is bad in every possible matchup and has no good use scenarios whatsoever.

Dynamaxing

As previously stated, dynamaxing is a very complex mechanic that requires sophisticated long term planning for proper usage. According to *Memezboi69*, the safest usage of this mechanic is to dynamax in a late stage of the game on a good matchup. He stated multiple reasons for this: Dynamaxing at a later stage of the game reduces the likeliness of an opponent being able to properly react due to the limited amount of Pokémon alive. Another flaw of *HerrGewitter* is the fact that defensive dynamaxing is not yet taken into consideration. If the opponent dynamaxes offensively, a good way to minimize damage received when no counter is available, is dynamaxing defensively and using the move *Max Guard* which protects the user from moves. If a Pokémon has access to a status move, this move will be replaced by *Max Guard* when dynamaxing. Lastly, dynamaxing can under some circumstances also be used in an earlier stage of the game if the current matchup is favorable.

Predictions

When play testing the bot, predictions were noticeable but only seemed to occur rarely. This is likely caused by the optimal move against the current Pokémon being equal to the best attack targeting the predicted switch. In order to compete with higher ranked players, more investigation on the prediction routine is required. Future work could also include switching in anticipation of the opponent withdrawing his active Pokémon.

Other special cases

This list contains more currently unhandled cases which will be addressed in future versions:

- The Pokémon *Ditto* can transform itself into the Pokémon of the current opponent.

- The Pokémon *Zoroark* can transform itself into another team member.
- The ability *Trace* changes the ability of a Pokémon to the ability of his opponent.

MinMax

The *MiniMax*-Algorithm described in paragraph 4.4 only supports changes in health but ignores other important factors such as boosts and status conditions. Therefore, the agent will not recognize the possibility to weaken a very strong physical attacker like *Garchomp* by burning it first and then defeating it with another Pokémon. A simple way to include brn into this algorithm is to multiply the expected damage dealt by a burned Pokémon with 0.5.

5.3. Results

5.3.1. Ranked results

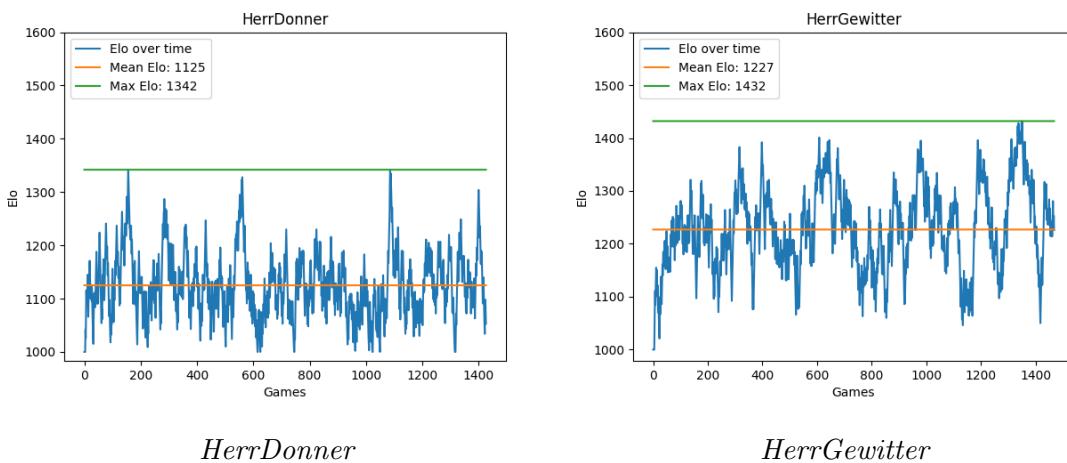


Figure 5.2.: Elo of *HerrDonner* and *HerrGewitter* after playing 1400 ranked games respectively

Both agents were evaluated at the same time over the span of multiple days by playing over 1400 ranked games each against human opponents on Pokémon Showdown. Figure 5.2 shows the Elo rating of *HerrDonner* and *HerrGewitter* over time. In order to make a comparison of both agents more easy, figure 5.3 shows the smoothed Elo of both agents over time. Smoothing was achieved by dividing the Elo history in chunks of size 20 and then plotting the average Elo of each chunk.

```

1  step = 20
2  smoothed_elo = []
3  for i in range(0, len(elo_history), step):
4      sec = elo_history[i: i + step]
5      smoothed_elo += [(i, sum(sec) / len(sec))]
```

Listing 5.1: Smoothing Elo values

The first thing to note is that *HerrGewitter* has a higher mean Elo (1227) and a higher max Elo (1432) than *HerrDonner* who achieved an average Elo of 1125 and peaked at an Elo of 1342. (**No footnote in table? How to properly format this?**) Table 5.1 shows the performance of different agents when directly competing against each other.

ToDo

It is very important to note that while the agent of (Huang & Lee, 2019) played against an older version of against (2022) than our bot did and is therefore not included in this

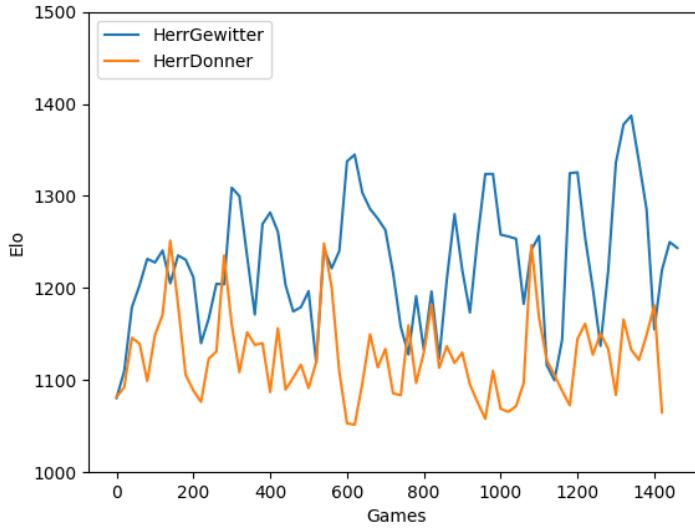


Figure 5.3.: Smoothed Elo of *HerrDonner* and *HerrGewitter* over the course of 1400 ranked games.

Table 5.1.: Battle results of 1,000 games between the agents.

Pmariglia: (pmariglia, 2022)

DQN: As described in (Huang & Lee, 2019)

	Random	MaxDmg	Pmariglia	DQN	Donner	Gewitter
Random	N / A	-	-	995 / 5	992 / 8	993 / 7
MaxDmg		N / A	-	929 / 71	906 / 94	951 / 49
Pmariglia			N / A	612 / 388 (*)	-	273 / 727
DQN				N / A	-	-
Donner					N / A	580 / 420
Gewitter						N / A

ToDo

table as this possibly leads to misleading interpretations. (**Currently included. Figure this out!**)

Each entry displays the results of a thousand games played between both agents. The first number indicates the amount of games the agent in the current column won. There are multiple things to point out here:

While *HerrDonner* and *HerrGewitter* achieved almost similar results when battling a random agent, *HerrDonner* performed notably worse against the *MaxDamage-Agent* which always picks the move with the highest base power. Also, *HerrGewitter* only managed to win 58% against *HerrDonner* despite notably better results against human players, described in section 5.4, possibly because *HerrGewitter* fails to properly exploit weaknesses of *HerrDonner* and has less weaknesses that humans can use to gain an advantage. At the time of writing, neither *poke-env* nor replays created by *Showdown* allow the retrieval of the Glicko-1 rating. After playing for 1,400 games, *HerrDonner* ended up with a Glicko-1 rating of 1,438 and *HerrGewitter* achieved a ranking of 1,509.

5.4. Investigating Replays

In addition to playing against humans on the *Showdown* ranked ladder, the bot was also evaluated by playing against *Jonathan Krebs* and *Markus Krug*. While *Markus* was able to

	<i>HerrDonner</i>	<i>HerrGewitter</i>
<i>Jonathan</i>	28 / 23	37 / 23
<i>Markus</i>	N / A	4 / 11

Table 5.2.: Performance of both agents against human players. *HerrGewitter* posed a harder challenge but is unable to gain a lead against *Markus*, an experienced human player.

reliably beat *HerrGewitter*, the agent provided a significant challenge and required careful planning to overcome. Replays retrieved from these games confirm assumptions made about the bot in the previous sections. *HerrGewitter* is very good at picking the most damaging move, but can be exploited using mechanics not yet implemented.

5.4.1. Pro Replays

Memzboi69, the top ranked player, agreed to play three games against *HerrGewitter*. While our agent lost all three games, we were able to gain some very interesting insights from these games:

Game One

In the first game, the bot made a huge mistake that ultimately lead to defeat. The agents *Shuckle* faced a *Mr. Mime (Galar)* which had access to *Rapid Spin*, a move that removes hazards. As *Shuckle* has access to the moves *Sticky Web* and *Stealth Rock*, the agent tried to set both hazards which were immediately removed. In the extra turn, *Mr. Mime (Galar)* boosted its speed stat and slowly killed his opponent using *Freeze-Dry*. This increased spe led to an advantage that the bot was not able to make up for. A future version of the agent will need to include a better hazard routine to prevent this kind of scenarios.

Game Two

Game two was lost due to a lack of counters against the opponents *Gengar*. Prior to sending in this Pokémon, *Memzboi69* set the entry hazard *Sticky Web* which decreases the spe by one stage on switch in. After boosting spa by two stages using *Nasty Plot*, *Gengar* was able to defeat the four remaining Pokémons of the agent. Further investigation revealed another flaw in the logic of the agent. When switching in the *Discover Phase* with no check or counter available, the agent ranks all possible options based on how much damage they are capable of dealing as described in section 4.3. If no Pokémon exists that is capable of surviving for more than two turns, the Pokémon with the lowest score will be sent into battle. In the given scenario, the agent sent out *Klefki* as his last Pokémon. This Pokémon always has access to either *Thunderwave*, a move that applies par or *Toxic* that poisons the opponent. While the chances of winning would still be very low when sending out this Pokémon, bringing in *Klefki* would have increased the odds. Additionally, the agent did not take the entry hazard into account which lead to the wrong assumption that *Zamazenta* is faster than *Gengar* and could therefore kill the opponent. This resulted in *Zamazenta* getting killed in one shot.

Game Three

The last battle was the closest and ended with *Memzboi69* having only two Pokémons remaining. The agent even managed to set up a burned *Obstagoon*, boosting atk to stage two and def to stage three. Unfortunately, this Pokémon was then killed by *Buzzwole* with the super effective move *Close Combat*. Even after a lot of investigation, it is hard to determine how this match could have been won by the agent, one possibility might have been switching in a *check* to *Buzzwole* and dynamaxing next turn.

6. Conclusion

(@Samuel: Conclusion bitte lesen) While simple rules like picking the most damaging move are sufficient to defeat *RandomPlayer* most of the time, more complex routines are required in order to rival human players.

ToDo

Previously existing search-based implementations managed to yield good results against baseline agents but are outperformed by the current state of the art DQN-Agent described by Huang and Lee, 2019. Additionally, a way to implement long term planning by modifying the *Minimax*-Algorithm was proposed. The heavy influence of rng results in some games being unbalanced favoring one of the two players heavily and therefore, a large quantity of games is required to properly compare different agents.

HerrGewitter failing to rival *pmariglia* is likely caused by its inferior item prediction algorithm and handling of edge cases. Nevertheless, the agent is capable of defeating less experienced humans. Further possible improvements to the agent include item prediction and the refinement of the *Minimax*-Algorithm by including status and boost. Lastly, a refined dynamaxing-routine and improved predictions of enemy actions are worth investigating.

Pointed out by Huang and Lee, 2019, the DQN-Agent might be improved by introducing “different network architectures or the addition of a recurrent element like an LSTM to better model human memory during the course of a game“. Experiments combining supervised learning and reinforcement learning similar to the *StarCraft II* AI AlphaStar, 2019 developed by *DeepMind* may improve the agent even further.

List of Figures

1.1.	Pikachu, one of the most popular Pokémon (Fandom, 2016)	1
1.2.	Screenshots of Pokémon Red, the earliest Pokémon game. The game was released in 1996 for the Game Boy and has since been re-released multiple times. Image source: (Nintendo, 2022)	2
1.3.	Screenshots of Pokémon Sword, the latest major Pokémon game released for the Nintendo Switch in 2019. While the graphics have evolved, the core concept of the game, exploring and collecting Pokémon, remained the same.	2
2.1.	Damage multiplier for all types. The damage dealt will be multiplied with two if a move is <i>super effective</i> (highlighted green), halved if the defender <i>resists</i> the attacking type (highlighted red) or set to zero if the move has no effect on the attacker (highlighted gray) (Pokémon-Database, 2022)	6
2.2.	Summary of the stats for <i>Charizard</i> . The left column shows the <i>base stats</i> of the Pokémon. To make comparison with other Pokémon easier, the second column shows these stats as bar chart. The last two columns show the possible ranges of final stats with ev and iv taken into account.	8
5.1.	Win rate of the baseline agents at a given board rating. Both agents move have a lower win rate on unfavorable board ratings while winning more games on a positive rating.	34
5.2.	Elo of <i>HerrDonner</i> and <i>HerrGewitter</i> after playing 1400 ranked games respectively	38
5.3.	Smoothed Elo of <i>HerrDonner</i> and <i>HerrGewitter</i> over the course of 1400 ranked games.	39
A.1.	Battle between two players on <i>Pokémon Showdown</i> . The <i>Mewtwo</i> of the player <i>Buckfae</i> is at full health while the opposing <i>Kommo-o</i> has 33% of his health remaining. Below the game window, all possible actions for the player is displayed. A log of previous actions is displayed on the right-hand side.	56
A.2.	The actor-critic neural network used by the authors of Huang and Lee, 2019. Embeddings are used to properly represent different game objects. The network has two outputs: a probability distribution over actions to take, and an estimate of player strength in the current state.	57
A.3.	Possible end game scenario. In the first line of a node are the two active Pokémon, the Pokémon that is expected to faint first is underlined. The lines below shows the remaining hp of all Pokémon for both teams. The last line shows the Value of the given state. A positive value indicates that the player is expected to win. If the node is outlined green, the first player has to bring in a new team member as his active Pokémon fainted, a red node indicates that the opponent has to send a new Pokémon into battle.	58

List of Tables

2.1.	Stage multipliers that are applied for atk, def, spa, spd and spe. For example, a Pokémon with atk raised by two stages will have his existing atk stat multiplied by two (Bulbapedia, 2022l)	13
2.2.	Damage dealt to Pokémon by <i>Stealth Rocks</i> (Bulbapedia, 2022m)	14
2.3.	Damage dealt to Pokémon by <i>Sharp Steel</i> (Bulbapedia, 2021j)	14
3.1.	Denotes how many games the Agent in the column won against the agent in the row RemptonGames, 2021	20
3.2.	Results of other approaches against a random agent (Lee & Togelius, 2017)	21
3.3.	Similar Pokémon within the embedding space of Chen and Lin, 2018	22
3.4.	Average epoch rewards after training convergence for opponent agents (Chen & Lin, 2018). Agent outperforms random and default agent but fails to defeat the <i>Minimax</i> -Agent	23
3.5.	Features used to describe a single Pokémon battle (Huang & Lee, 2019)	24
3.6.	Performance of the agent developed by Huang and Lee, 2019	25
5.1.	Battle results of 1,000 games between the agents. <i>Pmariglia</i> : (pmariglia, 2022) <i>DQN</i> : As described in (Huang & Lee, 2019)	39
5.2.	Performance of both agents against human players. <i>HerrGewitter</i> posed a harder challenge but is unable to gain a lead against <i>Markus</i> , an experienced human player	40

Listings

5.1. Smoothing Elo values	38
-------------------------------------	----

Acronyms

hp hit points
atk attack stat
def defense stat
spa special attack stat
spd special defense stat
spe speed stat
crit critical hit
iv individual values
ev effort values
brn burn
frz freeze
par paralysis
psn poison
slp sleep
rng random number generation
bfs breadth-first search
ppo proximal policy optimization
dmg damage
rps rock-paper-scissors
gen1 generation 1
gen6 generation 6
gen7 generation 7
gen8 generation 8
gen8randbats generation 8 random battle

Bibliography

- alopez247. (2017). Pokémon for data mining and machine learning. Retrieved January 24, 2022, from <https://www.kaggle.com/alopez247/pokemon>
- AlphaStar. (2019). Alphastar: Mastering the real-time strategy game starcraft ii. Retrieved February 3, 2022, from <https://deepmind.com/blog/article/alphastar-mastering-real-time-strategy-game-starcraft-ii>
- Bulbapedia. (2016). Category:moves that use stats from different categories. Retrieved February 2, 2022, from https://bulbapedia.bulbagarden.net/wiki/Category:Moves_that_use_stats_from_different_categories
- Bulbapedia. (2021a). Accuracy. Retrieved February 5, 2022, from <https://bulbapedia.bulbagarden.net/wiki/Accuracy>
- Bulbapedia. (2021b). Air balloon. Retrieved February 2, 2022, from https://bulbapedia.bulbagarden.net/wiki/Air_Balloon
- Bulbapedia. (2021c). Assault vest. Retrieved February 2, 2022, from https://bulbapedia.bulbagarden.net/wiki/Assault_Vest
- Bulbapedia. (2021d). Choice band. Retrieved February 2, 2022, from https://bulbapedia.bulbagarden.net/wiki/Choice_Band
- Bulbapedia. (2021e). Choice scarf. Retrieved February 2, 2022, from https://bulbapedia.bulbagarden.net/wiki/Choice_Scarf
- Bulbapedia. (2021f). Choice specs. Retrieved February 2, 2022, from https://bulbapedia.bulbagarden.net/wiki/Choice_Specs
- Bulbapedia. (2021g). Counter (move). Retrieved February 2, 2022, from [https://bulbapedia.bulbagarden.net/wiki/Counter_\(move\)](https://bulbapedia.bulbagarden.net/wiki/Counter_(move))
- Bulbapedia. (2021h). Court change (move). Retrieved February 2, 2022, from [https://bulbapedia.bulbagarden.net/wiki/Court_Change_\(move\)](https://bulbapedia.bulbagarden.net/wiki/Court_Change_(move))
- Bulbapedia. (2021i). Damage. Retrieved February 2, 2022, from <https://bulbapedia.bulbagarden.net/wiki/Damage>
- Bulbapedia. (2021j). G-max steelsurge (move). Retrieved February 2, 2022, from [https://bulbapedia.bulbagarden.net/wiki/G-Max_Steelsurge_\(move\)](https://bulbapedia.bulbagarden.net/wiki/G-Max_Steelsurge_(move))
- Bulbapedia. (2021k). Grounded. Retrieved February 2, 2022, from <https://bulbapedia.bulbagarden.net/wiki/Grounded>
- Bulbapedia. (2021l). Gyro ball (move). Retrieved February 2, 2022, from [https://bulbapedia.bulbagarden.net/wiki/Gyro_Ball_\(move\)](https://bulbapedia.bulbagarden.net/wiki/Gyro_Ball_(move))
- Bulbapedia. (2021m). Hail (weather condition). Retrieved February 3, 2022, from [https://bulbapedia.bulbagarden.net/wiki/Hail_\(weather_condition\)](https://bulbapedia.bulbagarden.net/wiki/Hail_(weather_condition))
- Bulbapedia. (2021n). Heavy-duty boots. Retrieved February 2, 2022, from https://bulbapedia.bulbagarden.net/wiki/Heavy-Duty_Boots
- Bulbapedia. (2021o). Life orb. Retrieved February 2, 2022, from https://bulbapedia.bulbagarden.net/wiki/Life_Orb
- Bulbapedia. (2021p). Light screen (move). Retrieved February 3, 2022, from [https://bulbapedia.bulbagarden.net/wiki/Light_Screen_\(move\)](https://bulbapedia.bulbagarden.net/wiki/Light_Screen_(move))

- Bulbapedia. (2021q). List of moves that cause entry hazards. Retrieved February 2, 2022, from https://bulbapedia.bulbagarden.net/wiki/List_of_moves_that_cause_entry_hazards
- Bulbapedia. (2021r). List of pokémon by base stats (generation vii). Retrieved February 2, 2022, from [https://bulbapedia.bulbagarden.net/wiki/List_of_Pok%C3%A9mon_by_base_stats_\(Generation_VII\)](https://bulbapedia.bulbagarden.net/wiki/List_of_Pok%C3%A9mon_by_base_stats_(Generation_VII))
- Bulbapedia. (2021s). Magic room (move). Retrieved February 3, 2022, from [https://bulbapedia.bulbagarden.net/wiki/Magic_Room_\(move\)](https://bulbapedia.bulbagarden.net/wiki/Magic_Room_(move))
- Bulbapedia. (2021t). Sandstorm (weather condition). Retrieved February 3, 2022, from [https://bulbapedia.bulbagarden.net/wiki/Sandstorm_\(weather_condition\)](https://bulbapedia.bulbagarden.net/wiki/Sandstorm_(weather_condition))
- Bulbapedia. (2021u). Switcheroo (move). Retrieved February 2, 2022, from [https://bulbapedia.bulbagarden.net/wiki/Switcheroo_\(move\)](https://bulbapedia.bulbagarden.net/wiki/Switcheroo_(move))
- Bulbapedia. (2021v). Toxic spikes (move). Retrieved February 2, 2022, from [https://bulbapedia.bulbagarden.net/wiki/Toxic_Spikes_\(move\)](https://bulbapedia.bulbagarden.net/wiki/Toxic_Spikes_(move))
- Bulbapedia. (2021w). Trick (move). Retrieved February 2, 2022, from [https://bulbapedia.bulbagarden.net/wiki/Trick_\(move\)](https://bulbapedia.bulbagarden.net/wiki/Trick_(move))
- Bulbapedia. (2021x). Weakness policy. Retrieved February 2, 2022, from https://bulbapedia.bulbagarden.net/wiki/Weakness_Policy
- Bulbapedia. (2021y). Wonder room (move). Retrieved February 3, 2022, from [https://bulbapedia.bulbagarden.net/wiki/Wonder_Room_\(move\)](https://bulbapedia.bulbagarden.net/wiki/Wonder_Room_(move))
- Bulbapedia. (2022a). Critical hit. Retrieved February 2, 2022, from https://bulbapedia.bulbagarden.net/wiki/Critical_hit
- Bulbapedia. (2022b). Defog (move). Retrieved February 2, 2022, from [https://bulbapedia.bulbagarden.net/wiki/Defog_\(move\)](https://bulbapedia.bulbagarden.net/wiki/Defog_(move))
- Bulbapedia. (2022c). Extreme speed (move). Retrieved February 3, 2022, from [https://bulbapedia.bulbagarden.net/wiki/Extreme_Speed_\(move\)](https://bulbapedia.bulbagarden.net/wiki/Extreme_Speed_(move))
- Bulbapedia. (2022d). Fire punch (move). Retrieved February 2, 2022, from [https://bulbapedia.bulbagarden.net/wiki/Fire_Punch_\(move\)](https://bulbapedia.bulbagarden.net/wiki/Fire_Punch_(move))
- Bulbapedia. (2022e). Focus sash. Retrieved February 2, 2022, from https://bulbapedia.bulbagarden.net/wiki/Focus_Sash
- Bulbapedia. (2022f). Gravity (move). Retrieved February 3, 2022, from [https://bulbapedia.bulbagarden.net/wiki/Gravity_\(move\)](https://bulbapedia.bulbagarden.net/wiki/Gravity_(move))
- Bulbapedia. (2022g). Pokémon champion. Retrieved February 3, 2022, from https://bulbapedia.bulbagarden.net/wiki/Pok%C3%A9mon_Champion
- Bulbapedia. (2022h). Rapid spin (move). Retrieved February 2, 2022, from [https://bulbapedia.bulbagarden.net/wiki/Rapid_Spin_\(move\)](https://bulbapedia.bulbagarden.net/wiki/Rapid_Spin_(move))
- Bulbapedia. (2022i). Reflect (move). Retrieved February 3, 2022, from [https://bulbapedia.bulbagarden.net/wiki/Reflect_\(move\)](https://bulbapedia.bulbagarden.net/wiki/Reflect_(move))
- Bulbapedia. (2022j). Roost (move). Retrieved February 5, 2022, from [https://bulbapedia.bulbagarden.net/wiki/Roost_\(move\)](https://bulbapedia.bulbagarden.net/wiki/Roost_(move))
- Bulbapedia. (2022k). Spikes (move). Retrieved February 2, 2022, from [https://bulbapedia.bulbagarden.net/wiki/Spikes_\(move\)](https://bulbapedia.bulbagarden.net/wiki/Spikes_(move))
- Bulbapedia. (2022l). Stat. Retrieved February 2, 2022, from <https://bulbapedia.bulbagarden.net/wiki/Stat>
- Bulbapedia. (2022m). Stealth rock (move). Retrieved February 2, 2022, from [https://bulbapedia.bulbagarden.net/wiki/Stealth_Rock_\(move\)](https://bulbapedia.bulbagarden.net/wiki/Stealth_Rock_(move))
- Bulbapedia. (2022n). Sticky web (move). Retrieved February 2, 2022, from [https://bulbapedia.bulbagarden.net/wiki/Sticky_Web_\(move\)](https://bulbapedia.bulbagarden.net/wiki/Sticky_Web_(move))
- Bulbapedia. (2022o). U-turn (move). Retrieved February 2, 2022, from [https://bulbapedia.bulbagarden.net/wiki/U-turn_\(move\)](https://bulbapedia.bulbagarden.net/wiki/U-turn_(move))
- Bulbapedia. (2022p). Weather. Retrieved February 2, 2022, from <https://bulbapedia.bulbagarden.net/wiki/Weather>

- Bulbapedia. (2022q). Will-o-wisp (move). Retrieved February 3, 2022, from [https://bulbapedia.bulbagarden.net/wiki/Will-O-Wisp_\(move\)](https://bulbapedia.bulbagarden.net/wiki/Will-O-Wisp_(move))
- Chen, K., & Lin, E. (2018). Gotta train 'em all: Learning to play pokémon showdown with reinforcement learning. Retrieved February 3, 2022, from https://cs230.stanford.edu/projects_fall_2018/reports/12447633.pdf
- Fandom. (2016). Ashs pikachu. Retrieved January 27, 2022, from https://pokemon.fandom.com/de/wiki/Ashs_Pikachu?action=history
- Grover, A., & Leskovec, J. (2016). Node2vec: Scalable feature learning for networks.
- Honko. (2022). Pokémon damage calculator (generation vi). Retrieved January 24, 2022, from <https://calc.pokemonshowdown.com/randoms.html?gen=6&mode=randoms>
- Huang, D., & Lee, S. (2019). A self-play policy optimization approach to battling pokémon. *2019 IEEE Conference on Games (CoG)*, 1–4. <https://doi.org/10.1109/CIG.2019.8848014>
- Konda, V., & Tsitsiklis, J. (2000). Actor-critic algorithms. In S. Solla, T. Leen, & K. Müller (Eds.), *Advances in neural information processing systems* (p. 7). MIT Press. <https://proceedings.neurips.cc/paper/1999/file/6449f44a102fde848669bdd9eb6b76fa-Paper.pdf>
- Lee, S., & Togelius, J. (2017). Showdown ai competition. *2017 IEEE Conference on Computational Intelligence and Games (CIG)*, 191–198. <https://doi.org/10.1109/CIG.2017.8080435>
- MattL. (2014). What are checks and counters? Retrieved February 3, 2022, from <https://www.smogon.com/smog/issue32/checks-and-counters>
- Mikolov, T., Sutskever, I., Chen, K., Corrado, G., & Dean, J. (2013). Distributed representations of words and phrases and their compositionality.
- Mitarai, R. (2021). Random battles competitive guide. Retrieved February 2, 2022, from <https://www.smogon.com/articles/randbats-guide-gen8>
- Nintendo. (2019a). Das kampf-stadion. Retrieved February 3, 2022, from <https://swordshield.pokemon.com/de-de/gameplay/pokemon-battle-stadium/>
- Nintendo. (2019b). Pokémon schwert. Retrieved February 3, 2022, from <https://www.nintendo.de/Spiele/Nintendo-Switch/Pokemon-Schwert-1522111.html>
- Nintendo. (2022). Pokémon rote edition. Retrieved February 3, 2022, from <https://www.nintendo.de/Spiele/Game-Boy/Pokemon-Rote-Edition-266109.html>
- OpenAI. (2018). Openai five. Retrieved February 2, 2022, from <https://blog.openai.com/openai-five/>
- pmariglia. (2022). Github/pmariglia-showdown. Retrieved February 2, 2022, from <https://github.com/pmariglia/showdown>
- Pokémon-Database. (2022). Pokémon types & type chart. Retrieved February 2, 2022, from <https://pokemondb.net/type>
- RemptonGames. (2021). Programming ai for pokemon showdown + bot battle royale! Retrieved February 3, 2022, from <https://www.youtube.com/watch?v=rhvj7CmTRkg>
- Sahovic, H. (2022a). Github/poke-env. Retrieved February 2, 2022, from <https://github.com/hsahovic/poke-env>
- Sahovic, H. (2022b). Poke-env: Baselines.py. Retrieved February 2, 2022, from https://github.com/hsahovic/poke-env/blob/master/src/poke_env/player/baselines.py
- Schulman, J., Moritz, P., Levine, S., Jordan, M., & Abbeel, P. (2018). High-dimensional continuous control using generalized advantage estimation.
- Schulman, J., Wolski, F., Dhariwal, P., Radford, A., & Klimov, O. (2017). Proximal policy optimization algorithms.
- Smogon. (2022a). Github/damage-calc. Retrieved February 2, 2022, from <https://github.com/smogon/damage-calc>
- Smogon. (2022b). Github/pokemon-showdown. Retrieved February 3, 2022, from <https://github.com/smogon/pokemon-showdown>

- TheThirdBuild. (2021). How an a.i. is becoming the world's best pokemon player. Retrieved February 3, 2022, from <https://www.youtube.com/watch?v=rhvj7CmTRkg>
- Veekun. (2021). Parasect. Retrieved February 2, 2022, from <https://veekun.com/dex/pokemon/parasect>
- Wikipedia. (2022a). Expectiminimax — Wikipedia, the free encyclopedia [[Online; accessed 24-January-2022]]. <http://en.wikipedia.org/w/index.php?title=Expectiminimax&oldid=952355961>
- Wikipedia. (2022b). Pokémon. Retrieved January 27, 2022, from <https://en.wikipedia.org/wiki/Pok%C3%A9mon>

Appendix

A. Figures

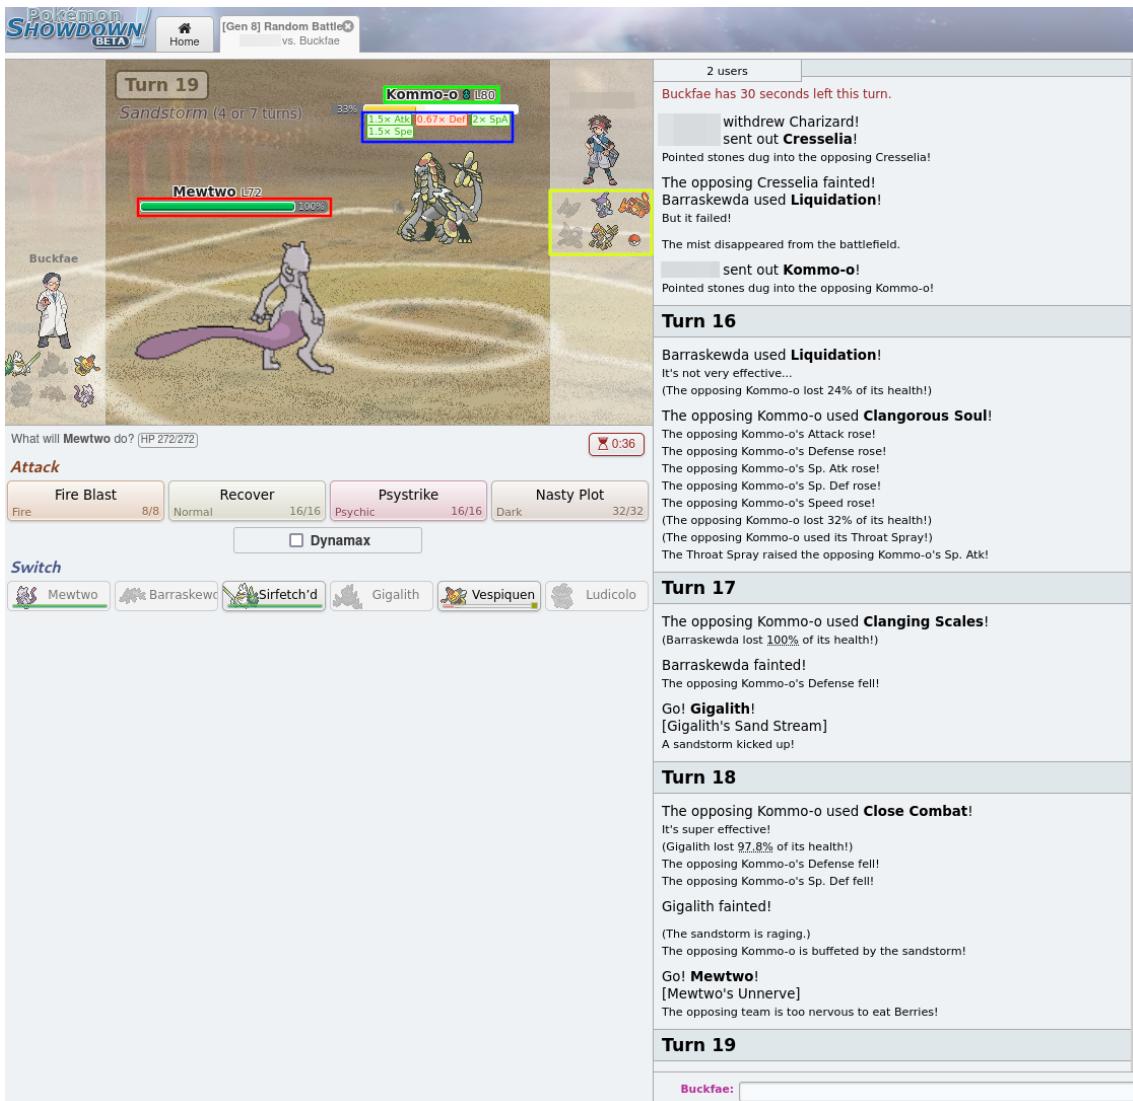


Figure A.1.: Battle between two players on *Pokémon Showdown*. The *Mewtwo* of the player *Buckfae* is at full health while the opposing *Kommo-o* has 33% of his health remaining. Below the game window, all possible actions for the player is displayed. A log of previous actions is displayed on the right-hand side.

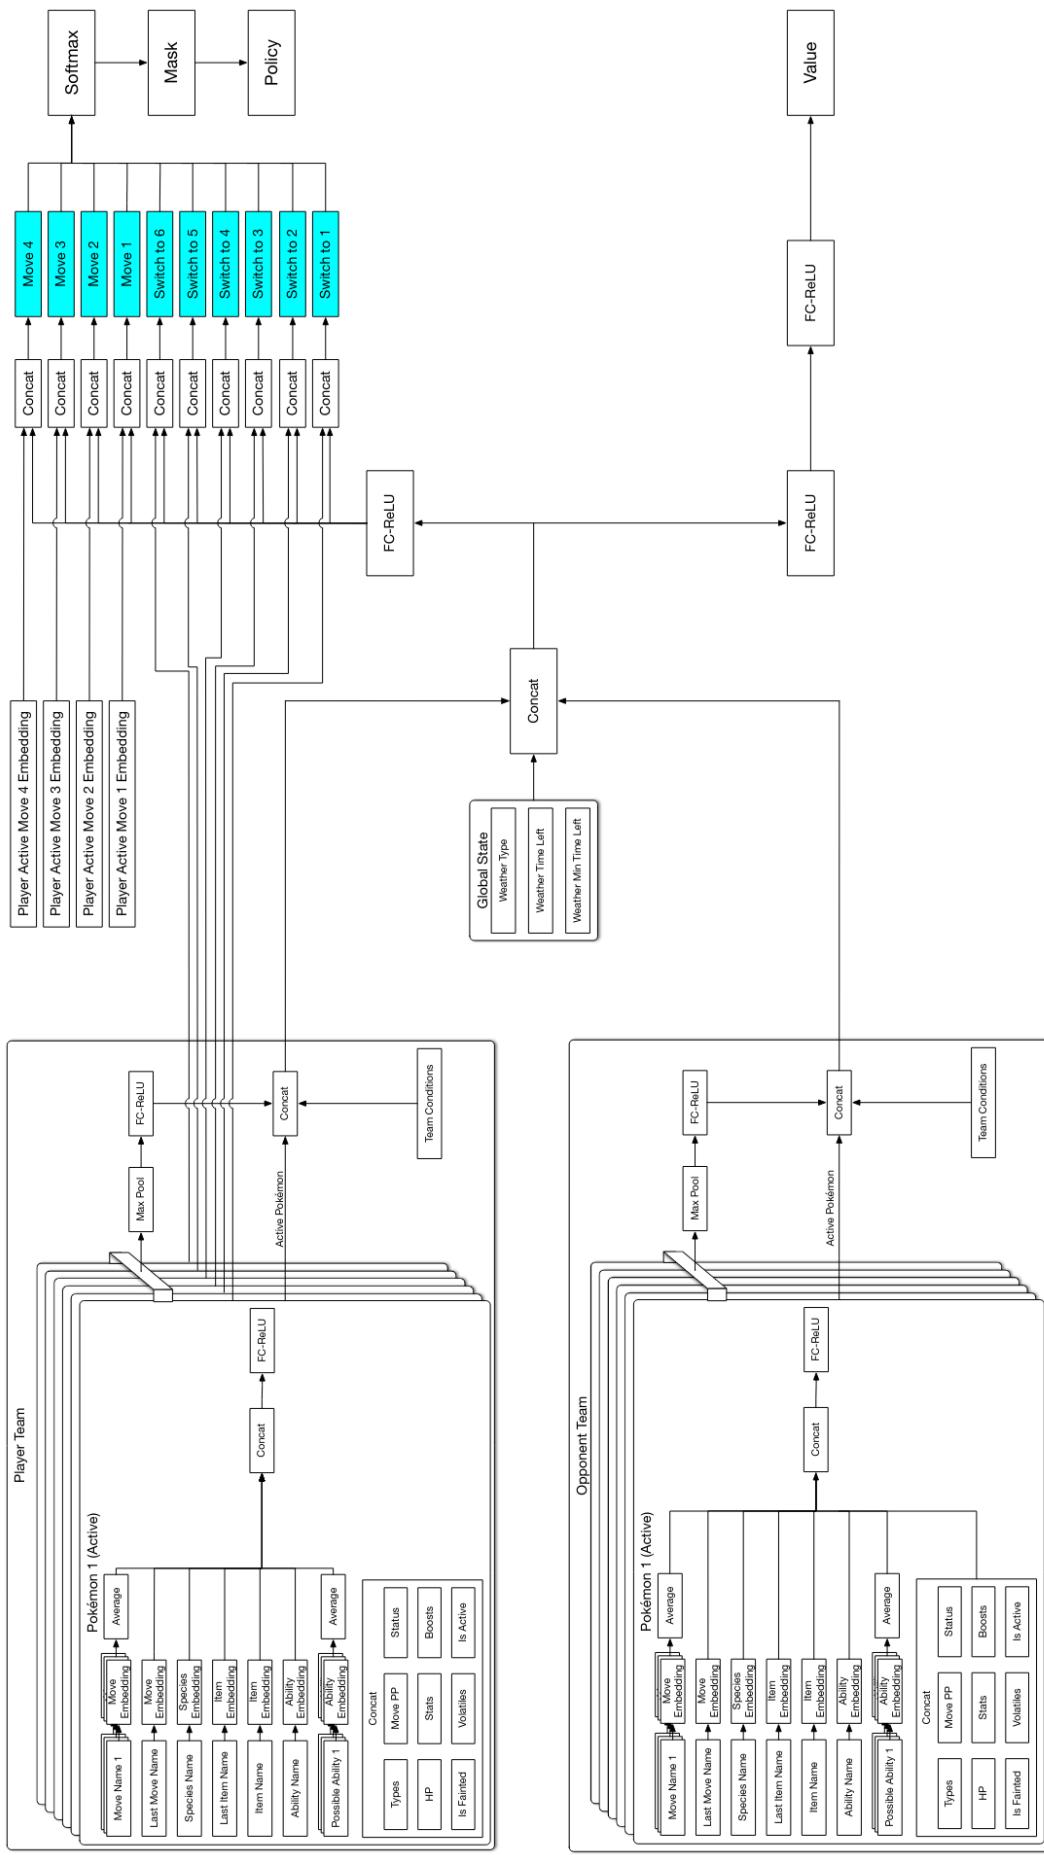


Figure A.2: The actor-critic neural network used by the authors of Huang and Lee, 2019. Embeddings are used to properly represent different game objects. The network has two outputs: a probability distribution over actions to take, and an estimate of player strength in the current state.

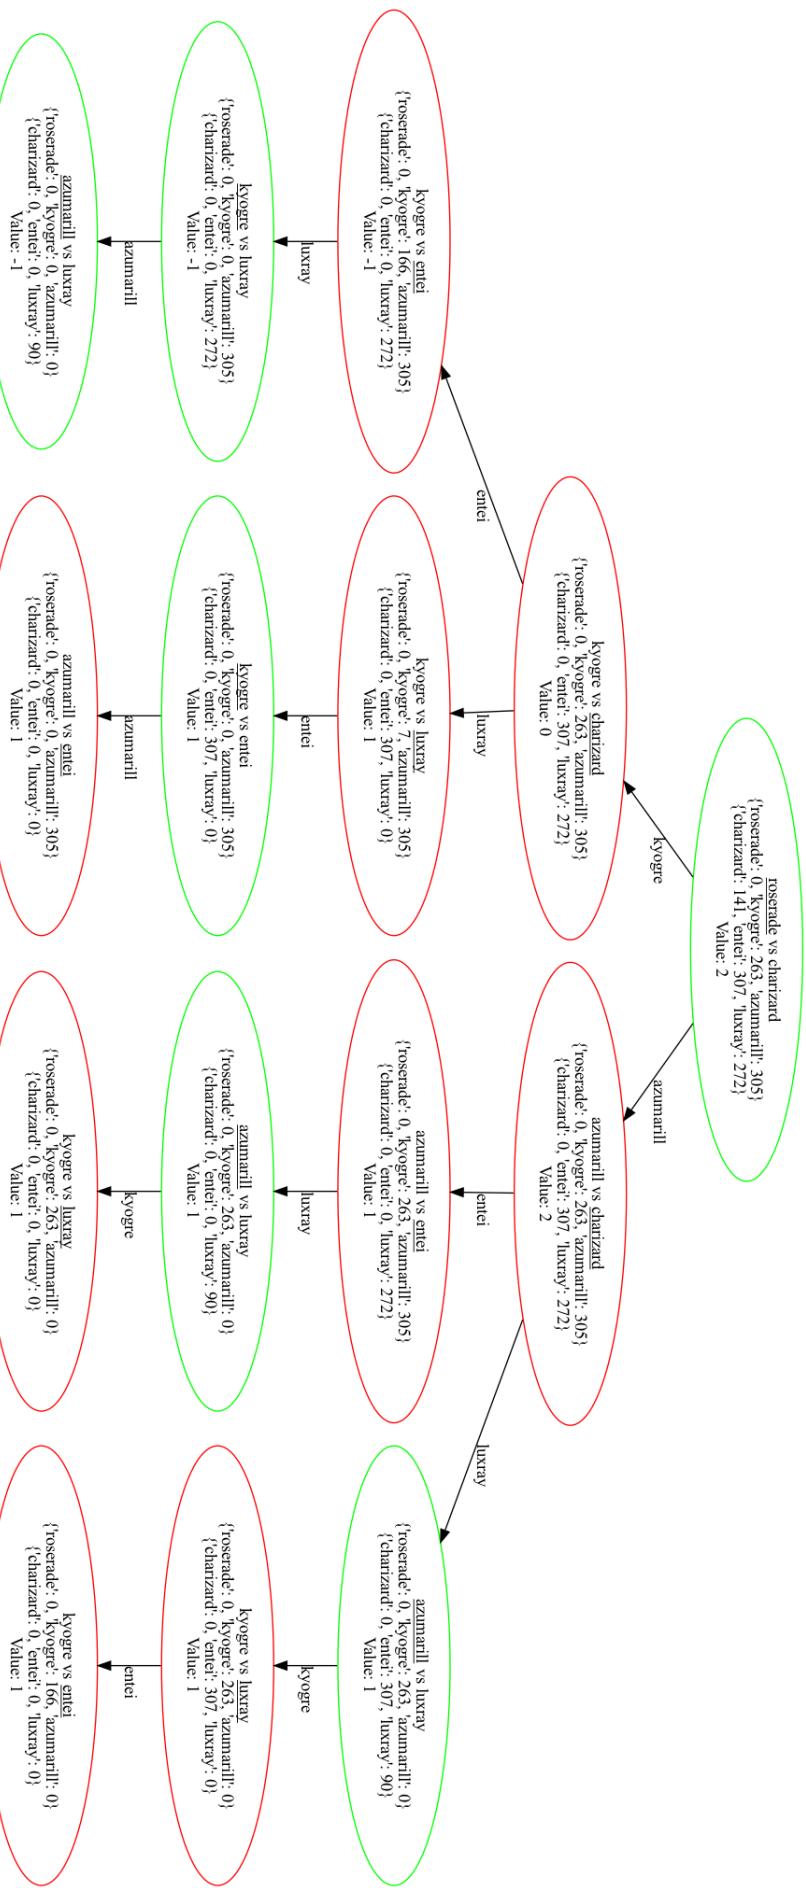


Figure A.3.: Possible end game scenario. In the first line of a node are the two active Pokémons, the Pokémon that is expected to faint first is underlined. The lines below shows the remaining hp of all Pokémons for both teams. The last line shows the Value of the given state. A positive value indicates that the player is expected to win. If the node is outlined green, the first player has to bring in a new team member as his active Pokémon fainted, a red node indicates that the opponent has to send a new Pokémon into battle.

Ich versichere, dass ich die vorstehende Arbeit selbstständig und ohne fremde Hilfe angefertigt und mich keiner anderer als der in den beigefügten Verzeichnissen angegebenen Hilfsmittel bedient habe. Alle Textstellen, die wörtlich oder sinngemäß aus Veröffentlichungen Dritter entnommen wurden, sind als solche kenntlich gemacht. Alle Quellen, die dem World Wide Web entnommen oder in einer digitalen Form verwendet wurden, sind der Arbeit beigefügt.

Weitere Personen waren an der geistigen Leistung der vorliegenden Arbeit nicht beteiligt. Insbesondere habe ich nicht die Hilfe eines Ghostwriters oder einer Ghostwriting-Agentur in Anspruch genommen. Dritte haben von mir weder unmittelbar noch mittelbar Geld oder geldwerte Leistungen für Arbeiten erhalten, die im Zusammenhang mit dem Inhalt der vorgelegten Arbeit stehen.

Der Durchführung einer elektronischen Plagiatsprüfung stimme ich hiermit zu. Die eingereichte elektronische Fassung der Arbeit ist vollständig. Mir ist bewusst, dass nachträgliche Ergänzungen ausgeschlossen sind.

Die Arbeit wurde bisher keiner anderen Prüfungsbehörde vorgelegt und auch nicht veröffentlicht. Ich bin mir bewusst, dass eine unwahre Erklärung zur Versicherung der selbstständigen Leistungserbringung rechtliche Folgen haben kann.

Würzburg, 07. February 2022

.....
(Julian Schubert)